

Tensors (with PyTorch)

Sept. 2023

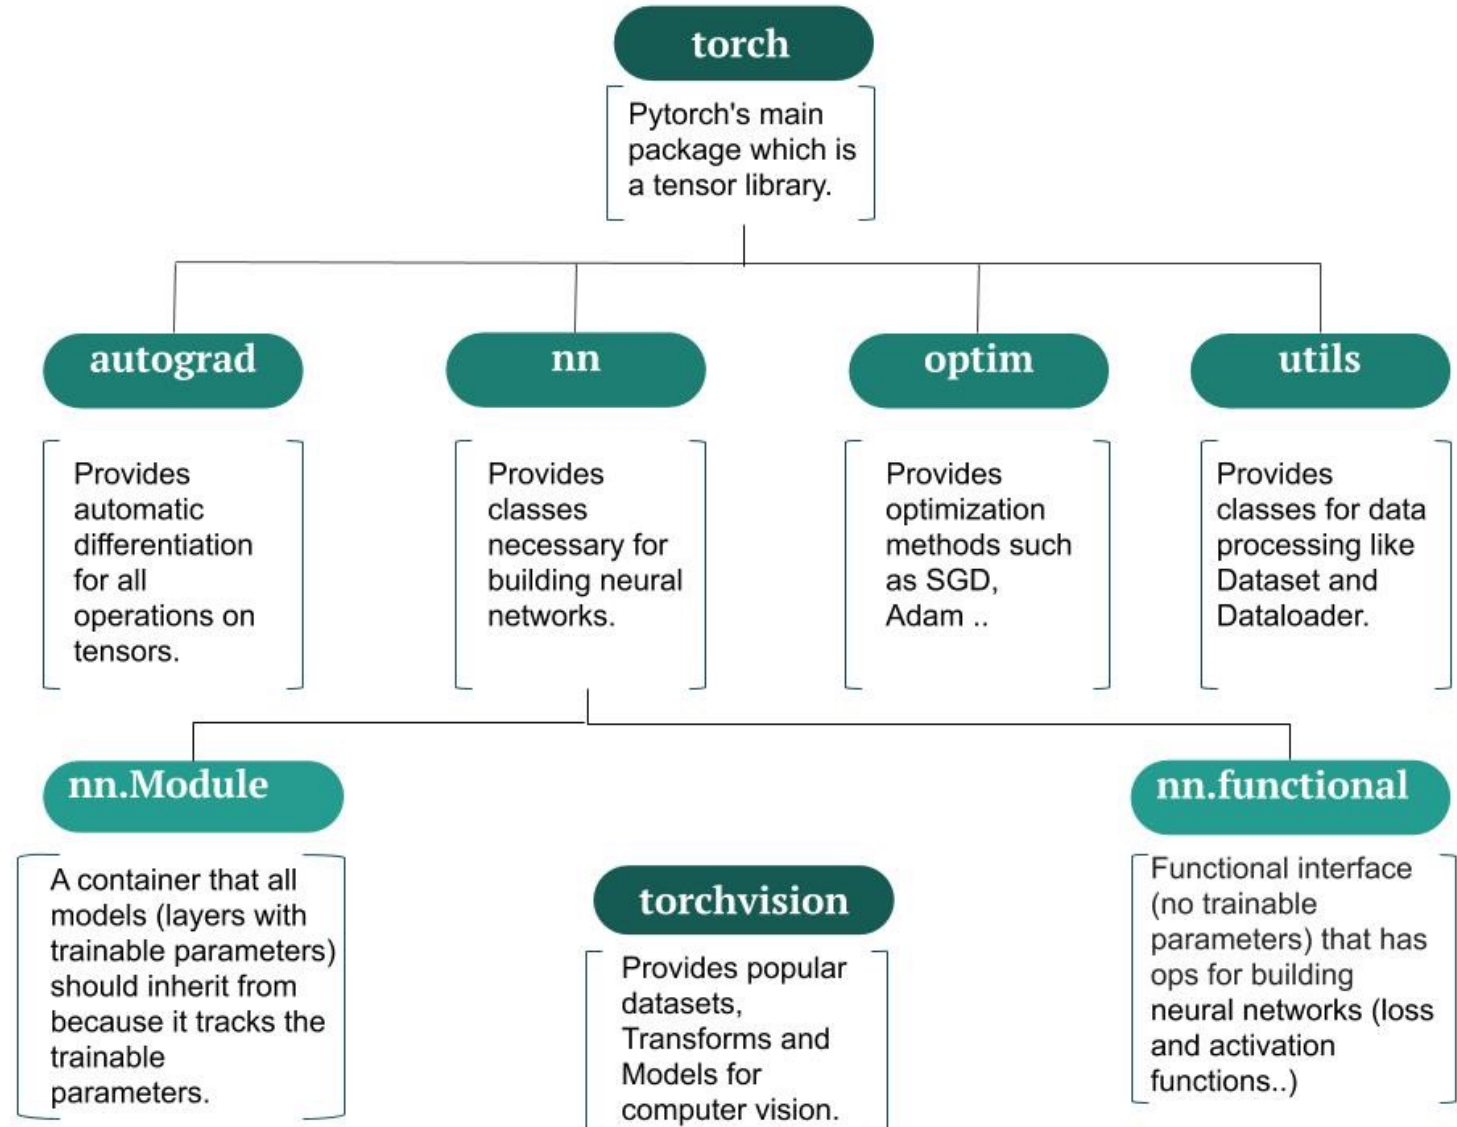
<http://link.koreatech.ac.kr>

PyTorch and Tensors 101

PyTorch Packages

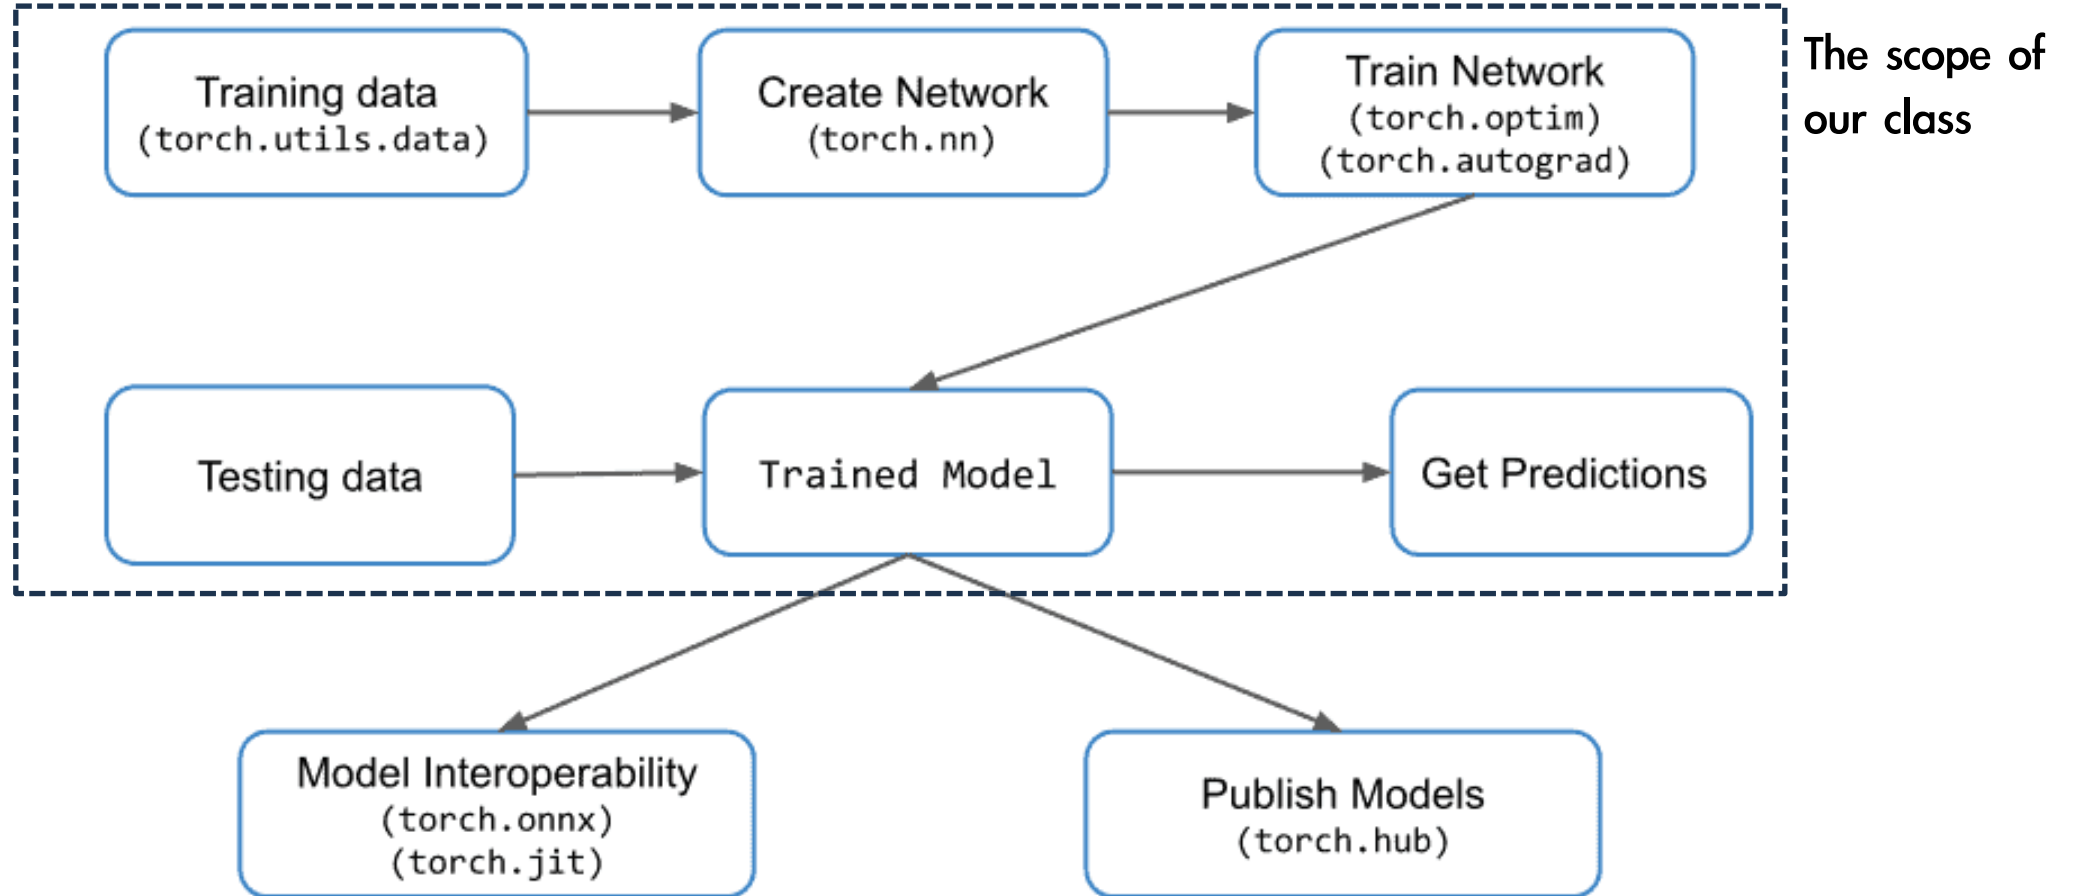
◆ PyTorch Packages

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
```



PyTorch Packages

◆ Basic PyTorch Workflow



Tensor

◆ Tensors

- Multi-dimensional array containing elements of a single data type
 - PyTorch tensors are very similar to NumPy ndarrays
 - PyTorch tensors are used on a GPU as well (this is not the case with NumPy arrays)


tensor = multidimensional array

vector




$$\mathbf{v} \in \mathbb{R}^{64}$$

matrix



$$\mathbf{X} \in \mathbb{R}^{8 \times 8}$$

tensor



$$\mathbf{X} \in \mathbb{R}^{4 \times 4 \times 4}$$

3

`A = torch.tensor(3)`

0 Dimensions

1.0
2.0
3.0

`B =
torch.tensor([1.0, 2.0, 3.0])`

1 Dimension

1.0 2.0
3.0 4.0

`C = torch.tensor([
[1.0, 2.0],
[3.0, 4.0]])`

2 Dimensions

5.0 6.0
7.0 8.0
1.0 2.0
3.0 4.0

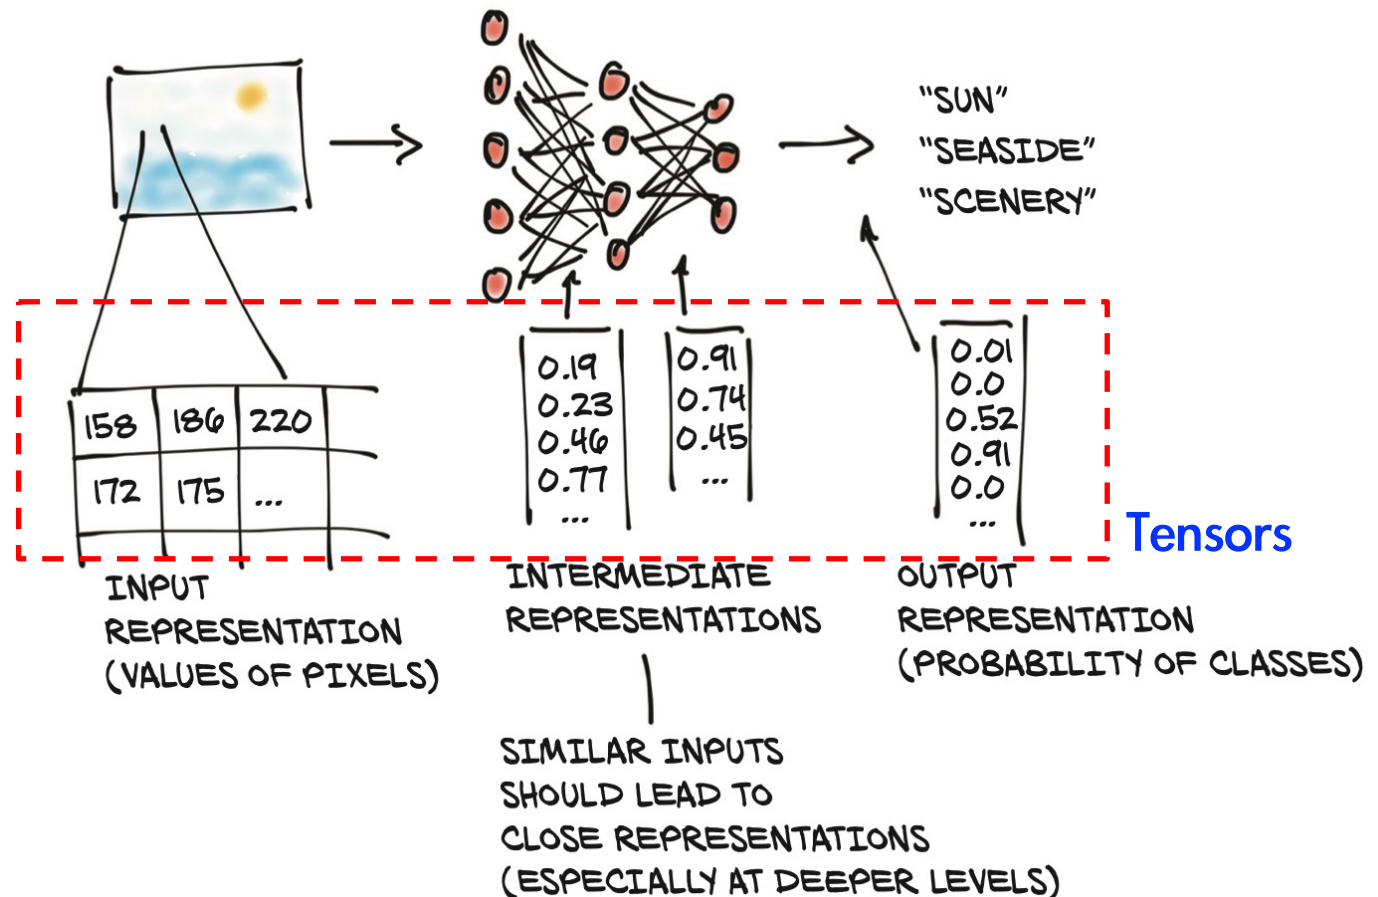
`D = torch.tensor([
[[1., 2.], [3., 4.]],
[[5., 6.], [7., 8.]]])`

3 Dimensions

Tensor

◆ Tensors

- A deep neural network learns how to transform an input representation (or input tensor) to an output representation (or output tensor)



Tensor

◇ Shape of Tensors

— check with `t.shape` or `t.size()`

- Check with `.shape()`




5

(5,)



dim 0



3

5


(3, 5)



dim 0



dim 1



4

5

3

(4, 5, 3)



dim 0



dim 1



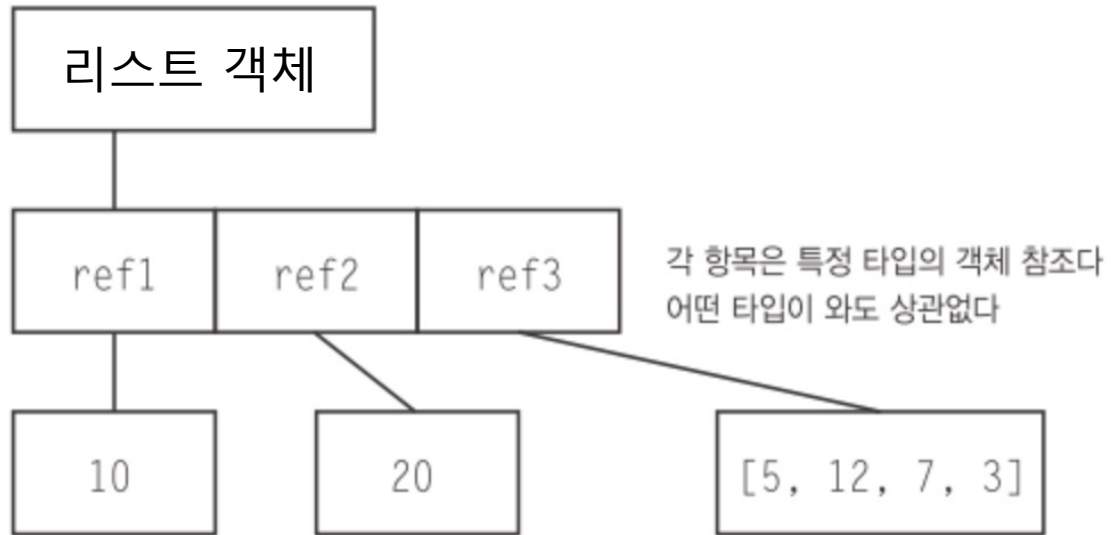
dim 2

Note: **dim** in PyTorch == **axis** in NumPy

Tensor

◆ Tensor Storage

파이썬 list 객체 저장 방식



파이썬 Numpy & PyTorch Tensor 객체 저장 방식



Tensor Initialization

Tensor Initialization

◆ Tensor Data Types

- Almost always `torch.float32` or `torch.int64` are used

- `torch.Tensor` is an alias for the default tensor type `torch.FloatTensor`

Data type	dtype	CPU tensor	GPU tensor
32-bit floating point	<code>torch.float32</code> or <code>torch.float</code>	<code>torch.FloatTensor</code>	<code>torch.cuda.FloatTensor</code>
64-bit floating point	<code>torch.float64</code> or <code>torch.double</code>	<code>torch.DoubleTensor</code>	<code>torch.cuda.DoubleTensor</code>
16-bit floating point	<code>torch.float16</code> or <code>torch.half</code>	<code>torch.HalfTensor</code>	<code>torch.cuda.HalfTensor</code>
8-bit integer (unsigned)	<code>torch.uint8</code>	<code>torch.ByteTensor</code>	<code>torch.cuda.ByteTensor</code>
8-bit integer (signed)	<code>torch.int8</code>	<code>torch.CharTensor</code>	<code>torch.cuda.CharTensor</code>
16-bit integer (signed)	<code>torch.int16</code> or <code>torch.short</code>	<code>torch.ShortTensor</code>	<code>torch.cuda.ShortTensor</code>
32-bit integer (signed)	<code>torch.int32</code> or <code>torch.int</code>	<code>torch.IntTensor</code>	<code>torch.cuda.IntTensor</code>
64-bit integer (signed)	<code>torch.int64</code> or <code>torch.long</code>	<code>torch.LongTensor</code>	<code>torch.cuda.LongTensor</code>
Boolean	<code>torch.bool</code>	<code>torch.BoolTensor</code>	<code>torch.cuda.BoolTensor</code>

Tensor Initialization

◆ Tensor initialization directly from data

— `torch.Tensor` class & `torch.tensor` function

- Every tensor has `device`, `dtype`, `shape`, and `requires_grad` attributes
- `torch.Tensor` is an alias for the default tensor type `torch.FloatTensor`

```
# torch.Tensor class
t1 = torch.Tensor([1, 2, 3], device='cpu')
print(t1.dtype) # >>> torch.float32
print(t1.device) # >>> cpu
print(t1.requires_grad) # >>> False
print(t1.size()) # torch.Size([3])
print(t1.shape) # torch.Size([3])
```

```
# if you have gpu device
t1_cuda = t1.to(torch.device('cuda'))
# or you can use shorthand
t1_cuda = t1.cuda()

t1_cpu = t1.cpu()
```

```
# torch.tensor function
t2 = torch.tensor([1, 2, 3], device='cpu')
print(t2.dtype) # >>> torch.int64
print(t2.device) # >>> cpu
print(t2.requires_grad) # >>> False
print(t2.size()) # torch.Size([3])
print(t2.shape) # torch.Size([3])
```

```
# if you have gpu device
t2_cuda = t2.to(torch.device('cuda'))
# or you can use shorthand
t2_cuda = t2.cuda()

t2_cpu = t2.cpu()
```

Tensor Initialization

◆ Tensor initialization directly from data

— torch.Tensor & torch.tensor always copies the given data

```
import torch

l1 = [1, 2, 3]
t1 = torch.Tensor(l1)

l2 = [1, 2, 3]
t2 = torch.tensor(l2)

l3 = [1, 2, 3]
t3 = torch.as_tensor(l3)

l1[0] = 100
l2[0] = 100
l3[0] = 100

print(t1) # >>> tensor([1., 2., 3.])
print(t2) # >>> tensor([1, 2, 3])
print(t3) # >>> tensor([1, 2, 3])
```

```
import torch
import numpy as np

l4 = np.array([1, 2, 3])
t4 = torch.Tensor(l4)

l5 = np.array([1, 2, 3])
t5 = torch.tensor(l5)

l6 = np.array([1, 2, 3])
t6 = torch.as_tensor(l6)

l4[0] = 100
l5[0] = 100
l6[0] = 100

print(t4) # >>> tensor([1., 2., 3.])
print(t5) # >>> tensor([1, 2, 3])
print(t6) # >>> tensor([100, 2, 3])
```

If you have a numpy array and want to avoid a copy, use `torch.as_tensor()`.

Tensor Initialization

◆ Tensor initialization with constant values

- `torch.zeros(*size)`
- `torch.zeros_like(input_tensor)`
- `torch.ones(*size)`
- `torch.ones_like(input_tensor)`
- `torch.empty(*size)`
 - filled with uninitialized data
- `torch.eye(n)`
 - return identity matrix with size $n \times n$

```
import torch

t1 = torch.ones(size=(5,)) # or torch.ones(5)
t1_like = torch.ones_like(input=t1)
print(t1) # >>> tensor([1., 1., 1., 1., 1.])
print(t1_like) # >>> tensor([1., 1., 1., 1., 1.])

t2 = torch.zeros(size=(6,)) # or torch.zeros(6)
t2_like = torch.zeros_like(input=t2)
print(t2) # >>> tensor([0., 0., 0., 0., 0., 0.])
print(t2_like) # >>> tensor([0., 0., 0., 0., 0., 0.])

t3 = torch.empty(size=(4,)) # or torch.zeros(4)
t3_like = torch.empty_like(input=t3)
print(t3) # >>> tensor([0., 0., 0., 0.])
print(t3_like) # >>> tensor([0., 0., 0., 0.])

t4 = torch.eye(n=3)
print(t4) # >>> tensor([[1., 0., 0.],
#                       [0., 1., 0.],
#                       [0., 0., 1.]])
```

Tensor Initialization

◆ Tensor initialization with random values (1/2)

— `torch.randint(low=0, high, size, ...)`



- filled with integer values that generate uniformly between low (inclusive) and high (exclusive)

— `torch.rand(*size, ...)`

- filled with float values ranging between 0 and 1 from a uniform distribution

— `torch.randn(*size, ...)`

- filled with a random float number from a standard normal distribution with a mean 0 and a variance of 1



```
import torch
```

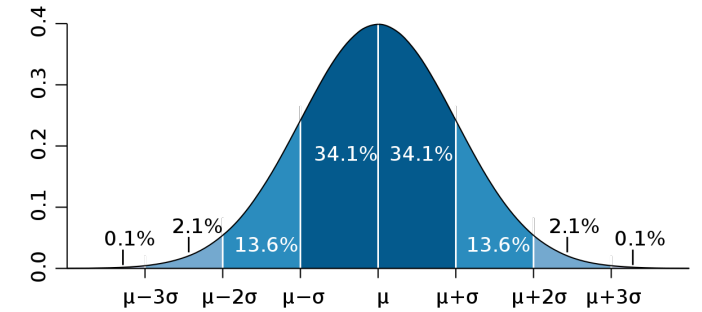
```
t1 = torch.randint(low=10, high=20, size=(1, 2))
t2 = torch.rand(size=(1, 3)) # or torch.rand(1, 3)
t3 = torch.randn(size=(1, 3)) # or torch.randn(1, 3)
print(t1) # >>> tensor([[11, 15]])
print(t2) # >>> tensor([[0.3704, 0.3847, 0.2096]])
print(t3) # >>> tensor([[-1.1459, -0.4099, -0.6727]])
```

Tensor Initialization

◆ Tensor initialization with random values (2/2)

— `torch.normal(mean, std, size, ...)`

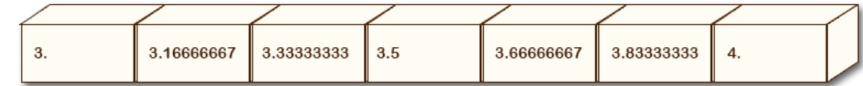
- filled with a random float number from a normal distribution whose mean and standard deviation are given



— `torch.linspace(start, end, steps, ...)`

- returns a one-dimensional tensor between start and end and it is equally spaced based on steps

`torch.linspace(3, 4, steps=7)`



— `torch.arange(start=0, end, steps=1,...)`

- returns a one-dimensional tensor with values from the interval `[start, end)` (start: inclusive, end: exclusive) taken with common difference `step` beginning from start

```
import torch
```

```
t4 = torch.normal(mean=10.0, std=1.0, size=(3, 2))
t5 = torch.linspace(start=0.0, end=5.0, steps=3)
t6 = torch.arange(5)
print(t4)  # >>> tensor([[11.4634, 12.1609],
                        #          [ 9.8654,  9.6206],
                        #          [11.4928, 10.3766]])
print(t5)  # >>> tensor([0.0000, 2.5000, 5.0000])
print(t6)  # >>> tensor([0, 1, 2, 3, 4])
```

Tensor Initialization

◆ Random seed

— `torch.manual_seed(seed)`

- it will set the seed of the random number generator to a fixed value, so that when you call for example `torch.rand(2)`, the results will be reproducible

```
import torch

torch.manual_seed(1729)
random1 = torch.rand(2, 3)
print(random1)
# >>> tensor([[0.3126, 0.3791, 0.3087],
#             [0.0736, 0.4216, 0.0691]])

random2 = torch.rand(2, 3)
print(random2)
# >>> tensor([[0.2332, 0.4047, 0.2162],
#             [0.9927, 0.4128, 0.5938]])
```

```
import torch

torch.manual_seed(1729)
random3 = torch.rand(2, 3)
print(random3)
# >>> tensor([[0.3126, 0.3791, 0.3087],
#             [0.0736, 0.4216, 0.0691]])

random4 = torch.rand(2, 3)
print(random4)
# >>> tensor([[0.2332, 0.4047, 0.2162],
#             [0.9927, 0.4128, 0.5938]])
```


Tensor Operations

Tensor Type Conversion

◆ Tensor type conversion

```
import torch
a = torch.ones((2, 3))
print(a.dtype) # >>> torch.float32

b = torch.ones((2, 3), dtype=torch.int16)
print(b)
# >>> tensor([[1, 1, 1],
#             [1, 1, 1]], dtype=torch.int16)

c = torch.rand((2, 3), dtype=torch.float64) * 20.
print(c)
# >>> tensor([[16.0387, 16.0498, 11.2471],
#             [18.4285, 12.5039, 0.7902]],
#             dtype=torch.float64)

d = c.to(torch.int32)
print(d)
# >>> tensor([[16, 16, 11],
#             [18, 12, 0]], dtype=torch.int32)
```

broadcasting

```
double_d = torch.ones(10, 2, dtype=torch.double)
short_e = torch.tensor([[1, 2]], dtype=torch.short)

double_d = torch.zeros(10, 2).double()
short_e = torch.ones(10, 2).short()

double_d = torch.zeros(10, 2).to(torch.double)
short_e = torch.ones(10, 2).to(dtype=torch.short)

print(double_d.dtype) # >>> torch.float64
print(short_e.dtype) # >>> torch.int16

double_f = torch.rand(5, dtype=torch.double)
short_g = double_f.to(torch.short)
print((double_f * short_g).dtype)
# >>> torch.float64
```

Tensor Operations

◆ Tensor operations

- [NOTE] Element-wise operations and Broadcasting
- `Torch.add(other)` or `torch.add(input, other)`
- `Torch.sub(other)` or `torch.sub(input, other)`
- `Torch.mul(other)` or `torch.mul(input, other)`
- `Torch.div(other)` or `torch.div(input, other)`
- `Torch.dot(other)` or `torch.dot(input, other)`
- `Torch.mm(other)` or `torch.mm(input, other)`
- `Torch.bmm(other)` or `torch.bmm(input, other)`
- `Torch.matmul(other)` or `torch.matmul(Tensor, Tensor)`
- `Torch.pow(exponent)` or `torch.pow(input, exponent)`

Tensor Operations

◆ Element-wise operations

– `torch.add(input, other) ⇔ input + other`

– `torch.sub(input, other) ⇔ input - other`

$$\begin{pmatrix} V_1 \\ V_2 \\ V_3 \\ V_4 \end{pmatrix} \odot \begin{pmatrix} W_1 \\ W_2 \\ W_3 \\ W_4 \end{pmatrix} = \begin{pmatrix} V_1 * W_1 \\ V_2 * W_2 \\ V_3 * W_3 \\ V_4 * W_4 \end{pmatrix}$$

$$\begin{pmatrix} X_{1,1} & X_{1,2} \\ X_{2,1} & X_{2,2} \end{pmatrix} \odot \begin{pmatrix} Y_{1,1} & Y_{1,2} \\ Y_{2,1} & Y_{2,2} \end{pmatrix} = \begin{pmatrix} X_{1,1} * Y_{1,1} & X_{1,2} * Y_{1,2} \\ X_{2,1} * Y_{2,1} & X_{2,2} * Y_{2,2} \end{pmatrix}$$

```
import torch

t1 = torch.ones(size=(2, 3))
t2 = torch.ones(size=(2, 3))

t3 = torch.add(t1, t2)
t4 = t1 + t2
print(t3) # >>> tensor([[2., 2., 2.],
                        #          [2., 2., 2.]])
print(t4) # >>> tensor([[2., 2., 2.],
                        #          [2., 2., 2.]])
```

```
import torch

t1 = torch.ones(size=(2, 3))
t2 = torch.ones(size=(2, 3))

t5 = torch.sub(t1, t2)
t6 = t1 - t2
print(t3) # >>> tensor([[0., 0., 0.],
                        #          [0., 0., 0.]])
print(t4) # >>> tensor([[0., 0., 0.],
                        #          [0., 0., 0.]])
```

Tensor Operations

◆ Element-wise operations

- `torch.mul(input, other) ⇔ input * other`
- `torch.div(input, other) ⇔ input / other`

a ₁	a ₂	a ₃
a ₄	a ₅	a ₆
a ₇	a ₈	a ₉

 \times

b ₁	b ₂	b ₃
b ₄	b ₅	b ₆
b ₇	b ₈	b ₉

 $=$

a ₁ b ₁	a ₂ b ₂	a ₃ b ₃
a ₄ b ₄	a ₅ b ₅	a ₆ b ₆
a ₇ b ₇	a ₈ b ₈	a ₉ b ₉

```
import torch

t1 = torch.ones(size=(2, 3))
t2 = torch.ones(size=(2, 3))

t7 = torch.mul(t1, t2)
t8 = t1 * t2
print(t7) # >>> tensor([[1., 1., 1.],
                        #          [1., 1., 1.]])
print(t8) # >>> tensor([[1., 1., 1.],
                        #          [1., 1., 1.]])
```

```
import torch

t1 = torch.ones(size=(2, 3))
t2 = torch.ones(size=(2, 3))

t9 = torch.div(t1, t2)
t10 = t1 / t2
print(t9) # >>> tensor([[1., 1., 1.],
                        #          [1., 1., 1.]])
print(t10) # >>> tensor([[1., 1., 1.],
                        #          [1., 1., 1.]])
```

Tensor Operations

◆ Multiple matrix multiplications (1/2)

— `torch.mul(input, other)`

- performs a element-wise multiplication with broadcasting
- (Tensor) by (Tensor or Number)

— `torch.dot(input, other)`

- Computes the dot product of two 1D tensors
- (1D tensor) by (1D tensor)

— `torch.mm(input, other)`

- performs a matrix multiplication without broadcasting
- (2D tensor, $n \times m$) by (2D tensor, $m \times p$) $\rightarrow (n \times p)$

— `torch.bmm(input, other)`

- performs a batch matrix multiplication without broadcasting
- ($b \times n \times m$) by ($b \times m \times p$) $\rightarrow (b \times n \times p)$

```
import torch
t1 = torch.dot(
    torch.tensor([2, 3]),
    torch.tensor([2, 1])
)
print(t1, t1.size())
# >>> tensor(7) torch.Size([1])

t2 = torch.randn(2, 3)
t3 = torch.randn(3, 2)
t4 = torch.mm(t2, t3)
print(t4, t4.size())
# >>> tensor([[2.8075, 2.3547],
#             [2.4447, 2.0428]]),
#             torch.Size([2, 2])

t5 = torch.randn(10, 3, 4)
t6 = torch.randn(10, 4, 5)
t7 = torch.bmm(t5, t6)
print(t7.size())
# >>> tensor.Size([10, 3, 5])
```

Tensor Operations

◆ Multiple matrix multiplications (2/2)

— `torch.matmul(Tensor, Tensor)`

⇔ `Tensor @ Tensor`

- Various tensor product with broadcasting
- it uses different modes depending on the input tensor shapes
- Supported operation
 - dot product
 - matrix product
 - batched matrix products

```
import torch

# vector x vector: dot
t1 = torch.randn(3)
t2 = torch.randn(3)
print(torch.matmul(t1, t2).size()) # >>> torch.Size([])

# matrix x vector: broadcasted dot
t3 = torch.randn(3, 4)
t4 = torch.randn(4)
print(torch.matmul(t3, t4).size()) # >>> torch.Size([3])

# batched matrix x vector: broadcasted dot
t5 = torch.randn(10, 3, 4)
t6 = torch.randn(4)
print(torch.matmul(t5, t6).size()) # >>> torch.Size([10, 3])

# batched matrix x batched matrix: bmm
t7 = torch.randn(10, 3, 4)
t8 = torch.randn(10, 4, 5)
print(torch.matmul(t7, t8).size()) # >>> torch.Size([10, 3, 5])

# batched matrix x matrix: bmm
t9 = torch.randn(10, 3, 4)
t10 = torch.randn(4, 5)
print(torch.matmul(t9, t10).size()) # >>> torch.Size([10, 3, 5])
```

Tensor Operations

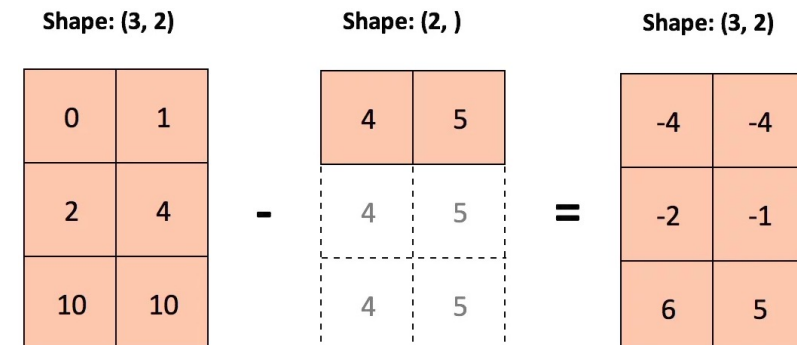

◆ Broadcasting

- the smaller tensor is "broadcast" across the larger tensor so that they have compatible shapes and element-wise operations can be performed on them
 - In broadcasting, the smaller array is found, the new axes are added as per the larger array, and data is added appropriately to the transformed array

```
import torch

t1 = torch.tensor([1.0, 2.0, 3.0])
t2 = 2.0
print(t1 * t2) # >>> tensor([2., 4., 6.])

t3 = torch.tensor([[0, 1], [2, 4], [10, 10]])
t4 = torch.tensor([4, 5])
print(t3 - t4) # >>> tensor([[ -4,  -4],
                              #      [-2,  -1],
                              #      [ 6,   5]])
```



Tensor Operations

◆ Broadcasting Examples

```
import torch
t5 = torch.tensor([[1., 2.], [3., 4.]])

print(t5 + 2.0)  # t5.add(2.0)
# >>> tensor([[3., 4.],
#             [5., 6.]])

print(t5 - 2.0)  # t5.sub(2.0)
# >>> tensor([[ -1.,  0.],
#             [ 1.,  2.]])

print(t5 * 2.0)  # t5.mul(2.0)
# >>> tensor([[2., 4.],
#             [6., 8.]])

print(t5 / 2.0)  # t5.div(2.0)
# >>> tensor([[0.5000, 1.0000],
#             [1.5000, 2.0000]])
```

```
import torch

def normalize(x):
    return x / 255

t6 = torch.randn(3, 28, 28)

print(normalize(t6).size())
# >>> torch.Size([3, 28, 28])
```

Tensor Operations

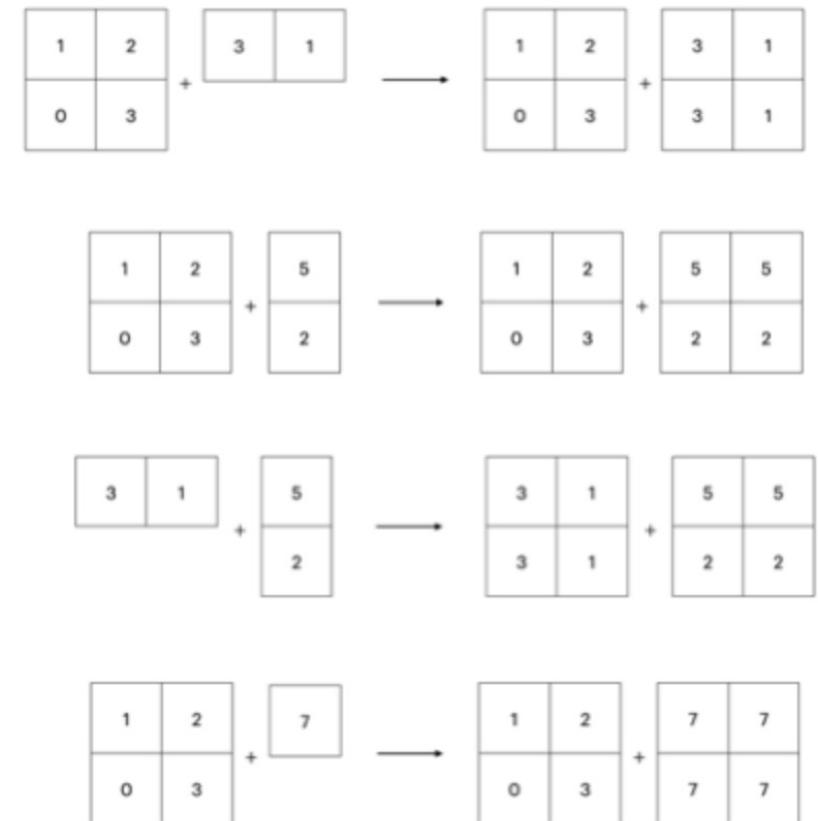
◆ Broadcasting Rules

- Each tensor must have at least one dimension - no empty tensors.
- Comparing the dimension sizes of the two tensors, going from last to first:
 - Each dimension must be equal, or
 - One of the dimensions must be of size 1, or
 - The dimension does not exist in one of the tensors

```
import torch

t7 = torch.tensor([[1, 2], [0, 3]]) # torch.Size([2, 2])
t8 = torch.tensor([[3, 1]])         # torch.Size([1, 2])
t9 = torch.tensor([[5], [2]])        # torch.Size([2, 1])
t10 = torch.tensor([7])              # torch.Size([1])

print(t7 + t8)    # >>> tensor([[4, 3], [3, 4]])
print(t7 + t9)    # >>> tensor([[6, 7], [2, 5]])
print(t8 + t9)    # >>> tensor([[8, 6], [5, 3]])
print(t7 + t10)   # >>> tensor([[8, 9], [7, 10]])
```



Tensor Operations

◆ Broadcasting Rules

- Comparing the dimension sizes of the two tensors, going from last to first

```
import torch

t11 = torch.ones(4, 3, 2)
t12 = t11 * torch.rand( 3, 2) # 3rd & 2nd dims identical to t11, dim 0 absent
print(t12.shape) # >>> torch.Size([4, 3, 2])

t13 = torch.ones(4, 3, 2)
t14 = t13 * torch.rand( 3, 1) # 3rd dim = 1, 2nd dim is identical to t13
print(t14.shape) # >>> torch.Size([4, 3, 2])

t15 = torch.ones(4, 3, 2)
t16 = t15 * torch.rand( 1, 2) # 3rd dim is identical to t15, 2nd dim is 1
print(t16.shape) # >>> torch.Size([4, 3, 2])

t17 = torch.ones(5, 3, 4, 1)
t18 = torch.rand( 3, 1, 1) # 2nd dim is identical to t17, 3rd and 4th dims are 1
print((t17 + t18).size()) # >>> torch.Size([5, 3, 4, 1])
```

Tensor Operations

◆ Broadcasting Rules

- Comparing the dimension sizes of the two tensors, going from last to first

```
import torch

t19 = torch.empty(5, 1, 4, 1)
t20 = torch.empty( 3, 1, 1)
print((t19 + t20).size())    # torch.Size([5, 3, 4, 1])

t21 = torch.empty(1)
t22 = torch.empty(3, 1, 7)
print((t21 + t22).size())    # torch.Size([3, 1, 7])

t23 = torch.ones(3, 3, 3)
t24 = torch.ones(3, 1, 3)
print((t23 + t24).size())    # torch.Size([3, 3, 3])

t25 = torch.empty(5, 2, 4, 1)
t26 = torch.empty(3, 1, 1)
print((t25 + t26).size())    # RuntimeError: The size of tensor a (2) must match
                             # the size of tensor b (3) at non-singleton dimension
```

Tensor Operations

◆ Pow() operation and its broadcasting

— `Torch.pow(exponent)` or `torch.pow(input, exponent)`

```
import torch

t27 = torch.ones(4) * 5
print(t27) # >>> tensor([ 5, 5, 5, 5])

t28 = torch.pow(t27, 2)
print(t28) # >>> tensor([ 25, 25, 25, 25])

exp = torch.arange(1., 5.) # tensor([ 1., 2., 3., 4.])
a = torch.arange(1., 5.) # tensor([ 1., 2., 3., 4.])

t29 = torch.pow(a, exp)
print(t29) # >>> tensor([ 1., 4., 27., 256.] )
```

Tensor Indexing & Slicing

Tensor Indexing & Slicing

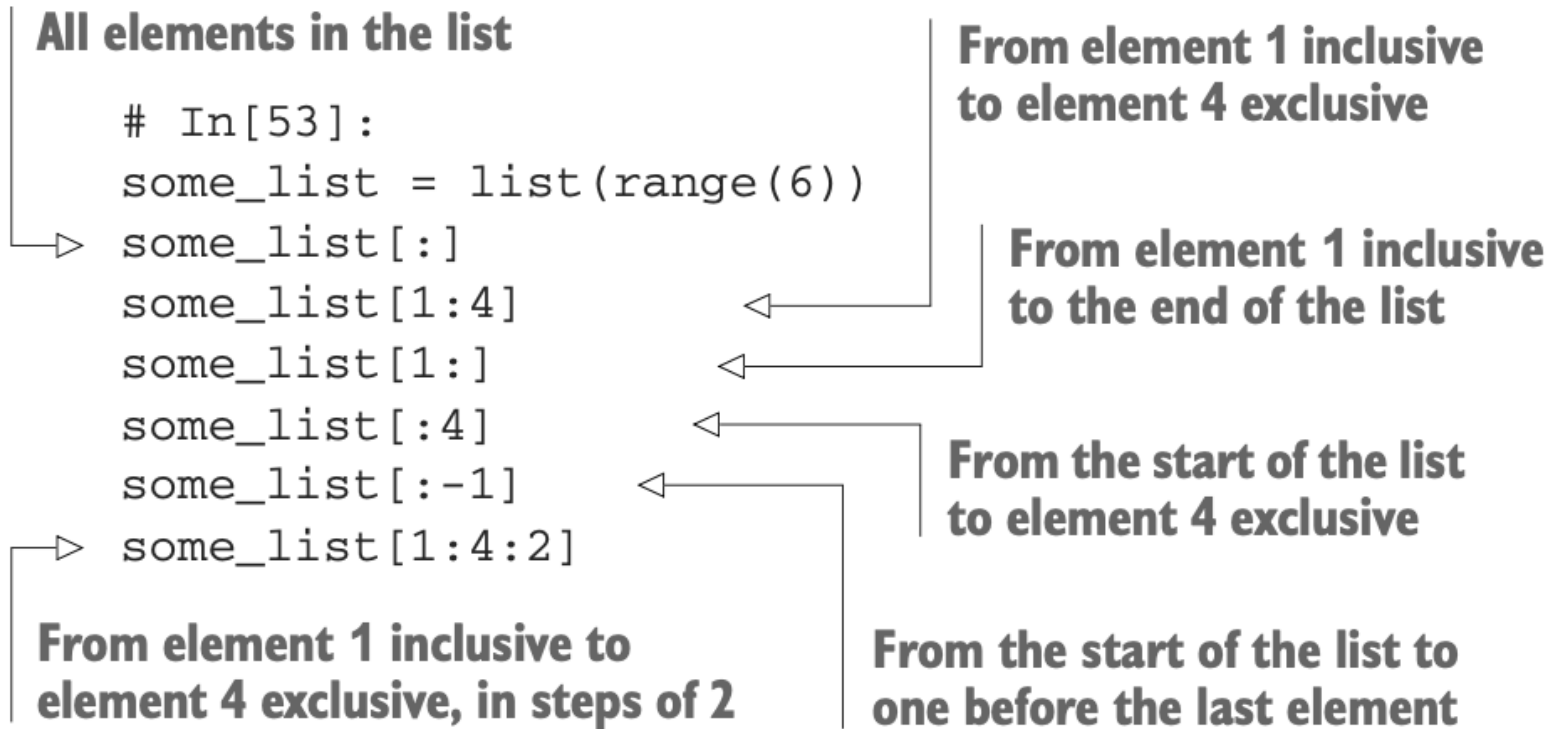
◆ Tensor Indexing & Slicing

- accessing or retrieving specific elements or slices from a tensor
- common methods
 - **Basic indexing**: accessing individual elements of a tensor by specifying the indices along each dimension
 - **Slicing**: Extracting a specific sub-tensor from a larger tensor by specifying ranges along each dimension
 - **Advanced indexing**: Some libraries support more advanced indexing methods, such as using boolean masks or arrays of indices to select elements from a tensor based on certain conditions

Tensor Indexing & Slicing

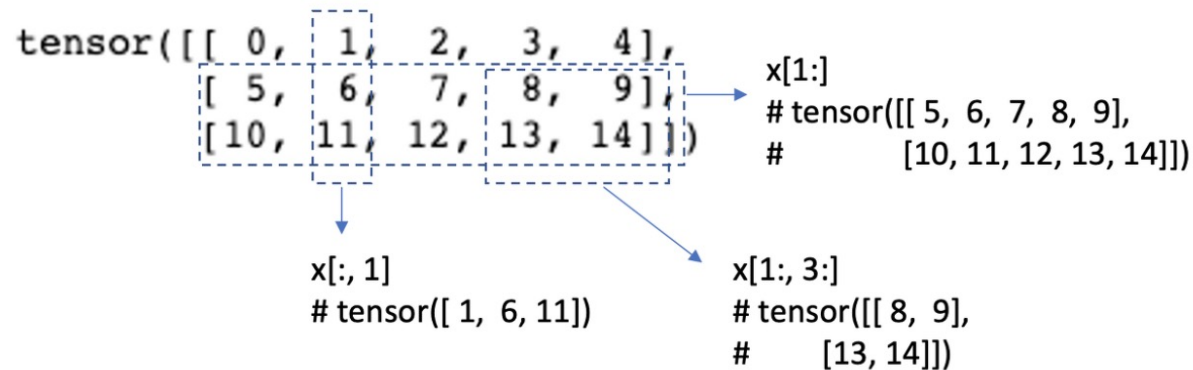
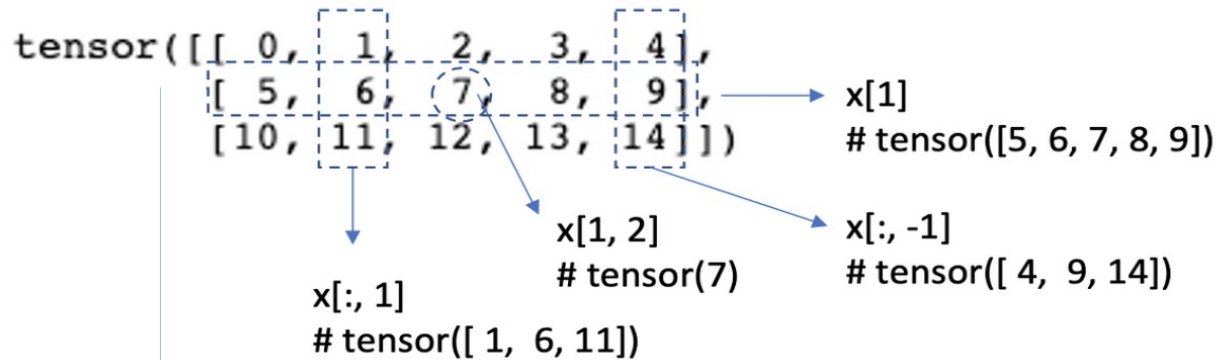
◆ Tensor Indexing & Slicing

- They are almost same as the classical python indexing & slicing



Tensor Indexing & Slicing

◆ Tensor Indexing & Slicing



```
import torch
```

```
x = torch.tensor(
    [[ 0,  1,  2,  3,  4],
     [ 5,  6,  7,  8,  9],
     [10, 11, 12, 13, 14]]
)
```

```
print(x[1]) # >>> tensor([5, 6, 7, 8, 9])
print(x[:, 1]) # >>> tensor([1, 6, 11])
print(x[1, 2]) # >>> tensor(7)
print(x[:, -1]) # >>> tensor([4, 9, 14])
```

```
print(x[1:]) # >>> tensor([[ 5,  6,  7,  8,  9],
                        #      [10, 11, 12, 13, 14]])
print(x[1:, 3:]) # >>> tensor([[ 8,  9],
                        #      [13, 14]])
```

Tensor Indexing & Slicing

◆ Indexing & Slicing - Examples

```
import torch

y = torch.zeros((6, 6))
y[1:4, 2] = 1
print(y)

# >>> tensor([[0., 0., 0., 0., 0., 0.],
#             [0., 0., 1., 0., 0., 0.],
#             [0., 0., 1., 0., 0., 0.],
#             [0., 0., 1., 0., 0., 0.],
#             [0., 0., 0., 0., 0., 0.],
#             [0., 0., 0., 0., 0., 0.]])

print(y[1:4, 1:4])

# >>> tensor([[0., 1., 0.],
#             [0., 1., 0.],
#             [0., 1., 0.]])
```

```
import torch

z = torch.tensor(
    [[1, 2, 3, 4],
     [2, 3, 4, 5],
     [5, 6, 7, 8]]
)
print(z[:2]) # >>> tensor([[1, 2, 3, 4],
#                        [2, 3, 4, 5]])

print(z[1:, 1:3]) # >>> tensor([[3, 4],
#                               [6, 7]])

print(z[:, 1:]) # >>> tensor([[2, 3, 4],
#                             [3, 4, 5],
#                             [6, 7, 8]])

z[1:, 1:3] = 0
print(z) # >>> tensor([[1, 2, 3, 4],
#                     [2, 0, 0, 5],
#                     [5, 0, 0, 8]])
```

Tensor Reshaping

Tensor Reshaping

◆ Tensor Reshaping Methods

- `Tensor.view(*shape)`
- `Tensor.reshape(*shape)` or `torch.reshape(input, shape)`
- `Tensor.unsqueeze(dim)` or `torch.unsqueeze(input, dim)`
- `Tensor.squeeze(dim)` or `torch.squeeze(input, dim)`
- `Tensor.flatten(start_dim=0, end_dim=-1)` or
`torch.flatten(input, start_dim=0, end_dim=-1)`
- `Tensor.permute(*dims)` or `torch.permute(input, dims)`
- `Tensor.transpose(dim0, dim1)` or `torch.transpose(input, dim0, dim1)`
- `Tensor.t()` or `torch.t()`

Tensor Reshaping

◆ `torch.view(input, *shape)` &
`torch.reshape(input, shape)`

- change the shape of a tensor without modifying its data
- The returned tensor will share the underling data with the original tensor
- But, `torch.reshape()` may return a copy for the non-contiguous tensor

```
import torch
import torch

t1 = torch.tensor([[1, 2, 3], [4, 5, 6]])
t2 = t1.view(3, 2)      # Shape becomes (3, 2)
t3 = t1.reshape(1, 6)   # Shape becomes (1, 6)
print(t2)
# >>> tensor([[1, 2],
#              [3, 4],
#              [5, 6]])
print(t3)
# >>> tensor([[1, 2, 3, 4, 5, 6]])

t4 = torch.arange(8).view(2, 4)
t5 = torch.arange(6).view(2, 3)
print(t4)
# >>> tensor([[0, 1, 2, 3],
#              [4, 5, 6, 7]])
print(t5)
# >>> tensor([[0, 1, 2],
#              [3, 4, 5]])
```

Tensor Reshaping

◆ `torch.unsqueeze()`

— `torch.unsqueeze()` adds a new dimension to the tensor at the specified position

Unsqueeze: expand a new dimension

```
>>> x = torch.zeros([2, 3])
```

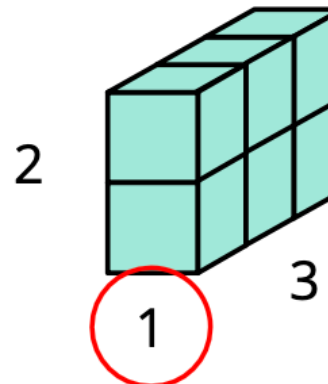
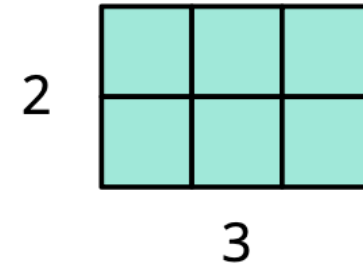
```
>>> x.shape
```

```
torch.Size([2, 3])
```

```
>>> x = x.unsqueeze(1) (dim = 1)
```

```
>>> x.shape
```

```
torch.Size([2, 1, 3])
```



Tensor Reshaping

◆ `torch.squeeze()`

- `torch.squeeze()` method removes all dimensions of size 1 or the specified dimension if its size is 1

Squeeze: remove the specified dimension with length = 1

```
>>> x = torch.zeros([1, 2, 3])
```

```
>>> x.shape
```



```
torch.Size([1, 2, 3])
```

```
>>> x = x.squeeze(0)
```

(dim = 0)

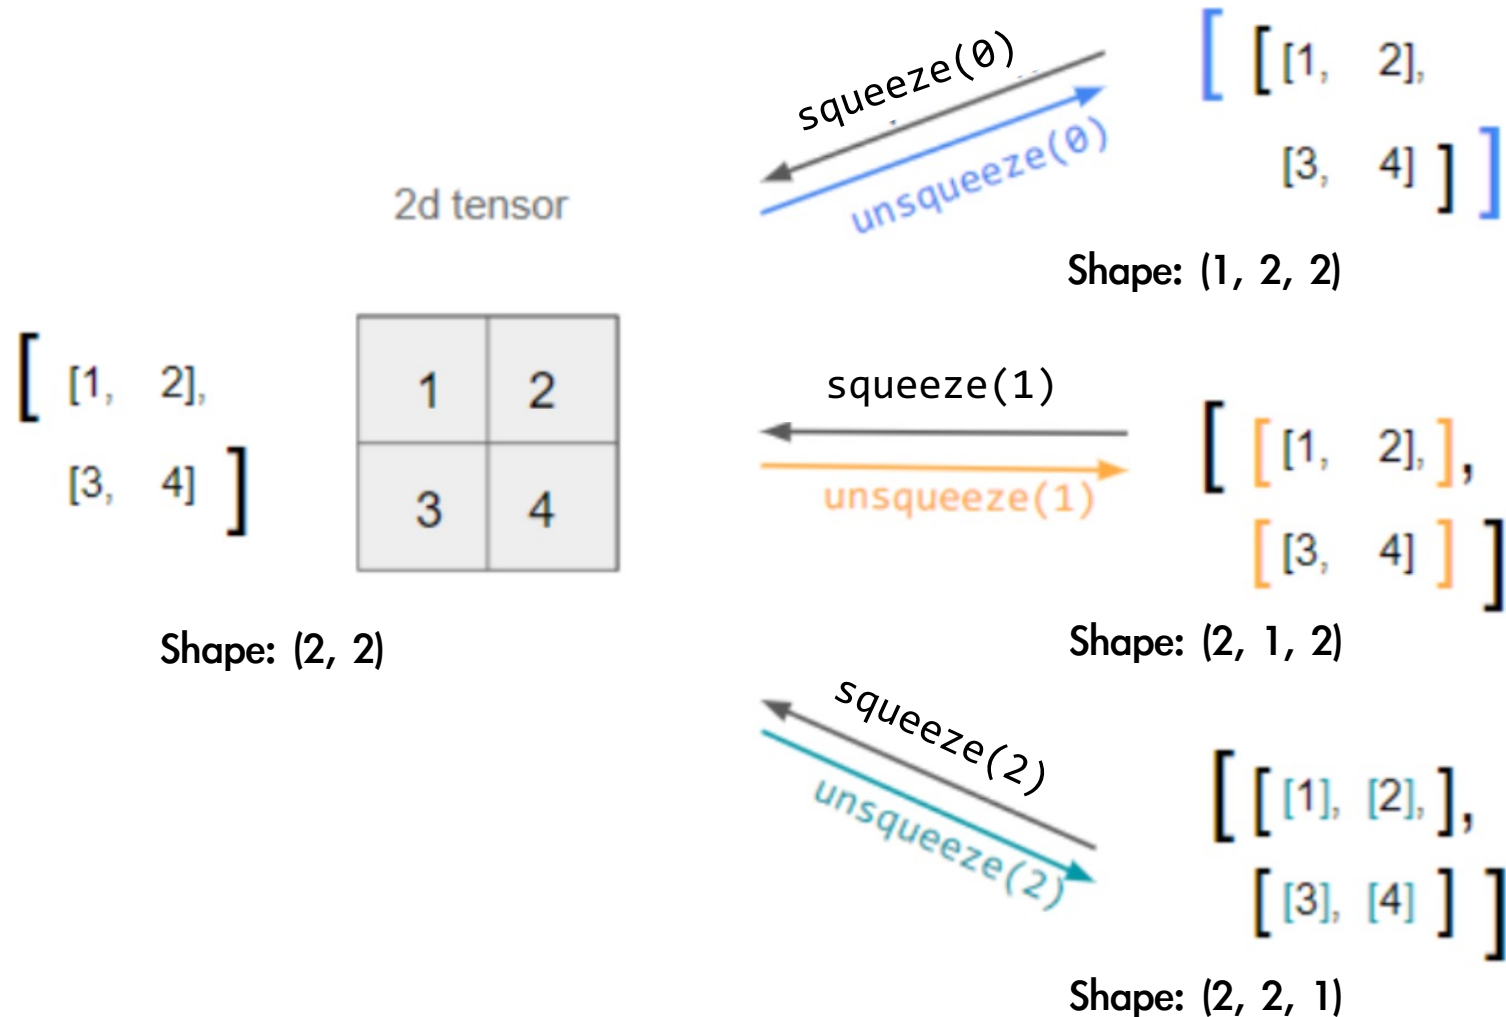
```
>>> x.shape
```

```
torch.Size([2, 3])
```



Tensor Reshaping

◆ `torch.unsqueeze()` & `torch.squeeze()`



Tensor Reshaping

◆ `torch.unsqueeze()` & `torch.squeeze()`

```
import torch

# Original tensor with shape (1, 3, 1)
t6 = torch.tensor([[[[1], [2], [3]]]])

# Remove all dimensions of size 1
t7 = t6.squeeze()    # Shape becomes (3,)

# Remove dimension at position 0
t8 = t6.squeeze(0)   # Shape becomes (3, 1)

print(t7)
# tensor([1, 2, 3])

print(t8)
# tensor([[1],
#         [2],
#         [3]])
```

```
import torch

# Original tensor with shape (3,)
t9 = torch.tensor([1, 2, 3])

# Add a new dimension at position 1
t10 = t9.unsqueeze(1) # Shape becomes (3, 1)
print(t10)
# >>> tensor([[1],
#             [2],
#             [3]])

t11 = torch.tensor(
    [[1, 2, 3],
     [4, 5, 6]]
)
t12 = t11.unsqueeze(1) # Shape becomes (2, 1, 3)
print(t12)
# >>> tensor([[[1, 2, 3]],
#             [[4, 5, 6]])
```

Tensor Reshaping

◆ `torch.flatten(input, start_dim=0, end_dim=-1)`

- Flattens input by reshaping it into a one-dimensional tensor
- If `start_dim` or `end_dim` are passed, only dimensions starting with `start_dim` and ending with `end_dim` are flattened

```
import torch

# Original tensor with shape (2, 3)
t13 = torch.tensor([[1, 2, 3], [4, 5, 6]])

# Flatten the tensor
t14 = t13.flatten() # Shape becomes (6,)

print(t14)
# >>> tensor([1, 2, 3, 4, 5, 6])
```

```
import torch

# Original tensor with shape (2, 2, 2)
t15 = torch.tensor([[[1, 2],
                     [3, 4]],
                    [[5, 6],
                     [7, 8]]])

t16 = torch.flatten(t15)
t17 = torch.flatten(t15, start_dim=1)

print(t16)
# >>> tensor([1, 2, 3, 4, 5, 6, 7, 8])

print(t17) # Shape becomes (2, 4)
# >>> tensor([[1, 2, 3, 4],
#             [5, 6, 7, 8]])
```

Tensor Reshaping

◆ `torch.permute(input, dims)` & `torch.transpose(input, dim0, dim1)`

- Returns a view of the original tensor input with its dimensions permuted
- In case of `torch.transpose`, the given dimensions `dim0` and `dim1` are swapped

```
import torch

t18 = torch.randn(2, 3, 5)
print(t18.shape) # >>> torch.Size([2, 3, 5])
print(torch.permute(t18, (2, 0, 1)).size()) # >>> torch.Size([5, 2, 3])

# Original tensor with shape (2, 3)
t19 = torch.tensor([[1, 2, 3], [4, 5, 6]])
t20 = torch.permute(t19, dims=(0, 1)) # Shape becomes (2, 3) still
t21 = torch.permute(t19, dims=(1, 0)) # Shape becomes (3, 2)
print(t20) # >>> tensor([[1, 2, 3],
                        #           [4, 5, 6]])
print(t21) # >>> tensor([[1, 4],
                        #           [2, 5],
                        #           [3, 6]])
t22 = torch.transpose(t19, 0, 1) # Shape becomes (3, 2)
print(t22) # >>> tensor([[1, 4],
                        #           [2, 5],
                        #           [3, 6]])
```

Tensor Reshaping

◆ `torch.transpose(input, dim0, dim1)` & `torch.t()`

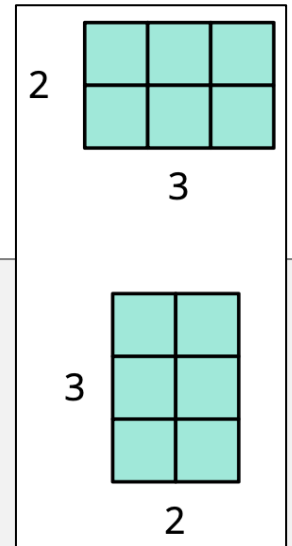
- In case of `torch.transpose()`, the given dimensions `dim0` and `dim1` are swapped
- In case of `torch.t()`, it expects input to be 2D tensor and transposes dimensions 0 and 1

```
import torch

# Original tensor with shape (2, 3)
t19 = torch.tensor([[1, 2, 3], [4, 5, 6]])

t22 = torch.transpose(t19, 0, 1) # Shape becomes (3, 2)
print(t22) # >>> tensor([[1, 4],
#                       [2, 5],
#                       [3, 6]])

t23 = torch.t(t19) # Shape becomes (3, 2)
print(t23) # >>> tensor([[1, 4],
#                       [2, 5],
#                       [3, 6]])
```



Tensor Stacking

Tensor Stacking

◆ Tensor Stacking Methods

— `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

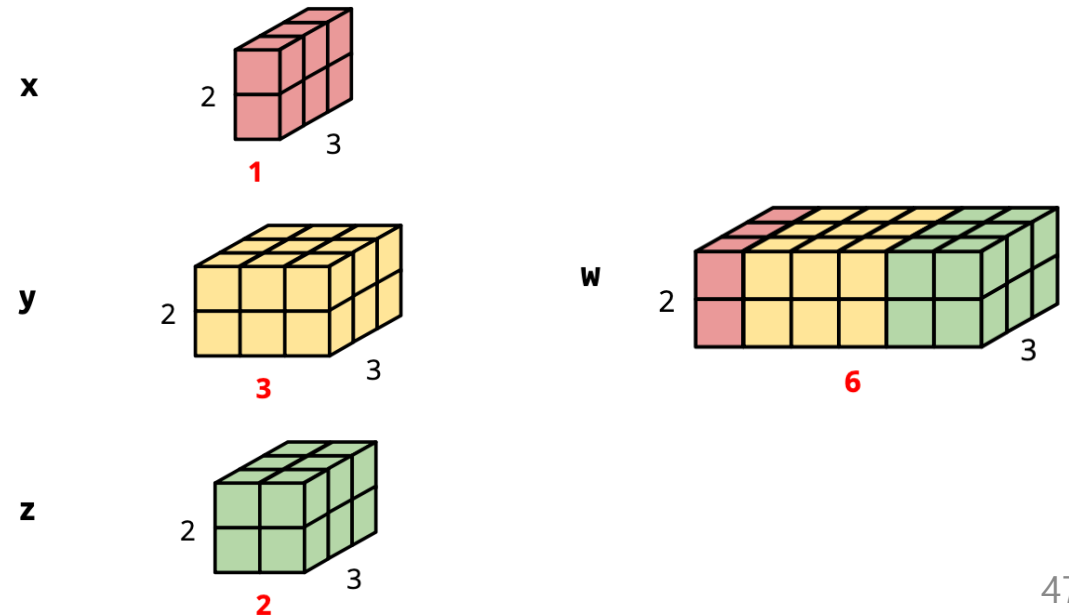
- It takes a sequence of tensors and concatenates them along the specified dimension, resulting in a new tensor
- `torch.concat(tensors, dim=0)` is alias of `torch.cat(tensors, dim=0)`

Tensor Stacking

◆ `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

- It takes a sequence of tensors and concatenates them along the specified dimension, resulting in a new tensor
- `torch.concat(tensors, dim=0)` is alias of `torch.cat(tensors, dim=0)`

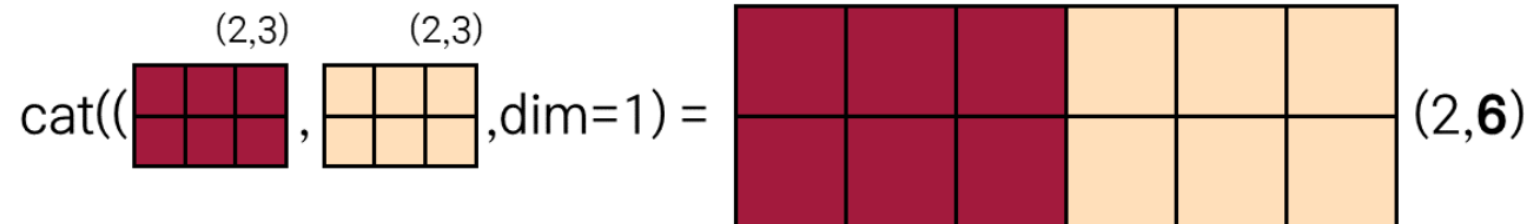
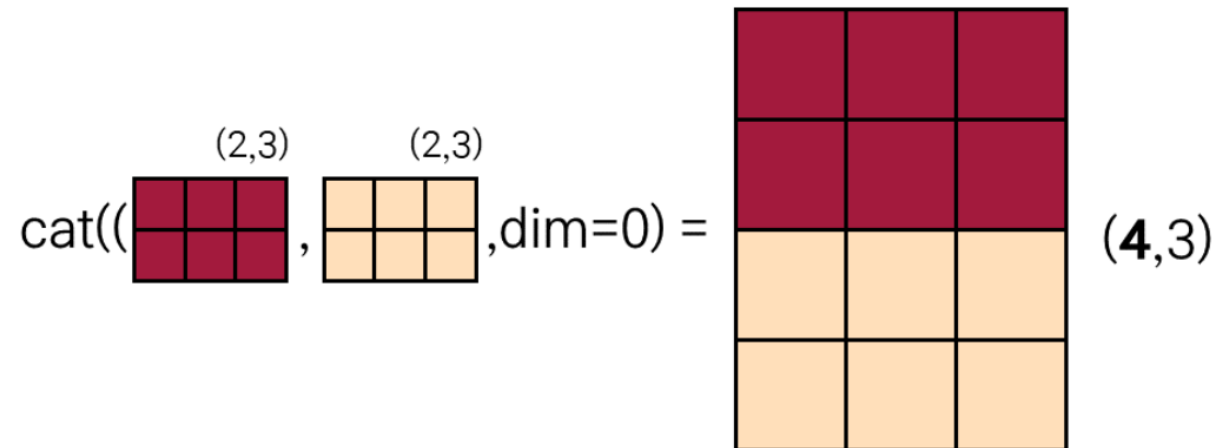
```
>>> x = torch.zeros([2, 1, 3])  
>>> y = torch.zeros([2, 3, 3])  
>>> z = torch.zeros([2, 2, 3])  
>>> w = torch.cat([x, y, z], dim=1)  
  
>>> w.shape  
torch.Size([2, 6, 3])
```



Tensor Stacking

◇ `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

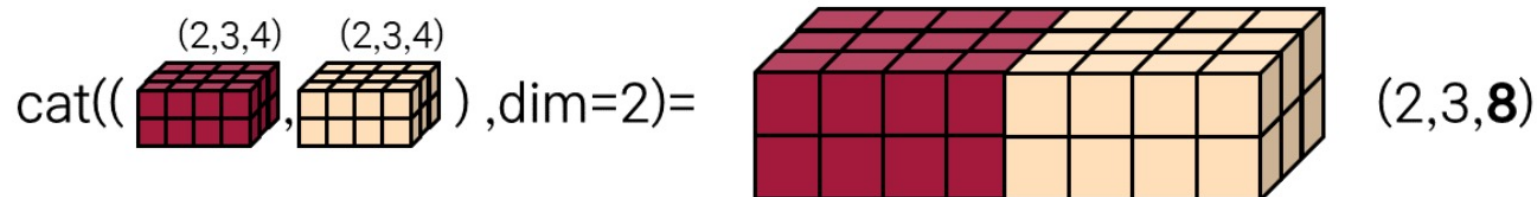
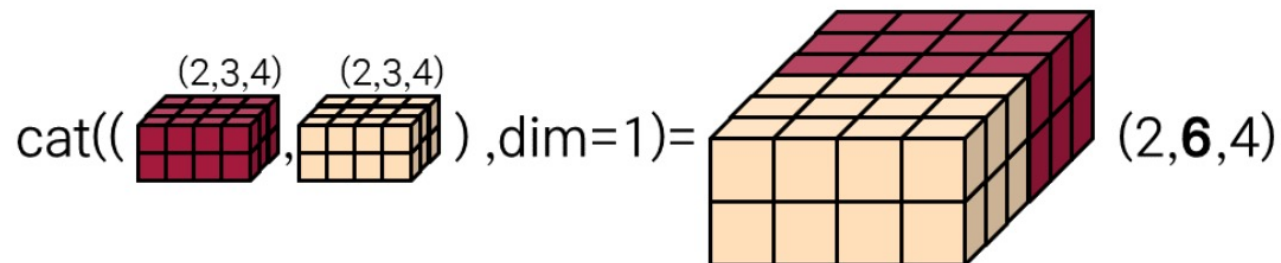
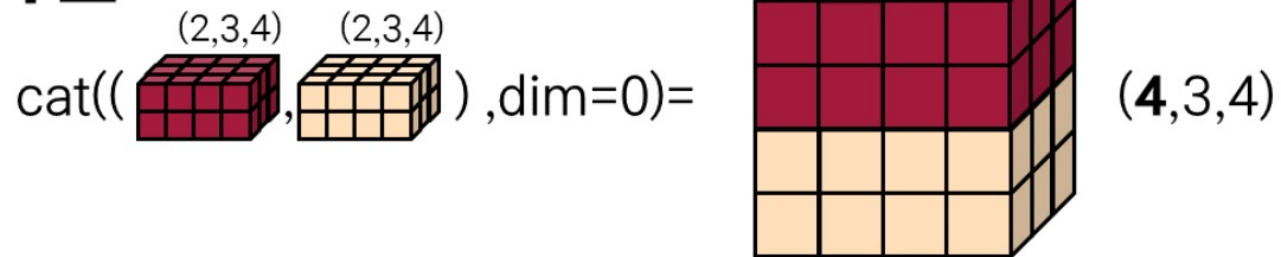
2차원



Tensor Stacking

◆ `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

3차원



Tensor Stacking

◆ `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

```
import torch

t1 = torch.zeros([2, 1, 3])
t2 = torch.zeros([2, 3, 3])
t3 = torch.zeros([2, 2, 3])

t4 = torch.cat([t1, t2, t3], dim=1)
print(t4.shape) # >>> torch.Size([2, 6, 3])

t5 = torch.arange(0, 3) # tensor([0, 1, 2])
t6 = torch.arange(3, 8) # tensor([3, 4, 5, 6, 7])

t7 = torch.cat((t5, t6), dim=0)
print(t7.shape) # >>> torch.Size([8])
print(t7) # >>> tensor([0, 1, 2, 3, 4, 5, 6, 7])
```

```
t8 = torch.arange(0, 6).reshape(2, 3)
t9 = torch.arange(6, 12).reshape(2, 3)

t10 = torch.cat((t8, t9), dim=0)
print(t10.size()) # >>> torch.Size([4, 3])
print(t10)
# >>> tensor([[ 0,  1,  2],
#              [ 3,  4,  5],
#              [ 6,  7,  8],
#              [ 9, 10, 11]])

t11 = torch.cat((t8, t9), dim=1)
print(t11.size()) # >>> torch.Size([2, 6])
print(t11)
# >>> tensor([[ 0,  1,  2,  6,  7,  8],
#              [ 3,  4,  5,  9, 10, 11]])
```

Tensor Stacking

◆ `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

```
import torch
```

```
t12 = torch.arange(0, 6).reshape(2, 3)
t13 = torch.arange(6, 12).reshape(2, 3)
t14 = torch.arange(12, 18).reshape(2, 3)
```

```
t15 = torch.cat((t12, t13, t14), dim=0)
print(t15.size()) # >>> torch.Size([6, 3])
print(t15)
```

```
# >>> tensor([[ 0,  1,  2],
#             [ 3,  4,  5],
#             [ 6,  7,  8],
#             [ 9, 10, 11],
#             [12, 13, 14],
#             [15, 16, 17]])
```

```
t16 = torch.cat((t12, t13, t14), dim=1)
print(t16.size()) # >>> torch.Size([2, 9])
print(t16)
```

```
# >>> tensor([[ 0,  1,  2,  6,  7,  8, 12, 13, 14],
#             [ 3,  4,  5,  9, 10, 11, 15, 16, 17]])
```

```
t17 = torch.arange(0, 6).reshape(1, 2, 3)
t18 = torch.arange(6, 12).reshape(1, 2, 3)
t19 = torch.cat((t17, t18), dim=0)
print(t19.size()) # >>> torch.Size([2, 2, 3])
print(t19)
```

```
# >>> tensor([[[ 0,  1,  2],
#               [ 3,  4,  5]],
#             [[ 6,  7,  8],
#               [ 9, 10, 11]]])
```

```
t20 = torch.cat((t17, t18), dim=1)
print(t20.size()) # >>> torch.Size([1, 4, 3])
print(t20)
```

```
# >>> tensor([[[ 0,  1,  2],
#               [ 3,  4,  5],
#               [ 6,  7,  8],
#               [ 9, 10, 11]]])
```

```
t21 = torch.cat((t17, t18), dim=2)
print(t21.size()) # >>> torch.Size([1, 2, 6])
print(t21)
```

```
# >>> tensor([[[ 0,  1,  2,  6,  7,  8],
#               [ 3,  4,  5,  9, 10, 11]]])
```

Tensor Stacking

◆ `torch.cat(tensors, dim=0)` or `torch.concat(tensors, dim=0)`

```
import torch

a = torch.tensor([1, 2, 3])
b = torch.tensor([4, 5, 6])

print(torch.cat([a, b], dim=0))
# >>> tensor([1, 2, 3, 4, 5, 6])

print(torch.cat([a, b], dim=1))
# >>> IndexError: Dimension out of range
# (expected to be in range of [-1, 0], but got 1)
```

Tensor Stacking

◆ Tensor Stacking Methods

— `torch.stack(tensors, dim=0)`

- It takes a sequence of tensors and stacks them along a new dimension, creating a new tensor with one additional dimension
 - 새로운 차원으로 확장하여 텐서 시퀀스를 병합
- It expects each tensor to be equal size

$$\begin{aligned}\text{cat}\left([(a, b), (a, b)], \text{dim} = 0\right) &= (2, a, b) \\ \text{cat}\left([(a, b), (a, b)], \text{dim} = 1\right) &= (a, 2, b) \\ \text{cat}\left([(a, b), (a, b)], \text{dim} = 2\right) &= (a, b, 2)\end{aligned}$$

Tensor Stacking

◆ Tensor Stacking Methods

– `torch.stack([a, b], dim=0)`


- $a: [2, 3] \rightarrow [1, 2, 3]$
- $b: [2, 3] \rightarrow [1, 2, 3]$
- ➔ $[2, 2, 3]$

– `torch.stack([a, b], dim=1)`

- $a: [2, 3] \rightarrow [2, 1, 3]$
- $b: [2, 3] \rightarrow [2, 1, 3]$
- ➔ $[2, 2, 3]$

– `torch.stack([a, b], dim=2)`

- $a: [2, 3] \rightarrow [2, 3, 1]$
- $b: [2, 3] \rightarrow [2, 3, 1]$
- ➔ $[2, 3, 2]$



```
tensor([[ 1,  2,  3],
        [ 4,  5,  6]])
tensor([[ 7,  8,  9],
        [10, 11, 12]])
```

```
tensor([[[ 1,  2,  3],
         [ 4,  5,  6]],
        [[ 7,  8,  9],
         [10, 11, 12]]])
```

```
tensor([[[ 1,  2,  3]],
        [[ 4,  5,  6]]],
       tensor([[[ 7,  8,  9]],
               [[10, 11, 12]]])
```

```
tensor([[[ 1,  2,  3],
         [ 7,  8,  9]],
        [[ 4,  5,  6],
         [10, 11, 12]]])
```

```
tensor([[[[ 1],
           [ 2],
           [ 3]],
         [[4],
          [5],
          [6]]]])
tensor([[[[ 7],
           [ 8],
           [ 9]],
         [[10],
          [11],
          [12]]]])
```

```
tensor([[[[ 1,  7],
           [ 2,  8],
           [ 3,  9]],
         [[ 4, 10],
          [ 5, 11],
          [ 6, 12]]]])
```

Tensor Stacking

◆ Tensor Stacking Methods

— `torch.stack([a, b], dim=0)`

- `a: [2, 3] → [1, 2, 3]`
 - `b: [2, 3] → [1, 2, 3]`
- `[2, 2, 3]`

— `torch.stack([a, b], dim=1)`

- `a: [2, 3] → [2, 1, 3]`
 - `b: [2, 3] → [2, 1, 3]`
- `[2, 2, 3]`

— `torch.stack([a, b], dim=2)`

- `a: [2, 3] → [2, 3, 1]`
 - `b: [2, 3] → [2, 3, 1]`
- `[2, 3, 2]`

```
import torch

t1 = torch.tensor([[1, 2, 3], [4, 5, 6]])
t2 = torch.tensor([[7, 8, 9], [10, 11, 12]])

t3 = torch.stack([t1, t2], dim=0)
t4 = torch.cat([t1.unsqueeze(dim=0), t2.unsqueeze(dim=0)], dim=0)
print(t3.shape, t3.equal(t4))
# >>> torch.Size([2, 2, 3]) True

t5 = torch.stack([t1, t2], dim=1)
t6 = torch.cat([t1.unsqueeze(dim=1), t2.unsqueeze(dim=1)], dim=1)
print(t5.shape, t5.equal(t6))
# >>> torch.Size([2, 2, 3]) True

t7 = torch.stack([t1, t2], dim=2)
t8 = torch.cat([t1.unsqueeze(dim=2), t2.unsqueeze(dim=2)], dim=2)
print(t7.shape, t7.equal(t8))
# >>> torch.Size([2, 3, 2]) True
```

Tensor Stacking

◆ `torch.stack(tensors, dim=0)`

```
t9 = torch.arange(0, 3)    # tensor([0, 1, 2])
t10 = torch.arange(3, 6)  # tensor([3, 4, 5])
```

```
print(t9.size(), t10.size())
# >>> torch.Size([3]) torch.Size([3])
```

```
t11 = torch.stack((t9, t10), dim=0)
print(t11.size()) # >>> torch.Size([2,3])
print(t11)
# >>> tensor([[0, 1, 2],
#             [3, 4, 5]])
```

```
t12 = torch.cat(
    (t9.unsqueeze(0), t10.unsqueeze(0)),
    dim=0
)
print(t11.equal(t12))
# >>> True
```

```
t13 = torch.stack((t9, t10), dim=1)
print(t13.size()) # >>> torch.Size([3,2])
print(t13)
# >>> tensor([[0, 3],
#             [1, 4],
#             [2, 5]])
```

```
t14 = torch.cat(
    (t9.unsqueeze(1), t10.unsqueeze(1)),
    dim=1
)
print(t13.equal(t14))
# >>> True
```


Tensor Stacking

◆ Tensor Stacking Methods

— `torch.vstack(tensors)`

- Stack tensors in sequence vertically (row wise)
 - the tensors should have the same number of columns

— `torch.hstack(tensors)`

- Stack tensors in sequence horizontally (column wise)
 - the tensors should have the same number of rows

Tensor Stacking

◆ torch.vstack(tensors)

```
import torch

t1 = torch.tensor([1, 2, 3])
t2 = torch.tensor([4, 5, 6])
t3 = torch.vstack((t1, t2))
print(t3)
# >>> tensor([[1, 2, 3],
#              [4, 5, 6]])

t4 = torch.tensor([[1], [2], [3]])
t5 = torch.tensor([[4], [5], [6]])
t6 = torch.vstack((t4, t5))
# >>> tensor([[1],
#              [2],
#              [3],
#              [4],
#              [5],
#              [6]])
```

```
t7 = torch.tensor([
    [[1, 2, 3], [4, 5, 6]],
    [[7, 8, 9], [10, 11, 12]]
])
print(t7.shape)
# >>> (2, 2, 3)

t8 = torch.tensor([
    [[13, 14, 15], [16, 17, 18]],
    [[19, 20, 21], [22, 23, 24]]
])
print(t8.shape)
# >>> (2, 2, 3)

t9 = torch.vstack([t7, t8])
print(t9.shape)
# >>> (4, 2, 3)

print(t9)
# >>> tensor([[[ 1,  2,  3],
#               [ 4,  5,  6]],
#              [[ 7,  8,  9],
#               [10, 11, 12]],
#              [[13, 14, 15],
#               [16, 17, 18]],
#              [[19, 20, 21],
#               [22, 23, 24]]])
```

Tensor Stacking

◆ `torch.hstack(tensors)`

```
import torch

t10 = torch.tensor([1, 2, 3])
t11 = torch.tensor([4, 5, 6])
t12 = torch.hstack((t10, t11))
print(t12)
# >>> tensor([1, 2, 3, 4, 5, 6])

t13 = torch.tensor([[1], [2], [3]])
t14 = torch.tensor([[4], [5], [6]])
t15 = torch.hstack((t13, t14))
print(t15)
# >>> tensor([[1, 4],
#             [2, 5],
#             [3, 6]])
```

```
t16 = torch.tensor([
    [[1, 2, 3], [4, 5, 6]],
    [[7, 8, 9], [10, 11, 12]]
])
print(t16.shape)
# >>> (2, 2, 3)

t17 = torch.tensor([
    [[13, 14, 15], [16, 17, 18]],
    [[19, 20, 21], [22, 23, 24]]
])
print(t17.shape)
# >>> (2, 2, 3)

t18 = torch.hstack([t16, t17])
print(t18.shape)
# >>> (2, 4, 3)

print(t18)
# >>> tensor([[[ 1,  2,  3],
#              [ 4,  5,  6],
#              [13, 14, 15],
#              [16, 17, 18]],
#             [[ 7,  8,  9],
#              [10, 11, 12],
#              [19, 20, 21],
#              [22, 23, 24]]])
```

Moving to Tensors into GPU

Moving tensors to the GPU

◆ Tensors at GPU

- Every PyTorch tensor can be transferred to (one of) the GPU(s) in order to perform massively parallel, fast computations

```
import torch

points = torch.tensor([4.0, 1.0, 5.0, 3.0, 2.0, 1.0])
points_gpu = points.to(device='cuda')
```

- We can create a tensor on the GPU

```
import torch

points_gpu = torch.tensor([[4.0, 1.0], [5.0, 3.0], [2.0, 1.0]], device='cuda')
# or
points_gpu = torch.tensor([[4.0, 1.0], [5.0, 3.0], [2.0, 1.0]], device='cuda:0')
```

Moving tensors to the GPU

◆ Tensors at GPU

- At this point, any operation performed on the tensor, such as multiplying all elements by a constant, is carried out on the GPU

```
# Multiplication performed on the GPU
points_gpu = 2 * points.to(device='cuda')
points_gpu = points_gpu + 4
```

- We can move the tensor back to the CPU

```
points_cpu = points_gpu.to(device='cpu')
```

- We can also use the shorthand methods

```
points_gpu = points.cuda() # Defaults to GPU index 0
points_gpu = points.cuda(0)
points_cpu = points_gpu.cpu()
```

Moving tensors to the GPU

◆ Tensors at GPU

- Check if your computer has NVIDIA GPU

```
if torch.cuda.is_available():  
    points_gpu = points.cuda()  
    # or  
    points_gpu = points.to('cuda')
```

- Multiple GPUs

- specify 'cuda:0', 'cuda:1', 'cuda:2', ...