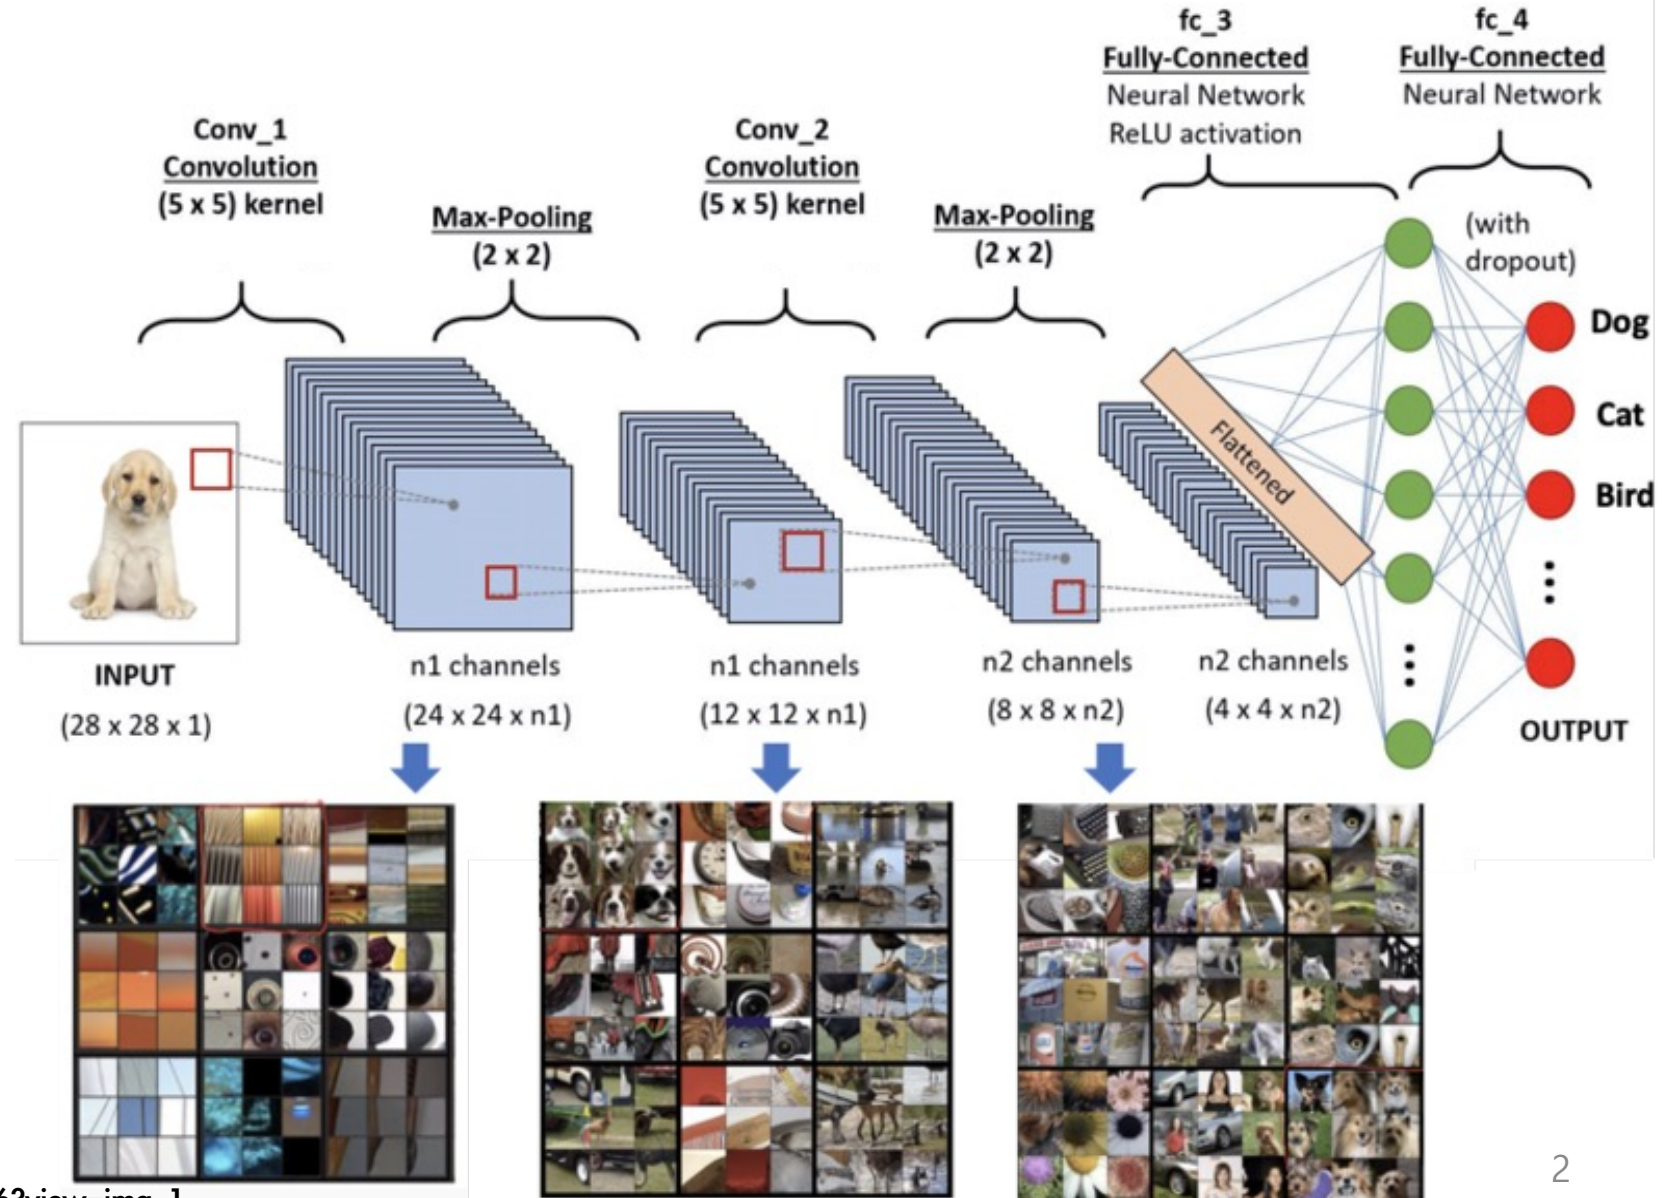


CNN Architecture & Its Examples

CNN Examples

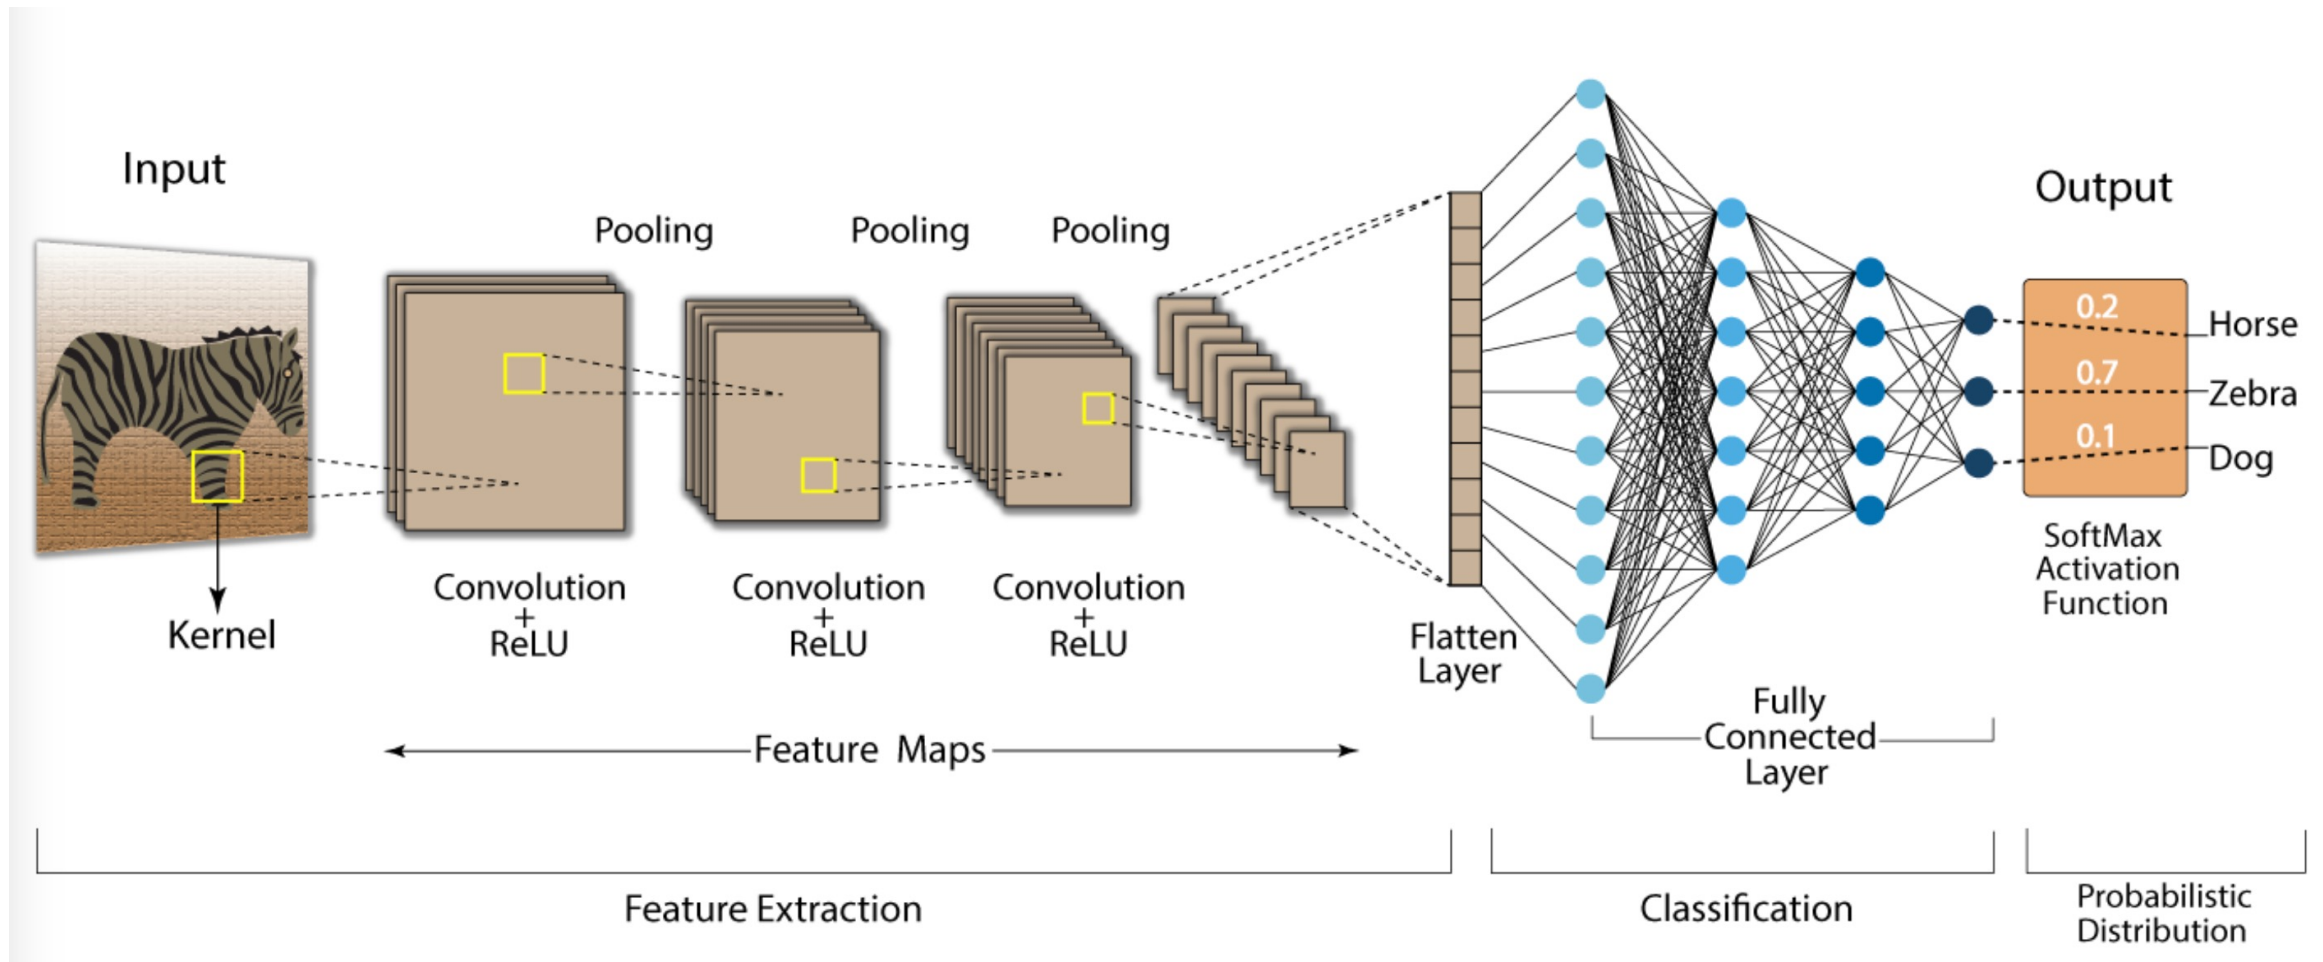
◆ A CNN Example

- Conv2D Layer
- Max-Pooling Layer
- Conv2D Layer
- Max-Pooling Layer
- Flatten Layer
- Fully-Connected (Dense or Linear) Layer
- Fully-Connected (Dense or Linear) Layer



CNN Examples

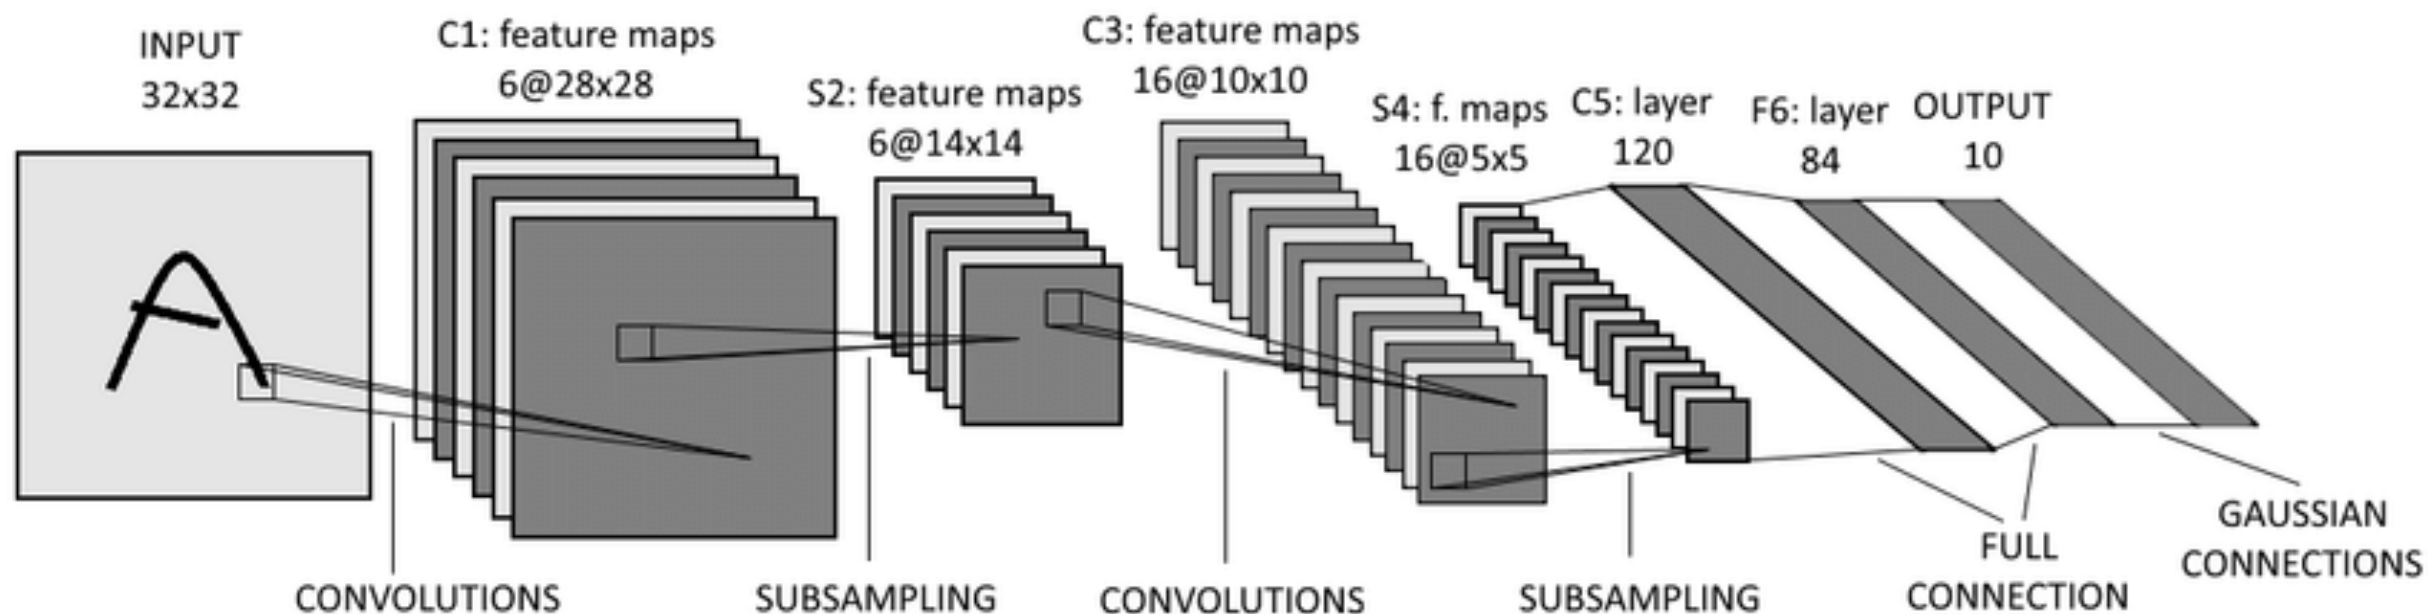
◆ A CNN Example



CNN Examples - LeNet-5

◆ LeNet-5 1998 (LeCun et al.) - 1/4

- The whole architecture is [Input - CONV - POOL - CONV - POOL - CONV - FC1 - FC2 - Output]




CNN Examples - LeNet-5

◆ LeNet-5 1998 (LeCun et al.) - 2/4

– C1 (Conv2D Layer): 6 filters of size 5x5

- Padding: 0, Stride: 1
- Input: 1x32x32 → output: 6x28x28
- Number of parameters: 1x5x5x6+6=156



$$(O_h, O_w) = \left(\frac{I_h - F_h + 2P_h}{S_h} + 1, \frac{I_w - F_w + 2P_w}{S_w} + 1 \right) = \left(\frac{32 - 5 + 2 \times 0}{1} + 1, \frac{32 - 5 + 2 \times 0}{1} + 1 \right) = (28, 28)$$

– S2 (Subsampling or Pooling Layer): 6 kernel of size 2x2

- Stride: 2
- Input: 6x28x28 → output: 6x14x14


$$(O_h, O_w) = \left(\frac{I_h - F_h}{S_h} + 1, \frac{I_w - F_w}{S_w} + 1 \right) = \left(\frac{28 - 2}{2} + 1, \frac{28 - 2}{2} + 1 \right) = (14, 14)$$

CNN Examples - LeNet-5

◆ LeNet-5 1998 (LeCun et al.) - 3/4

– C3 (Conv2D Layer): 16 filters of size 5x5

- Padding: 0, Stride: 1
- Input: 6x14x14 → output: 16x10x10
- Number of parameters: 6x5x5x16+16=2,416



$$(O_h, O_w) = \left(\frac{I_h - F_h + 2P_h}{S_h} + 1, \frac{I_w - F_w + 2P_w}{S_w} + 1 \right) = \left(\frac{14 - 5 + 2 \times 0}{1} + 1, \frac{14 - 5 + 2 \times 0}{1} + 1 \right) = (10, 10)$$

– S4 (Subsampling or Pooling Layer): 6 filters of size 2x2

- Stride: 2
- Input: 16x10x10 → output: 16x5x5

$$(O_h, O_w) = \left(\frac{I_h - F_h}{S_h} + 1, \frac{I_w - F_w}{S_w} + 1 \right) = \left(\frac{10 - 2}{2} + 1, \frac{10 - 2}{2} + 1 \right) = (5, 5)$$

CNN Examples - LeNet-5

◆ LeNet-5 1998 (LeCun et al.) - 4/4


– C5 (Conv2D Layer): 120 filters of size 5x5 ← [Flatten Layer](#)

- Padding: 0, Stride: 1
- Input: 16x5x5 → output: 120x1x1
- Number of parameters: 16x5x5x120+120 = 48,120

$$(O_h, O_w) = \left(\frac{I_h - F_h + 2P_h}{S_h} + 1, \frac{I_w - F_w + 2P_w}{S_w} + 1 \right) = \left(\frac{5 - 5 + 2 \times 0}{1} + 1, \frac{5 - 5 + 2 \times 0}{1} + 1 \right) = (1, 1)$$

– F6 (FC Layer)

- Input: 120 → output: 84
- Number of parameters: 120x84+84 = 10,164



– OUTPUT (FC Layer)

- Input: 84 → output: 10
- Number of parameters: 84x10+10 = 850

CNN Examples - AlexNet

◇ AlexNet (Alex Krizhevsky et al.)

— Paper title

- ImageNet Classification with Deep Convolutional Neural Networks, 2012

ImageNet Classification with Deep Convolutional Neural Networks

Alex Krizhevsky
University of Toronto
kriz@cs.utoronto.ca

Ilya Sutskever
University of Toronto
ilya@cs.utoronto.ca

Geoffrey E. Hinton
University of Toronto
hinton@cs.utoronto.ca



Neural Information Processing Systems

<https://proceedings.neurips.cc/paper/4824-...>

ImageNet Classification with Deep Convolutional Neural ...

A Krizhevsky 저술 120897회 인용 — We trained a large, deep convolutional neural network to **classify** the 1.2 million high-resolution images in the ImageNet LSVRC-2010 contest into the 1000 ...

페이지 9개

CNN Examples - AlexNet

mite: 진드기
leopard: 표범
Grille: 자동차 엔진 방열판

◆ AlexNet (Alex Krizhevsky et al.)

— "ImageNet" Data set

- A large-scale dataset of images that is widely used in computer vision research & machine learning
- Over 15M labeled high resolution images collected from web
- Roughly 22,000 categories

— ILSVRC


- The ImageNet Large Scale Visual Recognition Challenge (ILSVRC) is a competition that used this dataset to push the development of image classification algorithms
- 1.2M images in 1,000 categories
- Classification: make 5 guesses about the image label
- Many state-of-the-art deep neural networks, such as AlexNet, VGG, and ResNet, have been trained and evaluated on ImageNet data



CNN Examples - AlexNet

◆ AlexNet (Alex Krizhevsky et al.)


— ILSVRC History



CNN Examples - AlexNet

◆ AlexNet (Alex Krizhevsky et al.)

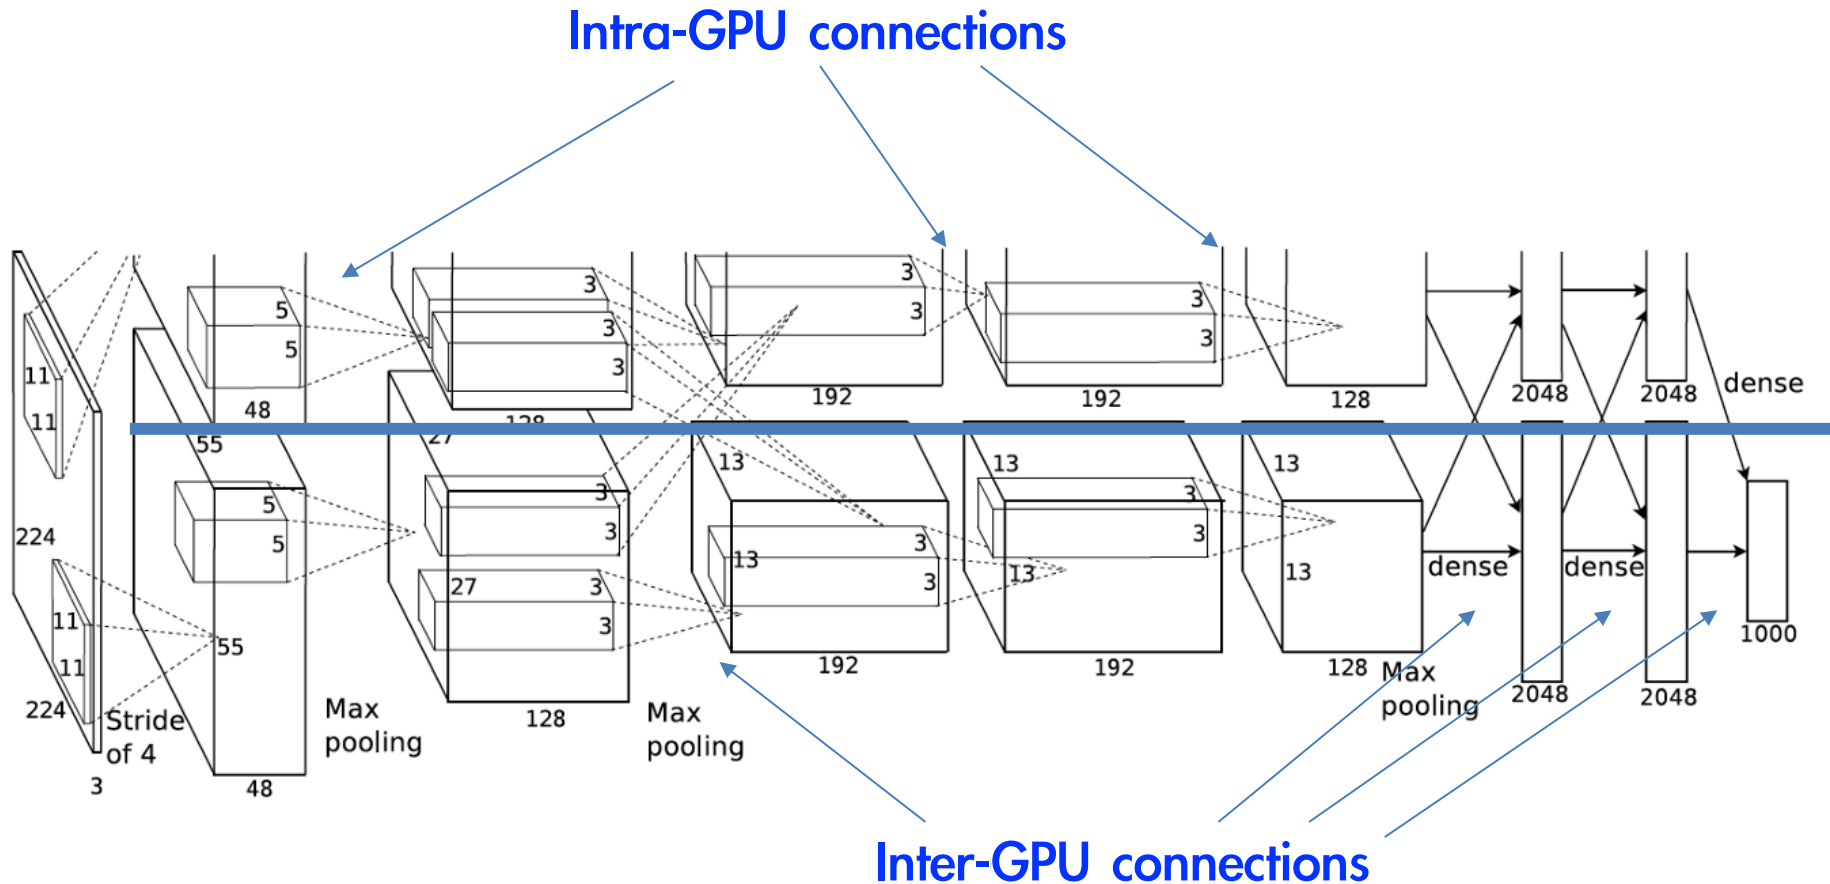
5 Convolutional Layers



3 Fully Connected Layers

CNN Examples


◆ AlexNet (Alex Krizhevsky et al.)



Top-1 and Top-5 error rates decreases by 1.7% & 1.2% respectively, comparing to the net trained with one GPU and half neurons!!

CNN Examples

◆ Input Layer & Conv-1 Layer & MaxPool-1 Layer @ AlexNet



[Input Layer]

Images: 3 x 227 x 227

AlexNet image size should be $227 \times 227 \times 3$, instead of $224 \times 224 \times 3$

[Conv-1 Layer]

96 kernels of size 11 x 11

Stride=4, Padding=0

Input: 3 x 227 x 227 → Output: 96 x 55 x 55 (= 2 x 48 x 55 x 55)

Number of parameters: 3 x 11 x 11 x 96 + 96 = 34,944

$$(O_h, O_w) = \left(\frac{I_h - F_h + 2P_h}{S_h} + 1, \frac{I_w - F_w + 2P_w}{S_w} + 1 \right) = \left(\frac{227 - 11 + 2 \times 0}{4} + 1, \frac{227 - 11 + 2 \times 0}{4} + 1 \right) = (55, 55)$$

[MaxPool Layer]


kernels of size 3 x 3, Stride=2

Input: 96 x 55 x 55 → Output: 96 x 27 x 27 (= 2 x 48 x 27 x 27)

$$(O_h, O_w) = \left(\frac{I_h - F_h}{S_h} + 1, \frac{I_w - F_w}{S_w} + 1 \right) = \left(\frac{55 - 3}{2} + 1, \frac{55 - 3}{2} + 1 \right) = (27, 27)$$

CNN Examples

◇ All Layers @ AlexNet

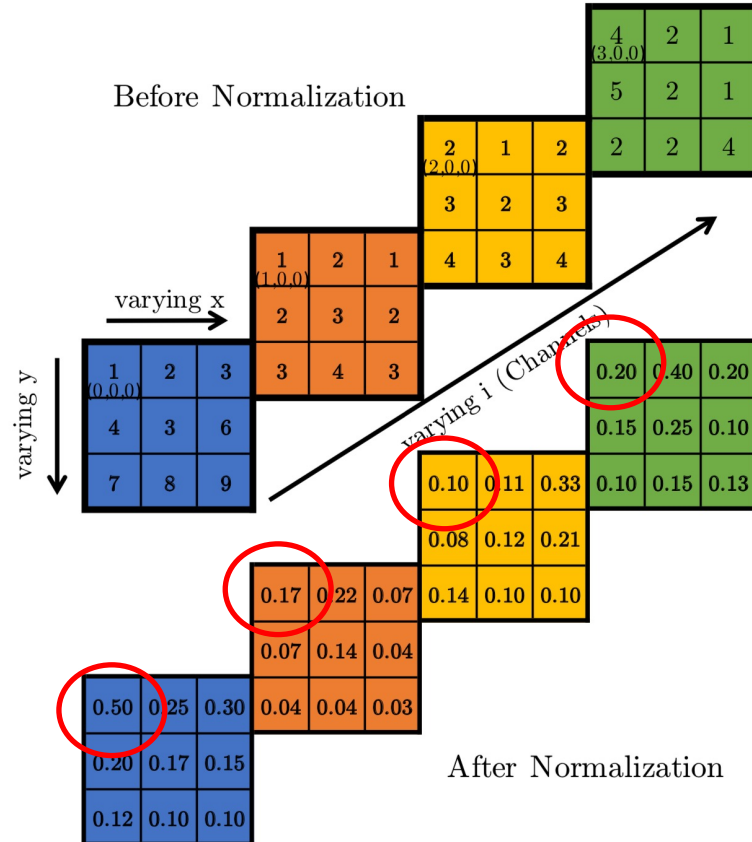


CNN Examples

◆ Local Response Normalization @ AlexNet

– Local Response Normalization

- Normalization is carried out in the channel dimension



Local Response Normalization
(size=3, alpha=1, beta=1, k=0)

$$b_{x,y}^i = a_{x,y}^i / \left(k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

i : Channel index, N : Number of channel (=4)

$$0.20 = \frac{4}{2^2 + 4^2} = \frac{4}{20}$$

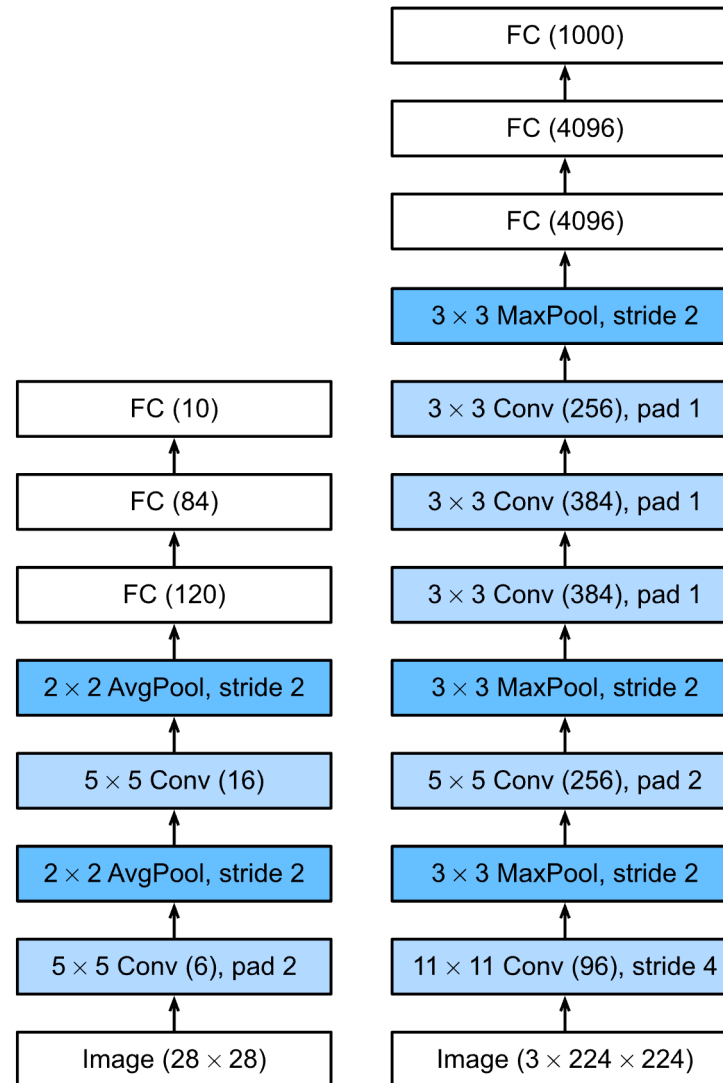
$$0.10 \approx \frac{2}{1^2 + 2^2 + 4^2} = \frac{2}{21}$$

$$0.17 \approx \frac{1}{1^2 + 1^2 + 2^2} = \frac{1}{6}$$

$$0.5 = \frac{1}{1^2 + 1^2}$$

CNN Examples

◆ LeNet (1998) vs. AlexNet (2012)



CNN Implementation

- MNIST CNN Train –
(Best Practice)**

MNIST CNN Train

◆ MNIST CNN Train

```
import torch
from torch import nn, optim
from datetime import datetime
import os
import wandb
from pathlib import Path
# BASE_PATH: /Users/yhhan/git/link_dl
BASE_PATH = str(Path(__file__).resolve().parent.parent.parent)
import sys
sys.path.append(BASE_PATH)
CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")
if not os.path.isdir(CHECKPOINT_FILE_PATH):
    os.makedirs(os.path.join(CURRENT_FILE_PATH, "checkpoints"))

import sys
sys.path.append(BASE_PATH)

from _01_code._06_fcn_best_practice.c_trainer import ClassificationTrainer
from _01_code._06_fcn_best_practice.f_mnist_train_fcn import get_data
from _01_code._06_fcn_best_practice.e_arg_parser import get_parser
```

MNIST CNN Train

◆ MNIST CNN Train

```
def get_cnn_model():  
    class MyModel(nn.Module):  
        def __init__(self, in_channels, n_output):  
            super().__init__()  
  
            self.model = nn.Sequential(  
                # B x 1 x 28 x 28 --> B x 6 x (28 - 5 + 1) x (28 - 5 + 1) = B x 6 x 24 x 24  
                nn.Conv2d(in_channels=in_channels, out_channels=6, kernel_size=(5, 5), stride=(1, 1)),  
                # B x 6 x 24 x 24 --> B x 6 x 12 x 12  
                nn.MaxPool2d(kernel_size=2, stride=2),  
                nn.ReLU(),  
                # B x 6 x 12 x 12 --> B x 16 x (12 - 5 + 1) x (12 - 5 + 1) = 16 x 8 x 8  
                nn.Conv2d(in_channels=6, out_channels=16, kernel_size=(5, 5), stride=(1, 1)),  
                # B x 16 x 8 x 8 --> B x 16 x 4 x 4  
                nn.MaxPool2d(kernel_size=2, stride=2),  
                nn.ReLU(),  
                nn.Flatten(),  
                nn.Linear(256, 128),  
                nn.ReLU(),  
                nn.Linear(128, n_output),  
            )
```

MNIST CNN Train

◆ MNIST CNN Train

```
def get_cnn_model():  
    class MyModel(nn.Module):  
        def __init__(self, in_channels, n_output):  
            ...  
  
        def forward(self, x):  
            x = self.model(x)  
            return x  
  
    # 1 * 28 * 28  
    my_model = MyModel(in_channels=1, n_output=10)  
  
    return my_model
```

MNIST CNN Train

◆ MNIST CNN Train

```
def main(args):
    run_time_str = datetime.now().astimezone().strftime('%Y-%m-%d_%H-%M-%S')

    config = {
        'epochs': args.epochs,
        'batch_size': args.batch_size,
        'validation_intervals': args.validation_intervals,
        'learning_rate': args.learning_rate,
        'early_stop_patience': args.early_stop_patience
    }
    project_name = "cnn_mnist"
    wandb.init(
        mode="online" if args.wandb else "disabled",
        project=project_name,
        notes="mnist experiment with cnn",
        tags=["cnn", "mnist"],
        name=run_time_str,
        config=config
    )
    print(args)
    print(wandb.config)
```

MNIST CNN Train

◆ MNIST CNN Train

```
def main(args):
    ...
    device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
    print(f"Training on device {device}.")

    train_data_loader, validation_data_loader, mnist_transforms = get_data(flatten=False)
    model = get_cnn_model()
    model.to(device)
    wandb.watch(model)

    from torchinfo import summary
    summary(model=model, input_size=(1, 1, 28, 28))
```

```
=====
Layer (type:depth-idx)                   Output Shape          Param #
=====
MyModel                                  [1, 10]               --
├─Sequential: 1-1                        [1, 10]               --
│   └─Conv2d: 2-1                        [1, 6, 24, 24]        156
│       └─MaxPool2d: 2-2                 [1, 6, 12, 12]        --
│           └─ReLU: 2-3                  [1, 6, 12, 12]        --
│               └─Conv2d: 2-4             [1, 16, 8, 8]         2,416
│                   └─MaxPool2d: 2-5      [1, 16, 4, 4]         --
│                       └─ReLU: 2-6      [1, 16, 4, 4]         --
│                           └─Flatten: 2-7 [1, 256]              --
│                               └─Linear: 2-8 [1, 128]              32,896
│                                   └─ReLU: 2-9 [1, 128]              --
│                                       └─Linear: 2-10 [1, 10]              1,290
=====
Total params: 36,758
Trainable params: 36,758
Non-trainable params: 0
Total mult-adds (M): 0.28
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.04
Params size (MB): 0.15
Estimated Total Size (MB): 0.19
=====
```

MNIST CNN Train

◆ MNIST CNN Train

```
def main(args):
    ...
    optimizer = optim.SGD(model.parameters(), lr=wandb.config.learning_rate)
    classification_trainer = ClassificationTrainer(
        project_name, model, optimizer, train_data_loader, validation_data_loader, mnist_transforms,
        run_time_str, wandb, device, CHECKPOINT_FILE_PATH
    )
    classification_trainer.train_loop()

    wandb.finish()

if __name__ == "__main__":
    parser = get_parser()
    args = parser.parse_args()
    main(args)

# python _01_code/_07_cnn/a_mnist_train_cnn.py --wandb -b 2048 -r 1e-3 -v 10
# python _01_code/_07_cnn/a_mnist_train_cnn.py --no-wandb -b 2048 -r 1e-3 -v 10
```

CNN Implementation

- MNIST CNN Test –
(Best Practice)**

MNIST CNN Train

◆ MNIST CNN Test

```
import numpy as np
import torch
import os

from matplotlib import pyplot as plt
from pathlib import Path

BASE_PATH = str(Path(__file__).resolve().parent.parent.parent) # BASE_PATH: /Users/yhhan/git/link_dl
CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")

import sys
sys.path.append(BASE_PATH)

from _01_code._07_cnn.a_mnist_train_cnn import get_cnn_model
from _01_code._06_fcn_best_practice.d_tester import ClassificationTester
from _01_code._06_fcn_best_practice.g_mnist_test_fcn import get_test_data
```

MNIST CNN Train

◆ MNIST CNN Test

```
def main():
    mnist_test_images, test_data_loader, mnist_transforms = get_test_data(flatten=False)

    test_model = get_cnn_model()
    classification_tester = ClassificationTester(
        "mnist", test_model, test_data_loader, mnist_transforms, CHECKPOINT_FILE_PATH
    )
    classification_tester.test()

    img, label = mnist_test_images[0]
    print("    LABEL:", label)
    plt.imshow(img)
    plt.show()

    output = classification_tester.test_single(
        torch.tensor(np.array(mnist_test_images[0][0])).unsqueeze(dim=0).unsqueeze(dim=0)
    ) # (1, 1, 28, 28)
    print("PREDICTION:", output)

if __name__ == "__main__":
    main()
```

CNN Implementation

- CIFAR10 CNN Train –
(Best Practice)**

CIFAR10 CNN Train

◆ CIFAR10 CNN Train

```
import torch
from torch import nn, optim
from datetime import datetime
import os
import wandb
from pathlib import Path
# # BASE_PATH: /Users/yhhan/git/link_dl
BASE_PATH = str(Path(__file__).resolve().parent.parent.parent)
import sys
sys.path.append(BASE_PATH)
CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")
if not os.path.isdir(CHECKPOINT_FILE_PATH):
    os.makedirs(os.path.join(CURRENT_FILE_PATH, "checkpoints"))

import sys
sys.path.append(BASE_PATH)

from _01_code._06_fcn_best_practice.c_trainer import ClassificationTrainer
from _01_code._06_fcn_best_practice.h_cifar10_train_fcn import get_data
from _01_code._06_fcn_best_practice.e_arg_parser import get_parser
```

CIFAR10 CNN Train

◆ CIFAR10 CNN Train

```
def get_cnn_model():  
    class MyModel(nn.Module):  
        def __init__(self, in_channels, n_output):  
            super().__init__()  
  
            self.model = nn.Sequential(  
                # B x 3 x 32 x 32 --> B x 6 x (32 - 5 + 1) x (32 - 5 + 1) = B x 6 x 28 x 28  
                nn.Conv2d(in_channels=in_channels, out_channels=6, kernel_size=(5, 5), stride=(1, 1)),  
                # B x 6 x 28 x 28 --> B x 6 x 14 x 14  
                nn.MaxPool2d(kernel_size=2, stride=2),  
                nn.ReLU(),  
                # B x 6 x 14 x 14 --> B x 16 x (14 - 5 + 1) x (14 - 5 + 1) = B x 16 x 10 x 10  
                nn.Conv2d(in_channels=6, out_channels=16, kernel_size=(5, 5), stride=(1, 1)),  
                # B x 16 x 10 x 10 --> B x 16 x 5 x 5  
                nn.MaxPool2d(kernel_size=2, stride=2),  
                nn.ReLU(),  
                nn.Flatten(),  
                nn.Linear(400, 128),  
                nn.ReLU(),  
                nn.Linear(128, n_output),  
            )
```

CIFAR10 CNN Train

◆ CIFAR10 CNN Train

```
def get_cnn_model():  
    class MyModel(nn.Module):  
        def __init__(self, in_channels, n_output):  
            ...  
  
        def forward(self, x):  
            x = self.model(x)  
            # print(x.shape, "!!!")  
            return x  
  
    # 3 * 32 * 32  
    my_model = MyModel(in_channels=3, n_output=10)  
  
    return my_model
```

CIFAR10 CNN Train

◆ CIFAR10 CNN Train

```
def main(args):
    run_time_str = datetime.now().astimezone().strftime('%Y-%m-%d_%H-%M-%S')

    config = {
        'epochs': args.epochs,
        'batch_size': args.batch_size,
        'validation_intervals': args.validation_intervals,
        'learning_rate': args.learning_rate,
        'early_stop_patience': args.early_stop_patience
    }
    project_name = "cnn_cifar10"
    wandb.init(
        mode="online" if args.wandb else "disabled",
        project=project_name,
        notes="cifar10 experiment with cnn",
        tags=["cnn", "cifar10"],
        name=run_time_str,
        config=config
    )
    print(args)
    print(wandb.config)
```

CIFAR10 CNN Train

◆ CIFAR10 CNN Train

```
def main(args):
```

```
...
```

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
print(f"Training on device {device}.")
```

```
train_data_loader, validation_data_loader, cifar10_transforms = get_data(flatten=False)
```

```
model = get_cnn_model()
```

```
model.to(device)
```

```
wandb.watch(model)
```

```
from torchinfo import summary
```

```
summary(model=model, input_size=(1, 3, 32, 32))
```

```
=====
```

Layer (type:depth-idx)	Output Shape	Param #
MyModel	[1, 10]	--
├─Sequential: 1-1	[1, 10]	--
│ └─Conv2d: 2-1	[1, 6, 28, 28]	456
│ └─MaxPool2d: 2-2	[1, 6, 14, 14]	--
│ └─ReLU: 2-3	[1, 6, 14, 14]	--
│ └─Conv2d: 2-4	[1, 16, 10, 10]	2,416
│ └─MaxPool2d: 2-5	[1, 16, 5, 5]	--
│ └─ReLU: 2-6	[1, 16, 5, 5]	--
│ └─Flatten: 2-7	[1, 400]	--
│ └─Linear: 2-8	[1, 128]	51,328
│ └─ReLU: 2-9	[1, 128]	--
│ └─Linear: 2-10	[1, 10]	1,290

```
=====
```

Total params: 55,490
Trainable params: 55,490
Non-trainable params: 0
Total mult-adds (M): 0.65

```
=====
```

Input size (MB): 0.01
Forward/backward pass size (MB): 0.05
Params size (MB): 0.22
Estimated Total Size (MB): 0.29

```
=====
```


CIFAR10 CNN Train

◆ CIFAR10 CNN Train

```
def main(args):
    ...
    optimizer = optim.SGD(model.parameters(), lr=wandb.config.learning_rate)

    classification_trainer = ClassificationTrainer(
        project_name, model, optimizer, train_data_loader, validation_data_loader, cifar10_transforms,
        run_time_str, wandb, device, CHECKPOINT_FILE_PATH
    )
    classification_trainer.train_loop()

    wandb.finish()

if __name__ == "__main__":
    parser = get_parser()
    args = parser.parse_args()
    main(args)
# python _01_code/_07_cnn/c_cifar10_train_cnn.py --wandb -b 2048 -r 1e-3 -v 10
# python _01_code/_07_cnn/c_cifar10_train_cnn.py --no-wandb -b 2048 -r 1e-3 -v 10
```

CNN Implementation

- CIFAR10 CNN Test –
(Best Practice)**

CIFAR10 CNN Test

◆ CIFAR10 CNN Test

```
import numpy as np
import torch
import os

from matplotlib import pyplot as plt
from pathlib import Path

BASE_PATH = str(Path(__file__).resolve().parent.parent.parent) # BASE_PATH: /Users/yhhan/git/link_d
1
CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")

import sys
sys.path.append(BASE_PATH)

from _01_code._07_cnn.c_cifar10_train_cnn import get_cnn_model
from _01_code._06_fcn_best_practice.d_tester import ClassificationTester
from _01_code._06_fcn_best_practice.i_cifar10_test_fcn import get_test_data
```

CIFAR10 CNN Test

◆ CIFAR10 CNN Test

```
def main():
    cifar10_test_images, test_data_loader, cifar10_transforms = get_test_data(flatten=False)

    test_model = get_cnn_model()
    classification_tester = ClassificationTester(
        "cifar10", test_model, test_data_loader, cifar10_transforms, CHECKPOINT_FILE_PATH
    )
    classification_tester.test()

    img, label = cifar10_test_images[0]
    print("    LABEL:", label)
    plt.imshow(img)
    plt.show()

    output = classification_tester.test_single(
        torch.tensor(np.array(cifar10_test_images[0][0])).permute(2, 0, 1).unsqueeze(dim=0)
    ) # (1, 3, 32, 32)
    print("PREDICTION:", output)

if __name__ == "__main__":
    main()
```

CNN Implementation

- CIFAR10 AlexNet Train –
(Best Practice)**

CIFAR10 AlexNet Train

◆ CIFAR10 AlexNet Train

```
def get_alexnet_model():
    class AlexNet(nn.Module):
        def __init__(self, in_channels=3, n_output=10):
            super().__init__()
            # The image in the original paper states that width and height are 224 pixels, but
            # the correct input size should be : (B x 3 x 227 x 227)
            self.cnn = nn.Sequential(
                # B x 3 x 32 x 32 -> B x 64 x (32 - 3 + 1) x (32 - 3 + 1) = B x 64 x 30 x 30
                nn.Conv2d(in_channels=in_channels, out_channels=64, kernel_size=(3, 3), stride=(1, 1)),
                nn.ReLU(),
                nn.LocalResponseNorm(size=3, alpha=0.0001, beta=0.75, k=2),
                # B x 64 x 30 x 30 -> B x 64 x ((30 - 2) / 2 + 1) x ((30 - 2) / 2 + 1) = B x 64 x 15 x 15
                nn.MaxPool2d(kernel_size=2, stride=2),

                # B x 64 x 15 x 15 -> B x 64 x (15 - 3 + 2 + 1) x (15 - 3 + 2 + 1) = B x 192 x 15 x 15
                nn.Conv2d(64, 192, (3, 3), (1, 1), padding=1),
                nn.ReLU(),
                nn.LocalResponseNorm(size=3, alpha=0.0001, beta=0.75, k=2),
                # B x 192 x 15 x 15 -> B x 192 x ((15 - 3) / 2 + 1) x ((15 - 3) / 2 + 1) = B x 192 x 7 x 7
                nn.MaxPool2d(kernel_size=3, stride=2),
```

CIFAR10 AlexNet Train

◆ CIFAR10 AlexNet Train

```
def get_alexnet_model():
    class AlexNet(nn.Module):
        def __init__(self, in_channels=3, n_output=10):
            super().__init__()
            ...
            self.cnn = nn.Sequential(
                ...
                # B x 192 x 7 x 7 -> B x 256 x ((7 - 3 + 2) / 1 + 1) x ((13 - 3 + 2) / 1 + 1) = B x 256 x 7 x 7
                nn.Conv2d(192, 256, (3, 3), (1, 1), padding=1),
                nn.ReLU(),

                # B x 256 x 7 x 7 -> B x 256 x ((7 - 3 + 2) / 1 + 1) x ((13 - 3 + 2) / 1 + 1) = B x 256 x 7 x 7
                nn.Conv2d(256, 256, (3, 3), (1, 1), padding=1),
                nn.ReLU(),

                # B x 256 x 7 x 7 -> B x 192 x ((7 - 2) / 1 + 1) x ((7 - 2) / 1 + 1) = B x 192 x 6 x 6
                nn.Conv2d(256, 192, (2, 2), (1, 1)),
                nn.ReLU(),

                # B x 192 x 6 x 6 -> B x 192 x ((6 - 2) / 2 + 1) x ((6 - 2) / 2 + 1) = B x 192 x 3 x 3
                nn.MaxPool2d(kernel_size=2, stride=2),
            )
```

CIFAR10 AlexNet Train

◆ CIFAR10 AlexNet Train

```
def get_alexnet_model():  
    class AlexNet(nn.Module):  
        def __init__(self, in_channels=3, n_output=10):  
            super().__init__()  
            ...  
  
            # classifier is just a name for linear layers  
            self.fcn = nn.Sequential(  
                nn.Dropout(p=0.5),  
                nn.Linear(in_features=192 * 3 * 3, out_features=512),  
                nn.ReLU(),  
                nn.Dropout(p=0.5),  
                nn.Linear(in_features=512, out_features=512),  
                nn.ReLU(),  
                nn.Linear(in_features=512, out_features=n_output),  
            )
```


CIFAR10 AlexNet Train

◆ CIFAR10 AlexNet Train

```
def get_alexnet_model():
    class AlexNet(nn.Module):

        ...

    def forward(self, x):
        """
        Pass the input through the net.
        """
        x = self.cnn(x)
        x = x.view(-1, 192 * 3 * 3) # reduce the dimensions for linear layer input
        return self.fcn(x)

my_model = AlexNet(in_channels=3, n_output=10)

return my_model
```

CIFAR10 AlexNet Train

◆ CIFAR10 AlexNet Train


```
from torchinfo import summary
```

```
model = get_alexnet_model()
```

```
summary(
    model=model,
    input_size=(1, 3, 32, 32),
    col_names=[
        "kernel_size",
        "input_size",
        "output_size",
        "num_params",
        "mult_adds"
    ]
)
```

Layer (type:depth-idx)	Kernel Shape	Input Shape	Output Shape	Param #	Mult-Adds
=====					
AlexNet	--	[1, 3, 32, 32]	[1, 10]	--	--
└Sequential: 1-1	--	[1, 3, 32, 32]	[1, 192, 3, 3]	--	--
└Conv2d: 2-1	[3, 3]	[1, 3, 32, 32]	[1, 64, 30, 30]	1,792	1,612,800
└ReLU: 2-2	--	[1, 64, 30, 30]	[1, 64, 30, 30]	--	--
└LocalResponseNorm: 2-3	--	[1, 64, 30, 30]	[1, 64, 30, 30]	--	--
└MaxPool2d: 2-4	2	[1, 64, 30, 30]	[1, 64, 15, 15]	--	--
└Conv2d: 2-5	[3, 3]	[1, 64, 15, 15]	[1, 192, 15, 15]	110,784	24,926,400
└ReLU: 2-6	--	[1, 192, 15, 15]	[1, 192, 15, 15]	--	--
└LocalResponseNorm: 2-7	--	[1, 192, 15, 15]	[1, 192, 15, 15]	--	--
└MaxPool2d: 2-8	3	[1, 192, 15, 15]	[1, 192, 7, 7]	--	--
└Conv2d: 2-9	[3, 3]	[1, 192, 7, 7]	[1, 256, 7, 7]	442,624	21,688,576
└ReLU: 2-10	--	[1, 256, 7, 7]	[1, 256, 7, 7]	--	--
└Conv2d: 2-11	[3, 3]	[1, 256, 7, 7]	[1, 256, 7, 7]	590,080	28,913,920
└ReLU: 2-12	--	[1, 256, 7, 7]	[1, 256, 7, 7]	--	--
└Conv2d: 2-13	[2, 2]	[1, 256, 7, 7]	[1, 192, 6, 6]	196,800	7,084,800
└ReLU: 2-14	--	[1, 192, 6, 6]	[1, 192, 6, 6]	--	--
└MaxPool2d: 2-15	2	[1, 192, 6, 6]	[1, 192, 3, 3]	--	--
└Sequential: 1-2	--	[1, 1728]	[1, 10]	--	--
└Dropout: 2-16	--	[1, 1728]	[1, 1728]	--	--
└Linear: 2-17	--	[1, 1728]	[1, 512]	885,248	885,248
└ReLU: 2-18	--	[1, 512]	[1, 512]	--	--
└Dropout: 2-19	--	[1, 512]	[1, 512]	--	--
└Linear: 2-20	--	[1, 512]	[1, 512]	262,656	262,656
└ReLU: 2-21	--	[1, 512]	[1, 512]	--	--
└Linear: 2-22	--	[1, 512]	[1, 10]	5,130	5,130
=====					
Total params: 2,495,114					
Trainable params: 2,495,114					
Non-trainable params: 0					
Total mult-adds (M): 85.38					
=====					
Input size (MB): 0.01					
Forward/backward pass size (MB): 1.07					
Params size (MB): 9.98					
Estimated Total Size (MB): 11.06					
=====					

- https://wandb.ai/link-koreatech/alexnet_cifar10/workspace?workspace=user-link-koreatech



CNN Implementation

- CIFAR10 AlexNet Test –
(Best Practice)**

CIFAR10 AlexNet Test

◆ CIFAR10 AlexNet Test

```
import numpy as np
import torch
import os

from matplotlib import pyplot as plt
from pathlib import Path

BASE_PATH = str(Path(__file__).resolve().parent.parent.parent) # BASE_PATH: /Users/yhhan/git/link_dl
CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")

import sys
sys.path.append(BASE_PATH)

from _01_code._06_fcn_best_practice.d_tester import ClassificationTester
from _01_code._06_fcn_best_practice.i_cifar10_test_fcn import get_cifar10_test_data
from _01_code._07_cnn.e_cifar10_train_alexnet import get_alexnet_model
```

CIFAR10 AlexNet Test

◆ CIFAR10 AlexNet Test

```
def main():
    cifar10_test_images, test_data_loader, cifar10_transforms = get_cifar10_test_data(flatten=False)

    test_model = get_alexnet_model()

    project_name = "alexnet_cifar10"
    classification_tester = ClassificationTester(
        project_name, test_model, test_data_loader, cifar10_transforms, CHECKPOINT_FILE_PATH
    )
    classification_tester.test()

    print()

    img, label = cifar10_test_images[0]
    print("    LABEL:", label)
    plt.imshow(img)
    plt.show()
```

CIFAR10 AlexNet Test

◆ CIFAR10 AlexNet Test

```
def main():
    ...

    # torch.tensor(np.array(cifar10_test_images[0][0])).permute(2, 0, 1).unsqueeze(dim=0).shape: (1, 3, 32, 32)
    output = classification_tester.test_single(
        torch.tensor(np.array(cifar10_test_images[0][0])).permute(2, 0, 1).unsqueeze(dim=0)
    )
    print("PREDICTION:", output)

if __name__ == "__main__":
    main()
```