

Recurrent Neural Networks (RNNs)

November 2023



<http://link.koreatech.ac.kr>

Time Series Data

시계열 데이터

◆시계열 데이터 (Time Series Data)

- d_t : 순서가 있는 연속 데이터 (Sequenced Data)
- $t = 1, 2, 3, \dots$
 - t 가 실제로 일정 간격의 시각에 대응되는 경우가 많음
 - 하지만, 어떤 경우에는 단순한 이벤트 기반 순서만 나타낼 수 있음
- 예: 음성, 동영상, 텍스트 등
- 연속적 데이터의 길이 T 는 고정 또는 가변
 - 종료 시점이 없을 수도 있음 \rightarrow 무한 데이터

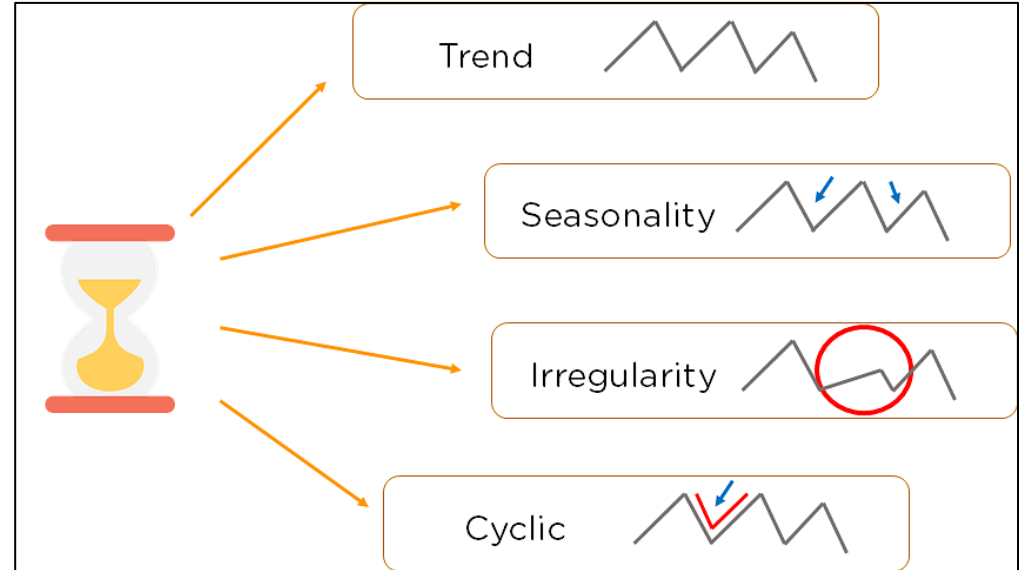


Date	Ozone ($\mu\text{g}/\text{m}^3$)	Temperature ($^{\circ}\text{C}$)	Relative humidity (%)	n deaths
1 Jan 2002	4.59	-0.2	75.7	199
2 Jan 2002	4.88	0.1	77.5	231
3 Jan 2002	4.71	0.9	81.3	210
4 Jan 2002	4.14	0.5	85.4	203
5 Jan 2002	2.01	4.3	93.5	224
6 Jan 2002	2.4	7.1	96.4	198
7 Jan 2002	4.08	5.2	93.5	180
8 Jan 2002	3.13	3.5	81.5	188
9 Jan 2002	2.05	3.2	88.3	168
10 Jan 2002	5.19	5.3	85.4	194
11 Jan 2002	3.59	3.0	92.6	223
12 Jan 2002	12.87	4.8	94.2	201

시계열 데이터 활용


◆시계열 해석(Time Series Analysis)

- 시계열을 해석하고 이해하는 데 쓰이는 여러 가지 방법을 연구
 - 이런 시계열이 어떤 법칙에서 생성되어서 나오는가에 대한 기본적인 질문을 이해



◆시계열 예측(time series prediction)


- 주어진 시계열을 보고 모델을 만들어서 미래에 일어날 것들을 예측하는 것



시계열 예측 예

◆ 단어 예측 문제

- 입력 연속렬 데이터 $x^1, x^2, x^3, x^4, \dots, x^t$ 로 부터 y^t 를 예측하는 문제
- 전체적인 문맥을 학습하여 다음에 올 단어를 높은 정확도로 예측



to be or not to be that is the question


word	to	be	or	...	is	the	?
input	x^1	x^2	x^3		x^{t-1}	x^t	
output		y^1	y^2		y^{t-2}	y^{t-1}	y^t

RNN Cell & RNN Layer

RNN Cell

◆ RNN Cell - Definition

- RNN cell includes the ability to maintain internal memory with feedback and therefore support temporal behavior
- RNN cell remains feed-forward still, but also maintains (internal) hidden state
 $H_{[t]}^{(l)} = (h_1, h_2, \dots, h_I)_{[t]}^{(l)}$ across each time step t
- The term "cell" captures the recurrent and sequential nature of the operation




RNN Cell

◆ RNN Cell

```
rnn_cell = nn.RNNCell(
    input_size=3, hidden_size=4,
    bias=True, nonlinearity='tanh', device=None
)

for name, parameter in rnn_cell.named_parameters():
    print(name, parameter.shape)

# >>> weight_ih torch.Size([4, 3]):  $W^{(l)}$ 
# >>> weight_hh torch.Size([4, 4]):  $W_H^{(l)}$ 
# >>> bias_ih torch.Size([4]):  $b^{(l)}$ 
# >>> bias_hh torch.Size([4]):  $b_H^{(l)}$ 
```



Why the "tanh" as the default in RNNs over "ReLU"?

- Stability in Recurrent Connections

- : The bounded nature of "tanh" can lead to more stable behavior in the recurrent connections of an RNN.
- : ReLU's unbounded output can cause the activations to explode, especially in deep or long RNNs.


$$H^{[t]} = \text{Activation} \left(\underset{\substack{\uparrow \\ [4, 4] \circ [4]}}{W_H^{(l)} H_{[t-1]}^{(l)}} + \underset{\substack{\uparrow \\ [4]}}{b_H^{(l)}} + \underset{\substack{\uparrow \\ [4, 3] \circ [3]}}{W^{(l)} X_{[t]}^{(l)}} + \underset{\substack{\uparrow \\ [4]}}{b^{(l)}} \right)$$

$$= [4, 4] \circ [4] + [4] + [4, 3] \circ [3] + [4]$$

$$= [4] + [4] + [4] + [4] = [4]$$

RNN Layer

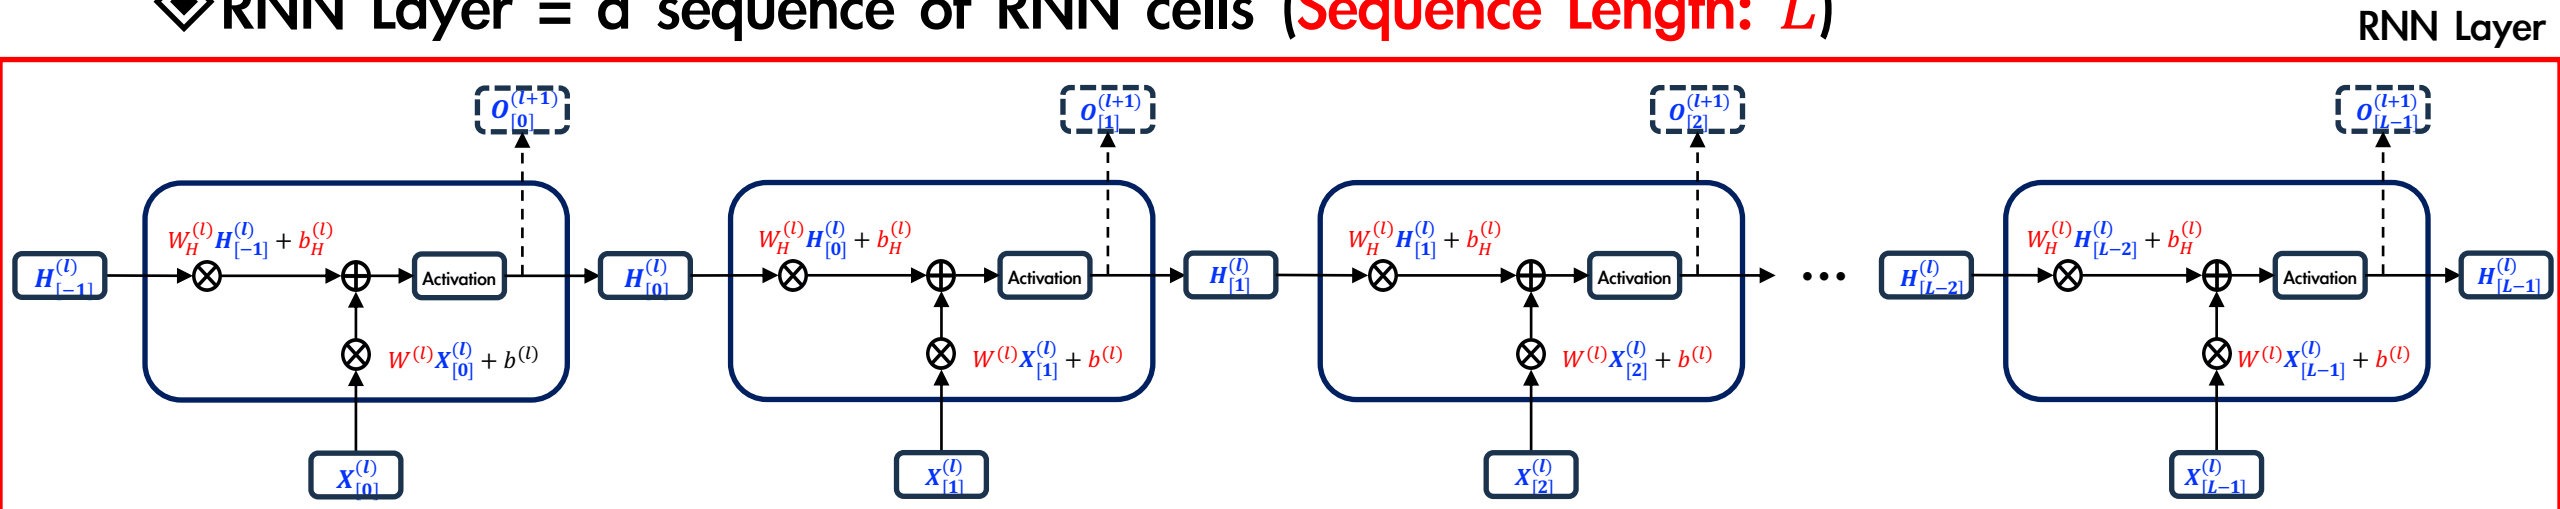
◆ RNN Layer = a sequence of RNN cells (Sequence Length: L)



	RNN Cell	RNN Layer
Definition	A single recurrent computation unit that processes one time step of the input sequence	A series of recurrent cells applied over the entire sequence
Units	Processes one time step where simple transformations are executed with weights and biases	Processes a sequence of time steps where simple transformations are executed with same weights and biases
Function	Produces an output and a hidden state	Processes the entire input sequence, producing a sequence of outputs (or hidden states)
Flow of Data	Processes input and the previous hidden state to produce a new hidden state and output	Data flows horizontally (from one time step to the next) and can also flow vertically in RNNs (from one layer to the next)

RNN Layer

◆ RNN Layer = a sequence of RNN cells (Sequence Length: L)



```
rnn_cell = nn.RNNCell(input_size=3, hidden_size=4)

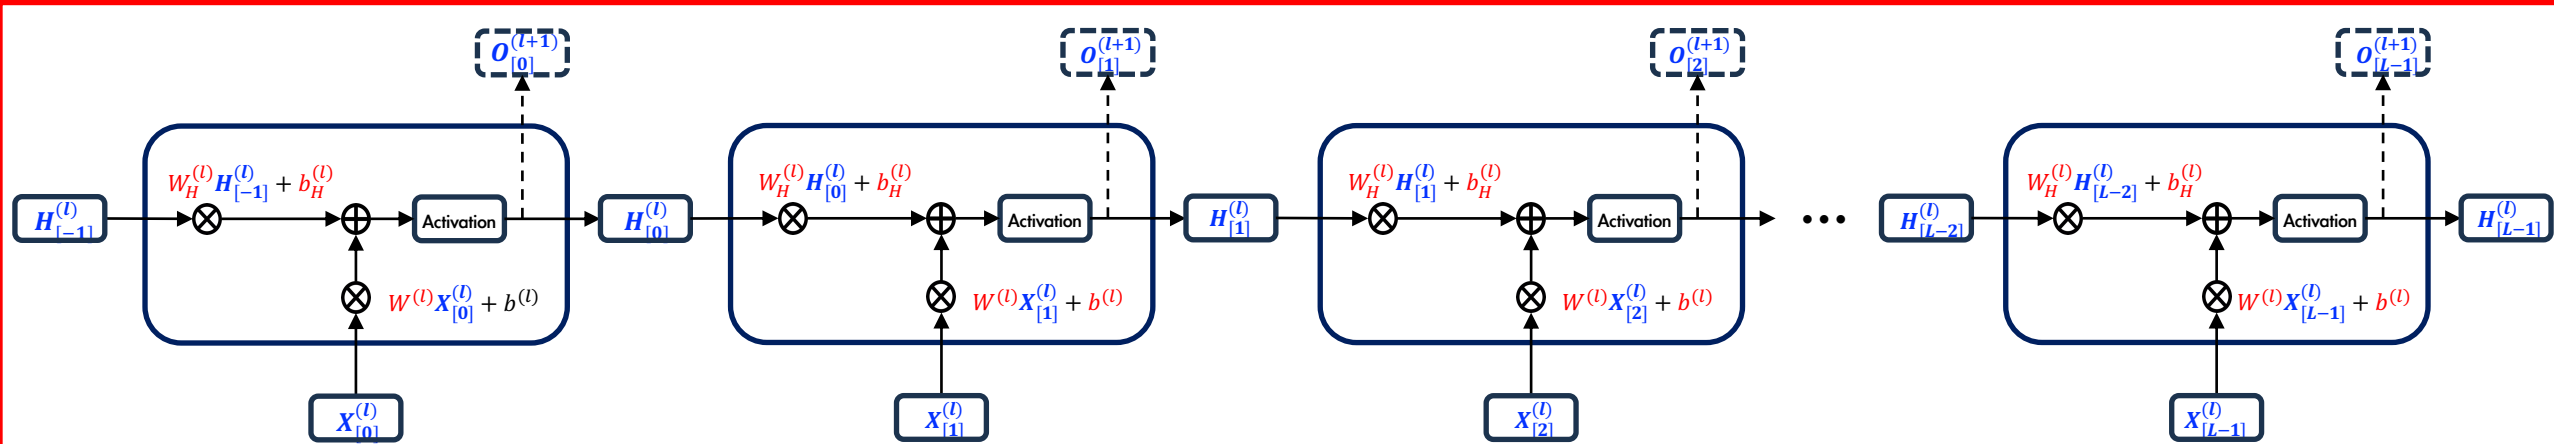
for name, parameter in rnn_cell.named_parameters():
    print(name, parameter.shape)

# >>> weight_ih torch.Size([4, 3]):  $W^{(l)}$ 
# >>> weight_hh torch.Size([4, 4]):  $W_H^{(l)}$ 
# >>> bias_ih torch.Size([4]):  $b^{(l)}$ 
# >>> bias_hh torch.Size([4]):  $b_H^{(l)}$ 
```

RNN Layer

◆ RNN Layer = a sequence of RNN cells (Sequence Length: L)

RNN Layer



```
rnn_cell = nn.RNNCell(input_size=3, hidden_size=4)
# sequence size (L): 6, input size : 3
input = torch.randn(6, 3)
```

```
# hidden size: 4
hx = torch.randn(4)
output = []
for i in range(6): # sequence size (N)
    hx = rnn_cell(input=input[i], hx=hx)
    output.append(hx)
```

```
for idx, out in enumerate(output):
    print(idx, output)
```

```
0 tensor([ 0.3005,  0.6700, -0.3974,  0.6741], grad_fn=<SqueezeBackward1>)
1 tensor([ 0.7542,  0.9379, -0.0069,  0.5461], grad_fn=<SqueezeBackward1>)
2 tensor([ 0.7395, -0.3393,  0.7367,  0.8133], grad_fn=<SqueezeBackward1>)
3 tensor([ 0.5698,  0.7913, -0.6019,  0.3906], grad_fn=<SqueezeBackward1>)
4 tensor([0.3682, 0.8590, 0.1441, 0.9075], grad_fn=<SqueezeBackward1>)
5 tensor([0.8370, 0.3408, 0.2225, 0.6136], grad_fn=<SqueezeBackward1>)
```

Sequence length (L) is not determined by the RNN cell construction

It will be determined by input data!!!


The RNN cell is rolled out based on the sequence length determined by the input data

RNN Cells & RNN Layers → RNN

RNN Stacking

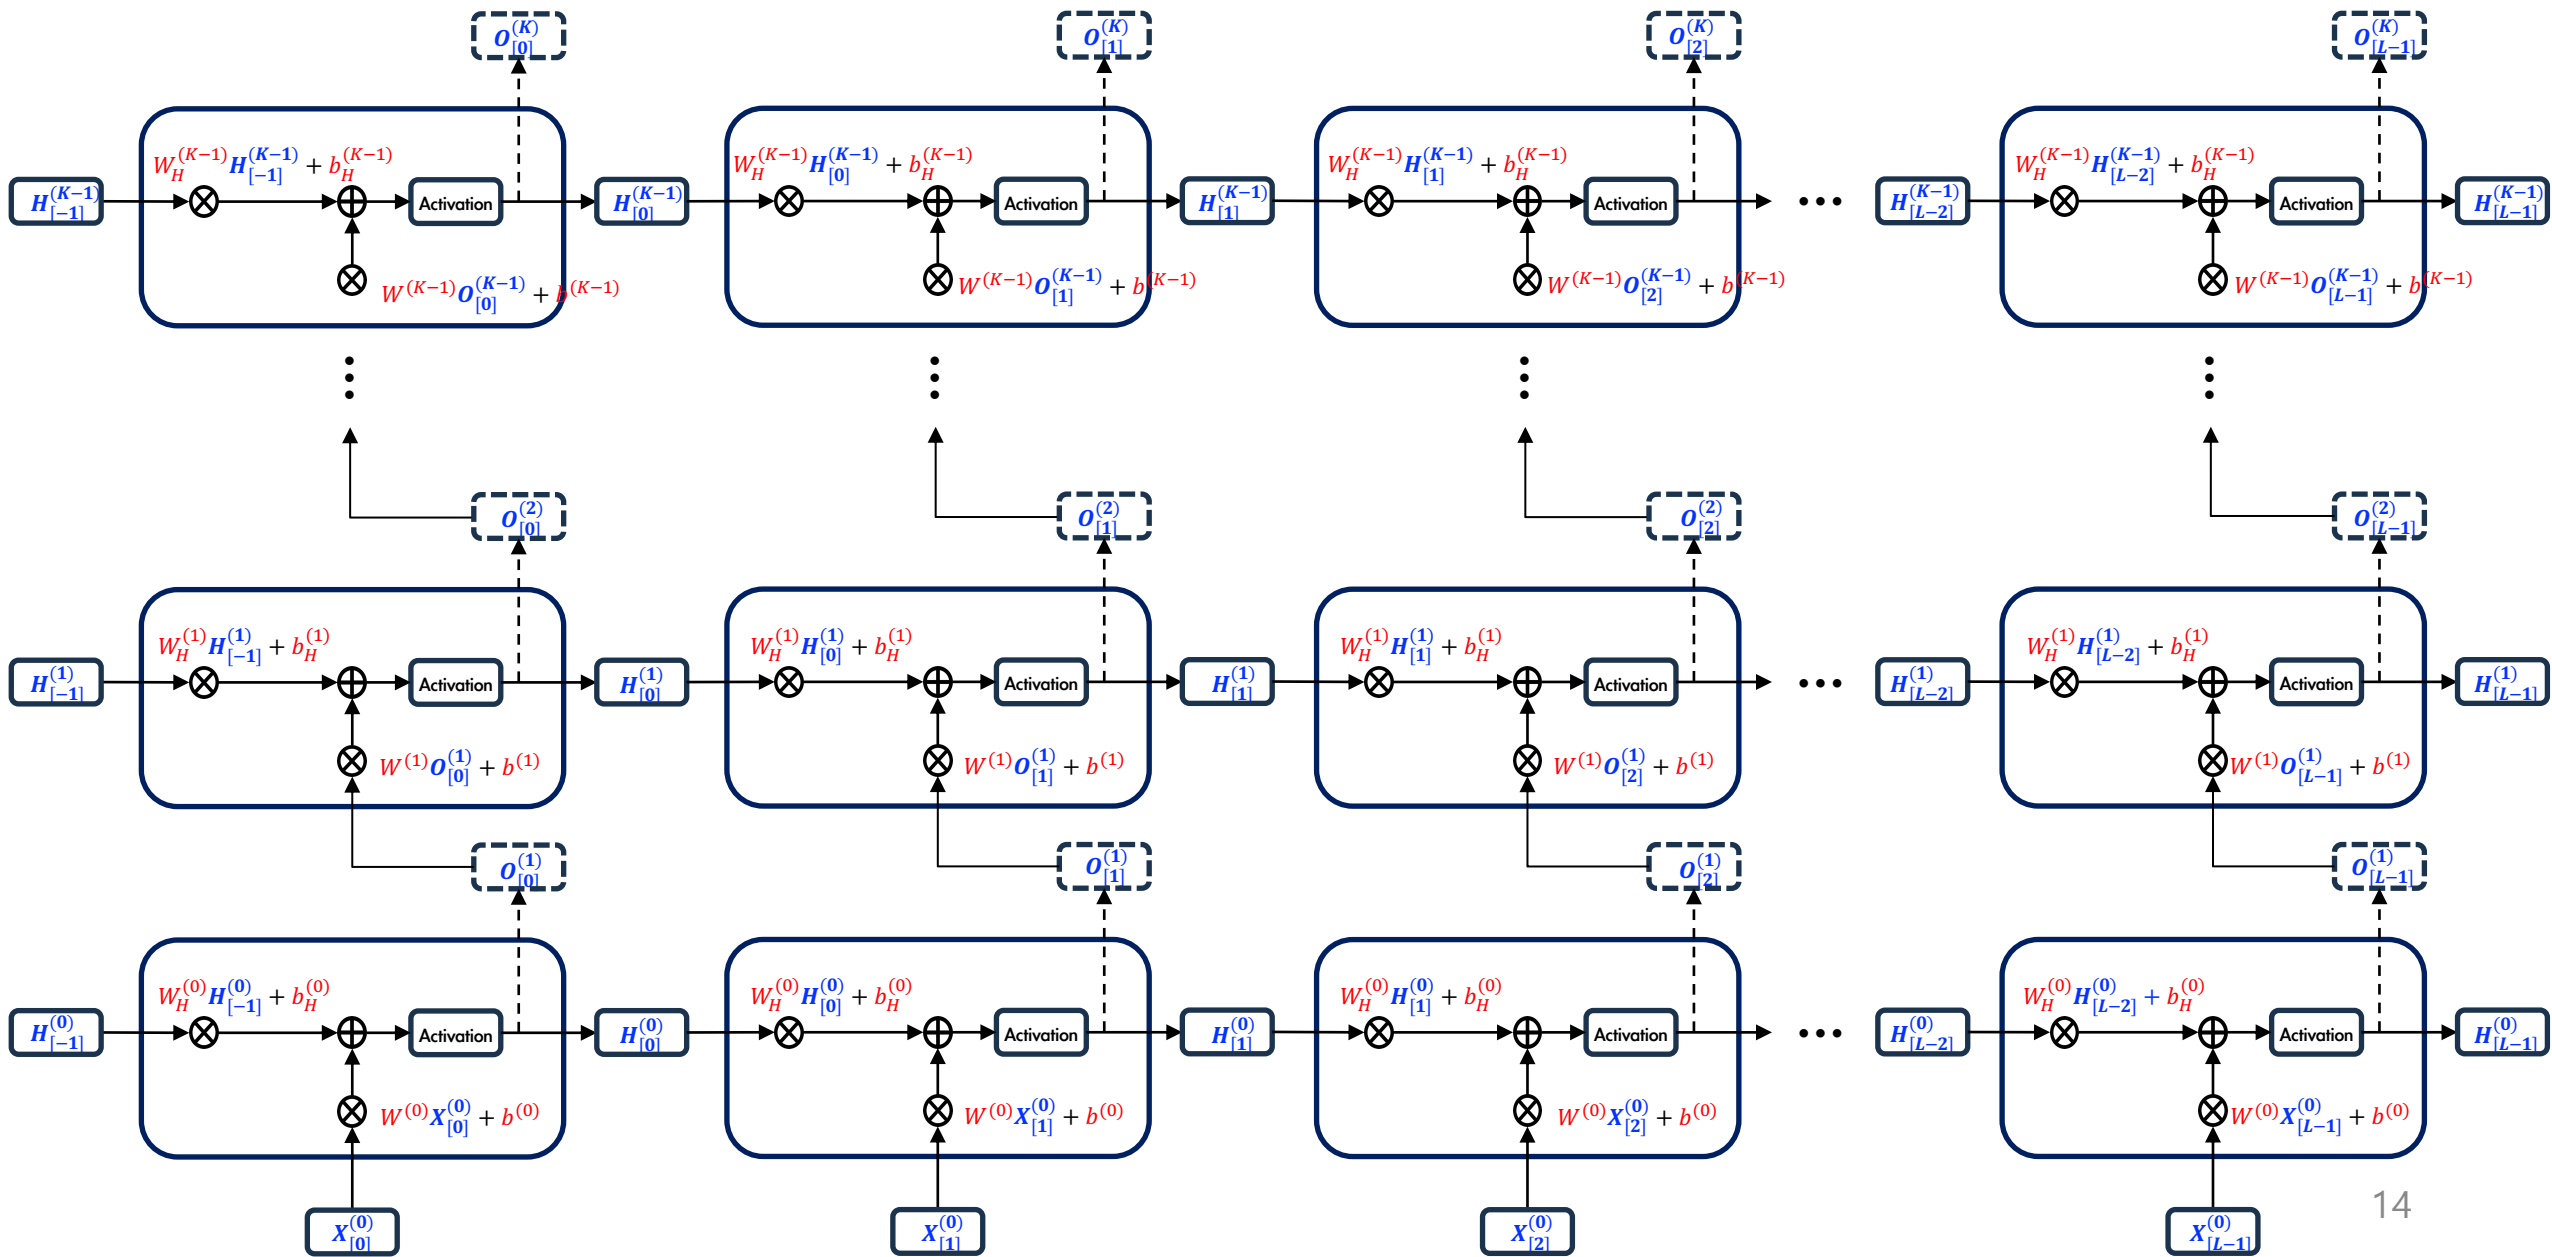
◆ RNN with Many Layers

While it is not theoretically clear what is the additional power gained by the deeper architecture, it was observed empirically that deep RNNs work better than shallower ones on some tasks.



http://cs231n.stanford.edu/slides/2018/cs231n_2018_lecture09.pdf


◆ **RNN** = a multi-layered sequence of RNN cells (**Num Layers: K** , Sequence Length: L)



RNN

◆ RNN with One Layer

Num Layers: 1



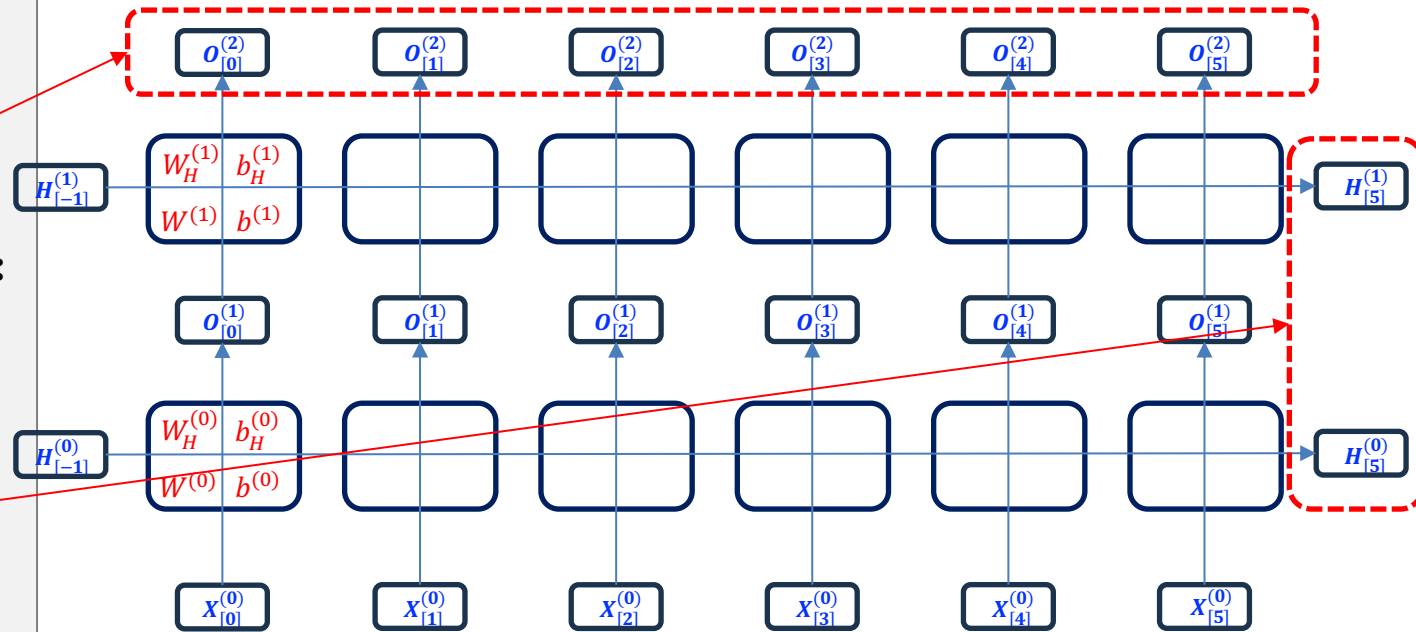
sequence size (L): 6, ~~batch size (N): ...~~, input size (F): 3

```
rnn1 = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=1  
    nonlinearity='tanh', bias=True,  
    batch_first=False, dropout=0.0,  
    bidirectional=False, device=None  
)  
  
for name, parameter in rnn1.named_parameters():  
    print(name, parameter.shape)  
# >>> weight_ih_l0 torch.Size([4, 3])  
# >>> weight_hh_l0 torch.Size([4, 4])  
# >>> bias_ih_l0 torch.Size([4])  
# >>> bias_hh_l0 torch.Size([4])  
  
input = torch.randn(6, 3)  
output, hidden_state = rnn1(input)  
for idx, out in enumerate(output):  
    print(idx, out)  
0 tensor([0.1388, 0.3009, 0.5137, 0.5397], grad_fn=<UnbindBackward0>)  
1 tensor([0.0844, 0.0305, 0.1594, 0.7344], grad_fn=<UnbindBackward0>)  
2 tensor([ 0.0070, -0.0689, 0.2532, 0.6872], grad_fn=<UnbindBackward0>)  
3 tensor([-0.0699, -0.0744, 0.1923, 0.6173], grad_fn=<UnbindBackward0>)  
4 tensor([-0.0539, 0.6625, 0.7565, 0.0793], grad_fn=<UnbindBackward0>)  
5 tensor([0.3101, 0.5119, 0.4400, 0.6907], grad_fn=<UnbindBackward0>)  
  
for idx, hidden in enumerate(hidden_state):  
    print(idx, hidden)  
0 tensor([0.3101, 0.5119, 0.4400, 0.6907], grad_fn=<UnbindBackward0>)
```

RNN

◆ RNN with Two Layers

```
rnn2 = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=2  
)  
  
for name, parameter in rnn2.named_parameters():  
    print(name, parameter.shape)  
  
# sequence size (L): 6, input size (F): 3  
input = torch.randn(6, 3)  
  
output, hidden_state = rnn2(input)  
  
for idx, out in enumerate(output):  
    print(idx, out)  
  
for idx, hidden in enumerate(hidden_state):  
    print(idx, hidden)
```



RNN knows how to handle the sequence of data, and rolls out the RNN cell automatically according to the sequence.

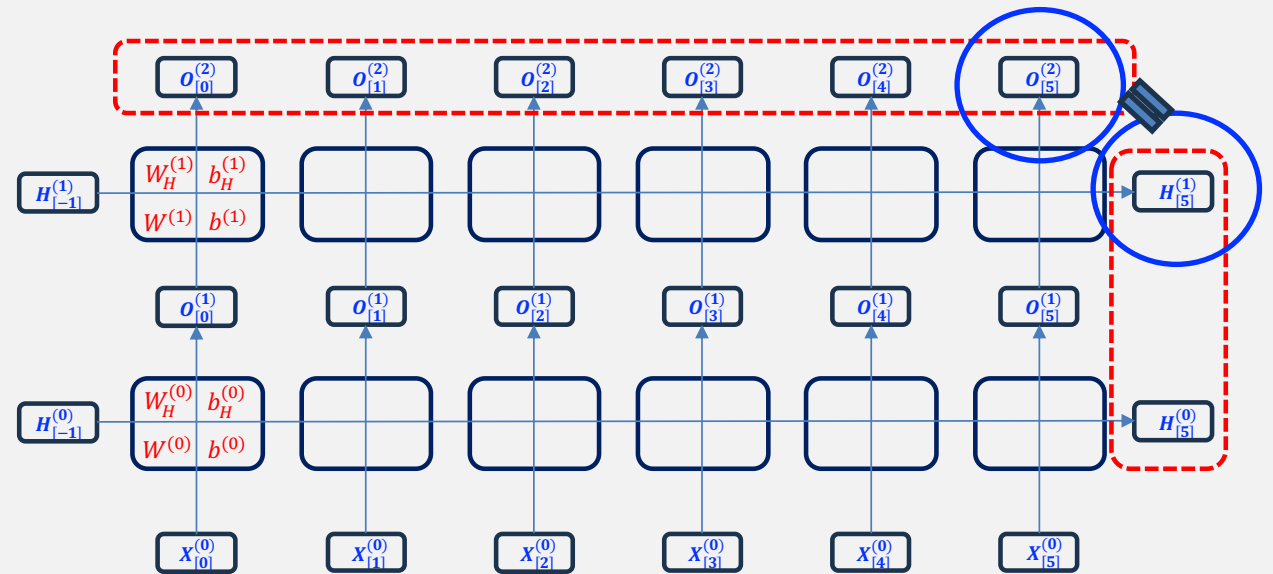
RNN

◆ RNN with Two Layers

```
rnn = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=2  
)  
  
for name, parameter in rnn.named_parameters():  
    print(name, parameter.shape)  
  
# sequence size (L): 6, input size (F): 3  
input = torch.randn(6, 3)  
  
output, hidden_state = rnn(input)  
  
for idx, out in enumerate(output):  
    print(idx, out)    # shape: torch.Size([4])  
  
for idx, hidden in enumerate(hidden_state):  
    print(idx, hidden) # shape: torch.Size([4])
```

0 tensor([0.1444, 0.5225, 0.0172, 0.3065], grad_fn=<UnbindBackward0>)
1 tensor([-0.3921, 0.4679, -0.0863, 0.2366], grad_fn=<UnbindBackward0>)
2 tensor([-0.4235, 0.2706, 0.2536, 0.1714], grad_fn=<UnbindBackward0>)
3 tensor([0.3311, 0.5493, 0.0872, 0.3736], grad_fn=<UnbindBackward0>)
4 tensor([-0.7800, 0.0835, 0.0847, -0.1501], grad_fn=<UnbindBackward0>)
5 tensor([-0.1697, -0.0223, 0.4692, 0.1598], grad_fn=<UnbindBackward0>)

0 tensor([0.2373, 0.6753, -0.6874, -0.7837], grad_fn=<UnbindBackward0>)
1 tensor([-0.1697, -0.0223, 0.4692, 0.1598], grad_fn=<UnbindBackward0>)



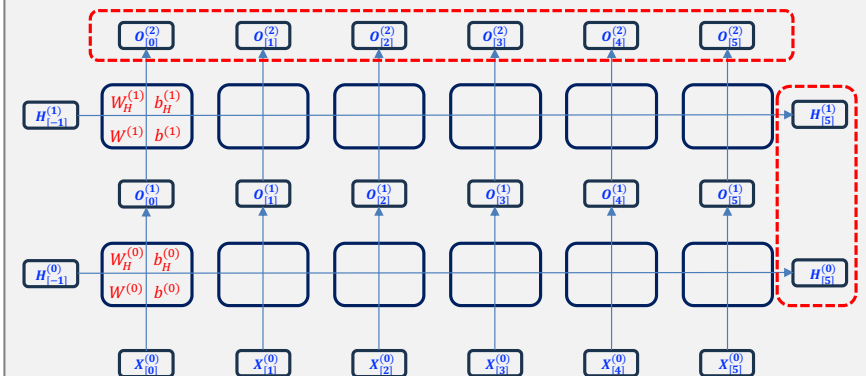
RNN

◆ Batch Inputs to RNN with Two Layers (Option 1) [Sequence, Batch, Input] $L \times N \times F$

```
rnn = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=2  
)  
  
# sequence size (L): 6, batch size (N): 10, input size (F): 3  
batch_input = torch.randn(6, 10, 3)  
  
batch_output, batch_hidden_state = rnn(batch_input)  
  
print(batch_output.shape)    # >>> torch.Size([6, 10, 4])  
for idx, out in enumerate(batch_output):  
    print(idx, out.shape)    # >>> idx torch.Size([10, 4])  
  
print()  
  
print(batch_hidden_state.shape) # >>> torch.Size([2, 10, 4])  
for idx, hidden in enumerate(batch_hidden_state):  
    print(idx, hidden.shape)  # >>> idx torch.Size([10, 4])
```

```
torch.Size([6, 10, 4])  
0 torch.Size([10, 4])  
1 torch.Size([10, 4])  
2 torch.Size([10, 4])  
3 torch.Size([10, 4])  
4 torch.Size([10, 4])  
5 torch.Size([10, 4])
```

```
torch.Size([2, 10, 4])  
0 torch.Size([10, 4])  
1 torch.Size([10, 4])
```



RNN

◆ Batch Inputs to RNN with Two Layers (Option 2) [Batch, Sequence, Input] $N \times L \times F$

```
rnn = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=2, batch_first=True  
)
```

```
# batch size (N): 10, sequence size (L): 6, input size (F): 3  
batch_input = torch.randn(10, 6, 3)
```

```
batch_output, batch_hidden_state = rnn(batch_input)
```

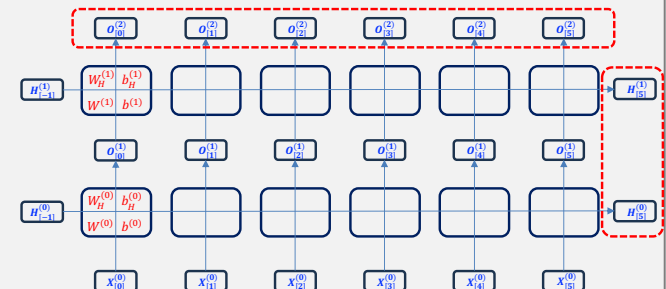
```
print(batch_output.shape)    # >>> torch.Size([10, 6, 4])  
for idx, out in enumerate(batch_output):  
    print(idx, out.shape)    # >>> idx torch.Size([6, 4])
```

```
print()
```

```
print(batch_hidden_state.shape) # >>> torch.Size([2, 10, 4])  
for idx, hidden in enumerate(batch_hidden_state):  
    print(idx, hidden.shape)   # >>> idx torch.Size([10, 4])
```

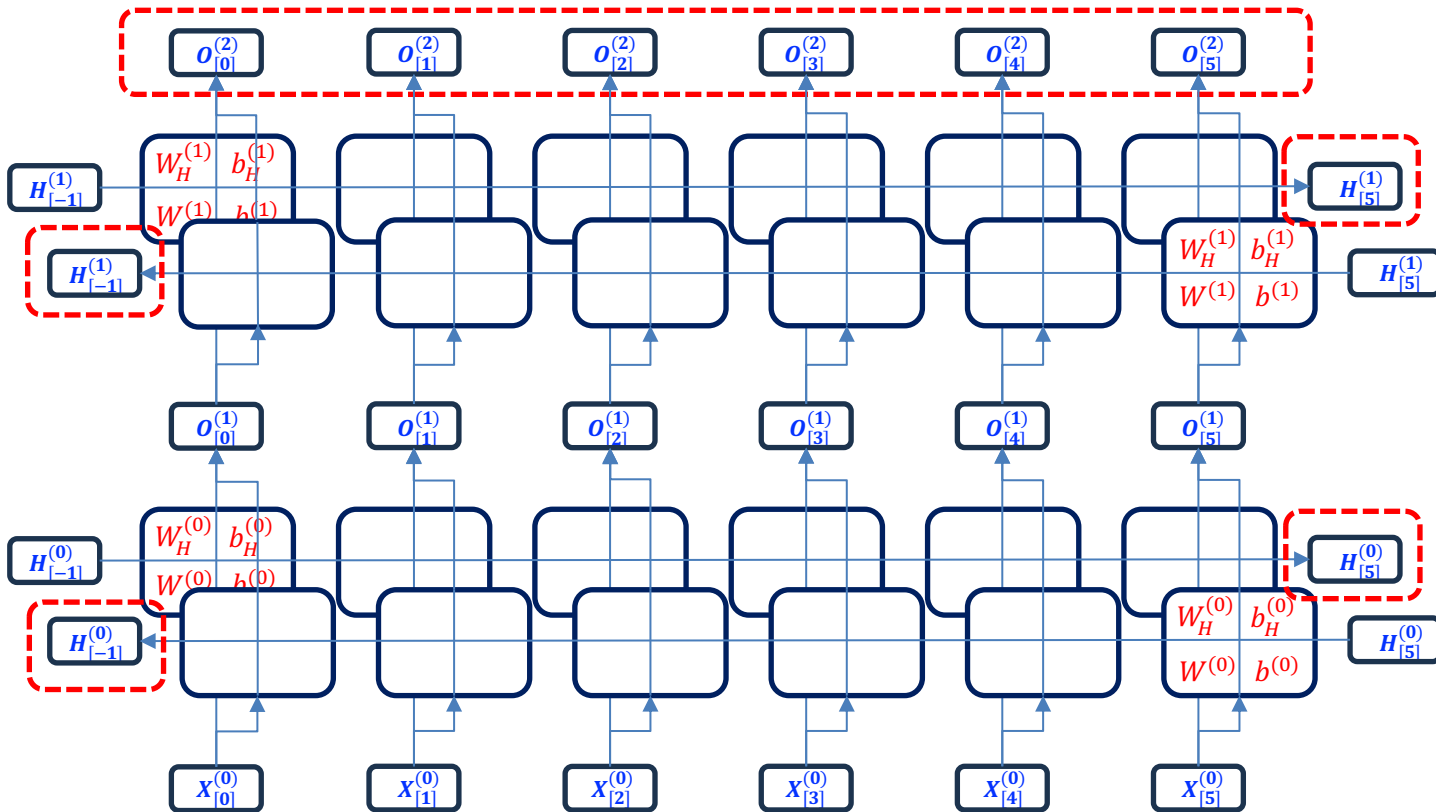
```
torch.Size([10, 6, 4])  
0 torch.Size([6, 4])  
1 torch.Size([6, 4])  
2 torch.Size([6, 4])  
3 torch.Size([6, 4])  
4 torch.Size([6, 4])  
5 torch.Size([6, 4])  
6 torch.Size([6, 4])  
7 torch.Size([6, 4])  
8 torch.Size([6, 4])  
9 torch.Size([6, 4])
```

```
torch.Size([2, 10, 4])  
0 torch.Size([10, 4])  
1 torch.Size([10, 4])
```



Bidirectional RNN

◆ Bidirectional RNN with Two Layers



```
rnn = nn.RNN(  
    input_size=3, hidden_size=4,  
    num_layers=2,  
    bidirectional=True  
)  
  
for name, parameter in rnn.named_parameters():  
    print(name, parameter.shape)
```

```
# >>> weight_ih_l0 torch.Size([4, 3])  
# >>> weight_hh_l0 torch.Size([4, 4])  
# >>> bias_ih_l0 torch.Size([4])  
# >>> bias_hh_l0 torch.Size([4])  
# >>> weight_ih_l0_reverse torch.Size([4, 3])  
# >>> weight_hh_l0_reverse torch.Size([4, 4])  
# >>> bias_ih_l0_reverse torch.Size([4])  
# >>> bias_hh_l0_reverse torch.Size([4])  
# >>> weight_ih_l1 torch.Size([4, 8])  
# >>> weight_hh_l1 torch.Size([4, 4])  
# >>> bias_ih_l1 torch.Size([4])  
# >>> bias_hh_l1 torch.Size([4])  
# >>> weight_ih_l1_reverse torch.Size([4, 8])  
# >>> weight_hh_l1_reverse torch.Size([4, 4])  
# >>> bias_ih_l1_reverse torch.Size([4])  
# >>> bias_hh_l1_reverse torch.Size([4])
```

Bidirectional RNN

◆ Bidirectional RNN with Two Layers

```
rnn = nn.RNN(
    input_size=3, hidden_size=4, num_layers=2,
    bidirectional=True
)
```

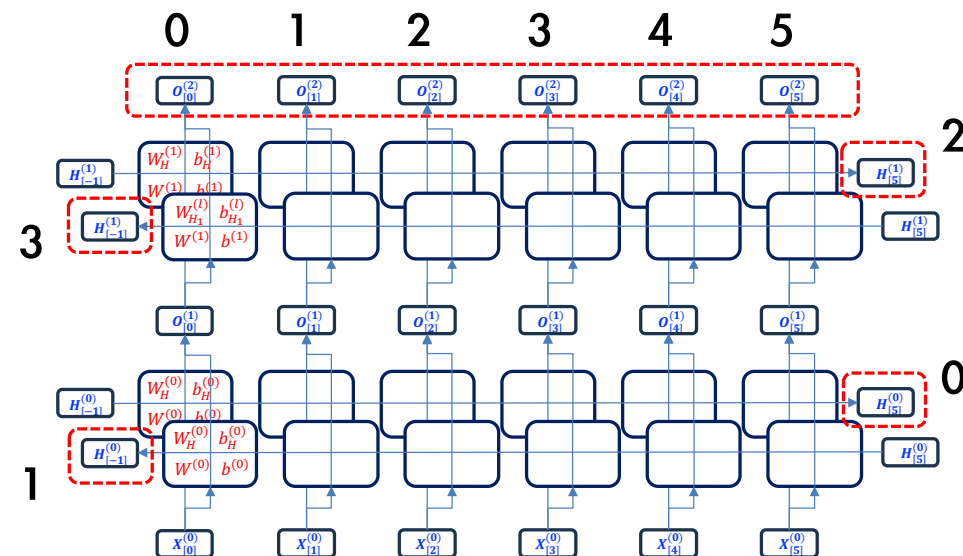
sequence size (L): 6, input size (F): 3

```
input = torch.randn(6, 3)
```

```
output, hidden_state = rnn(input)
```

```
for idx, out in enumerate(output):
    print(idx, out)    # shape: torch.Size([4])
```

```
for idx, hidden in enumerate(hidden_state):
    print(idx, hidden) # shape: torch.Size([4])
```



```
0 tensor([-0.2724, -0.7689, -0.4468, -0.1966,  0.1103,  0.3610,  0.6847,  0.0536],
grad_fn=<UnbindBackward0>)
```

```
1 tensor([ 0.4187, -0.5284, -0.0518, -0.7682,  0.6282,  0.5067,  0.6140, -0.2717],
grad_fn=<UnbindBackward0>)
```

```
2 tensor([ 0.0062,  0.4153,  0.9006,  0.1024, -0.0874,  0.7522,  0.8035,  0.6523],
grad_fn=<UnbindBackward0>)
```

```
3 tensor([ 0.0162,  0.2409,  0.5205, -0.7116,  0.1083,  0.2039,  0.6002,  0.6289],
grad_fn=<UnbindBackward0>)
```

```
4 tensor([ 0.3742,  0.1368,  0.6304, -0.6833, -0.0602,  0.0595,  0.7382,  0.5925],
grad_fn=<UnbindBackward0>)
```

```
5 tensor([ 0.2831,  0.4516,  0.6628, -0.4940, -0.6240, -0.0086,  0.6180,  0.5068],
grad_fn=<UnbindBackward0>)
```

```
0 tensor([ 0.4010, -0.9737, -0.5143,  0.5953], grad_fn=<UnbindBackward0>)
```

```
1 tensor([-0.2416,  0.3461,  0.4018, -0.5601], grad_fn=<UnbindBackward0>)
```

```
2 tensor([ 0.2831,  0.4516,  0.6628, -0.4940], grad_fn=<UnbindBackward0>)
```

```
3 tensor([0.1103, 0.3610, 0.6847, 0.0536], grad_fn=<UnbindBackward0>)
```

Bidirectional RNN

◆ Batch Inputs to Bidirectional RNN with Two Layers (Option 1) [Sequence, Batch, Input] $L \times N \times F$

```
rnn = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=2, bidirectional=True  
)
```

```
# sequence size (L): 6, batch size (N): 10, input size (F): 3  
batch_input = torch.randn(6, 10, 3)
```

```
batch_output, batch_hidden_state = rnn(batch_input)
```

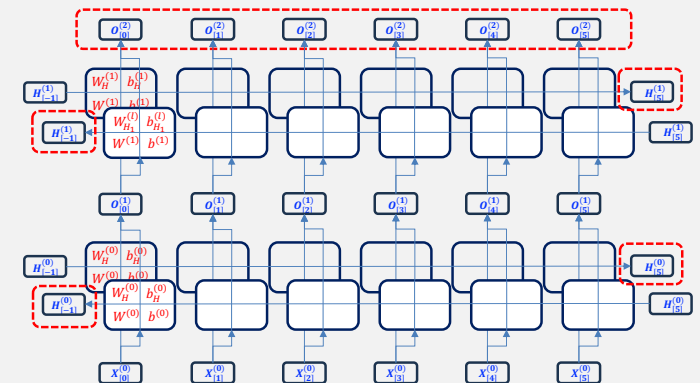
```
print(batch_output.shape)    # >>> torch.Size([6, 10, 4])  
for idx, out in enumerate(batch_output):  
    print(idx, out.shape)    # >>> idx torch.Size([10, 4])
```

```
print()
```

```
print(batch_hidden_state.shape) # >>> torch.Size([2, 10, 4])  
for idx, hidden in enumerate(batch_hidden_state):  
    print(idx, hidden.shape)   # >>> idx torch.Size([10, 4])
```

```
torch.Size([6, 10, 8])  
0 torch.Size([10, 8])  
1 torch.Size([10, 8])  
2 torch.Size([10, 8])  
3 torch.Size([10, 8])  
4 torch.Size([10, 8])  
5 torch.Size([10, 8])
```

```
torch.Size([4, 10, 4])  
0 torch.Size([10, 4])  
1 torch.Size([10, 4])  
2 torch.Size([10, 4])  
3 torch.Size([10, 4])
```



Bidirectional RNN

◆ Batch Inputs to Bidirectional RNN with Two Layers (Option 2) [Batch, Sequence, Input] $N \times L \times F$

```
rnn = nn.RNN(  
    input_size=3, hidden_size=4, num_layers=2, batch_first=True,  
    bidirectional=True  
)
```

```
# batch size (N): 10, sequence size (L): 6, input size (F): 3  
batch_input = torch.randn(10, 6, 3)
```

```
batch_output, batch_hidden_state = rnn(batch_input)
```

```
print(batch_output.shape) # >>> torch.Size([10, 6, 4])
```

```
for idx, out in enumerate(batch_output):
```

```
    print(idx, out.shape) # >>> idx torch.Size([6, 4])
```

```
print()
```

```
print(batch_hidden_state.shape) # >>> torch.Size([2, 10, 4])
```

```
for idx, hidden in enumerate(batch_hidden_state):
```

```
    print(idx, hidden.shape) # >>> idx torch.Size([10, 4])
```

```
torch.Size([10, 6, 8])
```

```
0 torch.Size([6, 8])
```

```
1 torch.Size([6, 8])
```

```
2 torch.Size([6, 8])
```

```
3 torch.Size([6, 8])
```

```
4 torch.Size([6, 8])
```


```
5 torch.Size([6, 8])
```

```
6 torch.Size([6, 8])
```

```
7 torch.Size([6, 8])
```

```
8 torch.Size([6, 8])
```

```
9 torch.Size([6, 8])
```



```
torch.Size([4, 10, 4])
```

```
0 torch.Size([10, 4])
```

```
1 torch.Size([10, 4])
```

```
2 torch.Size([10, 4])
```

```
3 torch.Size([10, 4])
```


RNN Operation Mode

Operation Mode for an RNN

◆ Modes of operation for an RNN

— Many-to-One (1/2)

- a sequence of inputs generates a single output
- This is often used in sentiment analysis, where a sequence of words is classified as expressing a positive or negative sentiment
- Time series prediction such as "bikes rental count prediction" or "bitcoin price prediction"



- An example dataset

Input dim

	Open	High	Low	Volume	Close
0	828.659973	833.450012	828.349976	1247700	831.659973
1	823.020020	828.070007	821.655029	1597800	828.070007
2	819.929993	824.400024	818.979980	1281700	824.159973
3	819.359985	823.000000	818.469971	1304000	818.979980
4	819.000000	823.000000	816.000000	1053600	820.450012
5	816.000000	820.958984	815.489990	1198100	819.239990
6	811.700012	815.250000	809.780029	1129100	813.669983
7	809.510010	810.659973	804.539978	989700	809.559998
8	807.000000	811.840027	803.190002	1155300	808.380005
9	803.989990	810.500000	801.780029	1235200	806.969971

Sequence

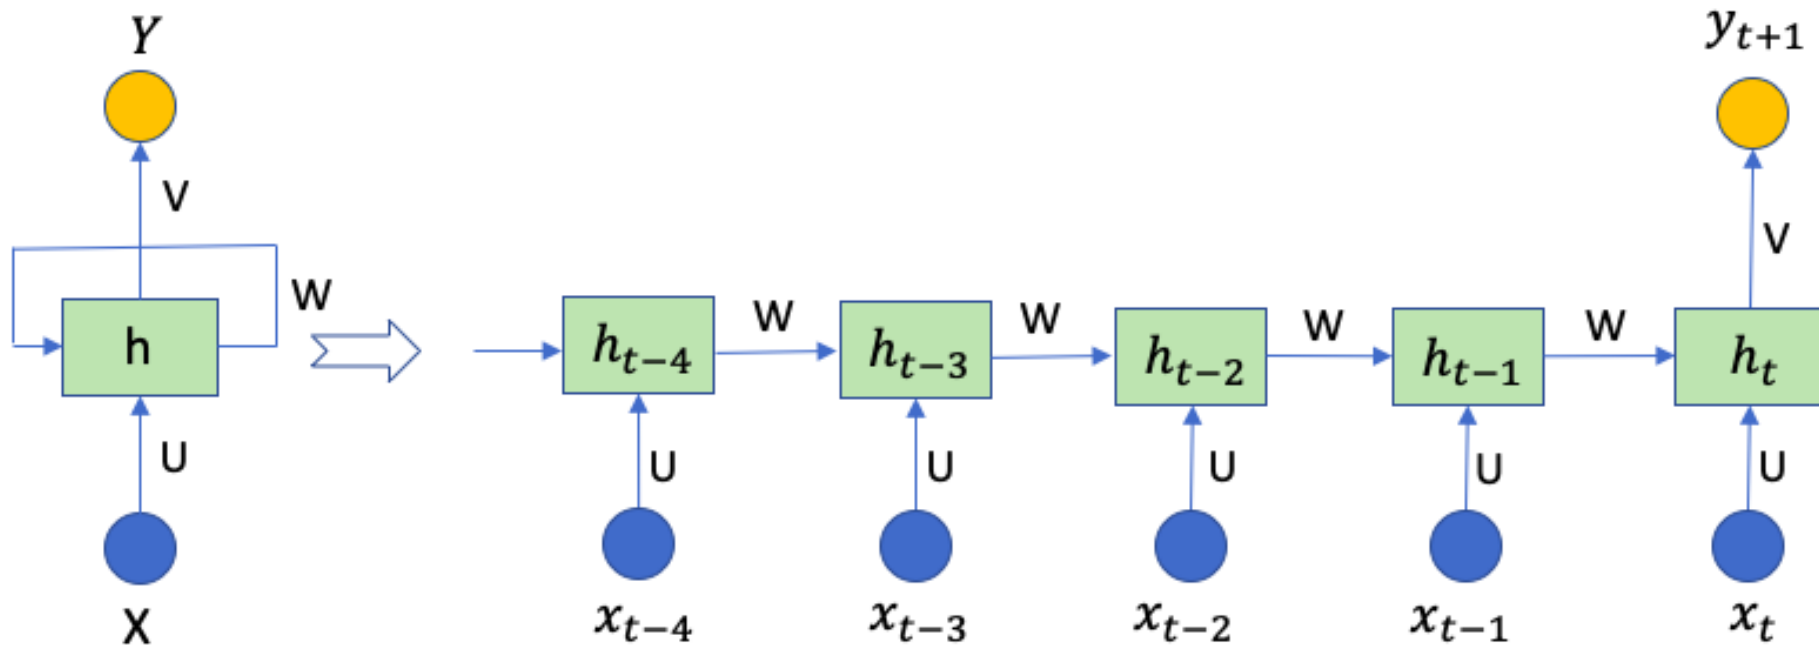
Output dim

Operation Mode for an RNN

◆ Modes of operation for an RNN

– Many-to-One (2/2)

- Model Construction and Usage




Operation Mode for an RNN

◆ Modes of operation for an RNN

— One-to-Many

- This mode is used in applications where one input may correspond to a sequence of outputs, such as image captioning, where an image input generates a sequence of words as an output
- An example dataset



```
[1, 4, 7, 10, 13, 16, 19, 22, 25, 28, 31, 34, 37, 40, 43]  
[[2, 3], [5, 6], [8, 9], [11, 12], [14, 15], [17, 18], [20, 21], [23, 24], [26, 27], [29, 30], [32, 33], [35, 36], [38, 39], [41, 42], [44, 45]]
```


Operation Mode for an RNN

◆ Modes of operation for an RNN

— Many-to-Many Synchronous (Sequence-to-Sequence or Seq2Seq Model) (1/2)

- a sequence of inputs is mapped to a sequence of outputs, used for tasks such as video classification, where every frame of the video is labeled
- An example dataset

many to many



Raw dataset

Feature 1	Feature 2	Target (Y)
1	2	10
3	4	20
5	6	30
7	8	40
9	10	50
11	12	60
13	14	70
15	16	80
17	18	90
19	20	100
21	22	110
23	24	120
25	26	130
27	28	140
29	30	150
...

Sequence data

Feature 1	Feature 2	Target (Y)
1	2	50
3	4	60
5	6	70
7	8	80
9	10	90

Feature 1	Feature 2	Target (Y)
3	4	60
5	6	70
7	8	80
9	10	90
11	12	100

Feature 1	Feature 2	Target (Y)
5	6	70
7	8	80
9	10	90
11	12	100
13	14	110

Sequential dataset

Feature 1	Feature 2	Target (Y)
5	6	70
7	8	80
9	10	90
11	12	100
13	14	110

...

Source: Author


[Batch, Sequence, Input]

$$N \times L \times F = (3, 5, 2)$$

Operation Mode for an RNN

◆ Modes of operation for an RNN

- Many-to-Many Synchronous (or Sequence-to-Sequence or Seq2Seq Model) (2/2)
 - Model Construction and Usage

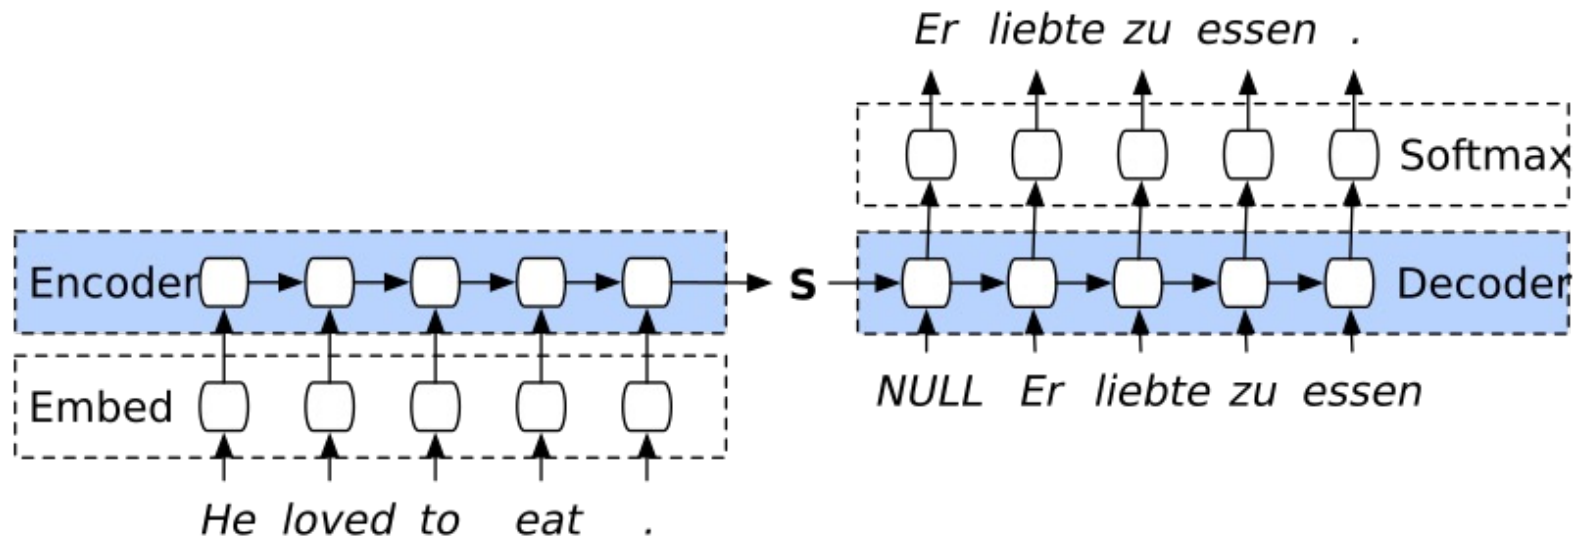



Operation Mode for an RNN

◆ Modes of operation for an RNN

— Many-to-Many Asynchronous (or Encoder-Decoder Model)

- It is used for tasks like machine translation, where an input sentence in one language is translated into a sentence in another language with potentially different word order
- An example



RNN Best Practice

- Hourly Bikes Sharing -

RNN with PyTorch

◇ Regression Trainer - init

```
from datetime import datetime
import torch
from torch import nn

from _01_code._06_fcn_best_practice.c_trainer import EarlyStopping
from _01_code._99_common_utils.utils import strfdelta

class CustomRegressionTrainer:
    def __init__(
        self, project_name, model, optimizer, train_data_loader, validation_data_loader, transforms,
        run_time_str, wandb, device, checkpoint_file_path
    ):
        self.project_name = project_name
        self.model = model
        self.optimizer = optimizer
        self.train_data_loader = train_data_loader
        self.validation_data_loader = validation_data_loader
        self.transforms = transforms
```


RNN with PyTorch

◆ Regression Trainer - init

```
class CustomRegressionTrainer:
    def __init__(
        self, project_name, model, optimizer, train_data_loader, validation_data_loader, transforms,
        run_time_str, wandb, device, checkpoint_file_path
    ):
        ...

        self.run_time_str = run_time_str
        self.wandb = wandb
        self.device = device
        self.checkpoint_file_path = checkpoint_file_path

        # Use a built-in loss function
        self.loss_fn = nn.MSELoss()
```

RNN with PyTorch

◆ Regression Trainer - do_train

```
class CustomRegressionTrainer:
    ...
    def do_train(self):
        self.model.train() # Explained at 'Diverse Techniques' section

        loss_train = 0.0
        num_trains = 0

        for train_batch in self.train_data_loader:
            input_train, target_train = train_batch
            input_train = input_train.to(device=self.device)
            target_train = target_train.to(device=self.device)

            if self.transforms:
                input_train = self.transforms(input_train)

            output_train = self.model(input_train)
```

RNN with PyTorch

◆ Regression Trainer - do_train

```
class CustomRegressionTrainer:
    ...
    def do_train(self):
        ...
        for train_batch in self.train_data_loader:
            ...
            loss = self.loss_fn(output_train.squeeze(dim=-1), target_train)
            loss_train += loss.item()

            num_trains += 1

            self.optimizer.zero_grad()
            loss.backward()
            self.optimizer.step()

        train_loss = loss_train / num_trains

    return train_loss
```

RNN with PyTorch

◆ Regression Trainer - do_validation

```
class CustomRegressionTrainer:
    ...
    def do_validation(self):
        self.model.eval()    # Explained at 'Diverse Techniques' section

        loss_validation = 0.0
        num_validations = 0

        with torch.no_grad():
            for validation_batch in self.validation_data_loader:
                input_validation, target_validation = validation_batch
                input_validation = input_validation.to(device=self.device)
                target_validation = target_validation.to(device=self.device)

            if self.transforms:
                input_validation = self.transforms(input_validation)
```

RNN with PyTorch

◆ Regression Trainer - do_validation

```
class CustomRegressionTrainer:
    ...
    def do_validation(self):
        ...
        with torch.no_grad():
            for validation_batch in self.validation_data_loader:
                ...

                output_validation = self.model(input_validation)
                loss_validation += self.loss_fn(output_validation.squeeze(dim=-1), target_validation).item()

                num_validations += 1

        validation_loss = loss_validation / num_validations

        return validation_loss
```

RNN with PyTorch

◆ Regression Trainer - train_loop

```
class CustomRegressionTrainer:
    ...
    def train_loop(self):
        early_stopping = EarlyStopping(
            patience=self.wandb.config.early_stop_patience, delta=self.wandb.config.early_stop_delta,
            project_name=self.project_name, checkpoint_file_path=self.checkpoint_file_path,
            run_time_str=self.run_time_str
        )
        n_epochs = self.wandb.config.epochs
        training_start_time = datetime.now()

        for epoch in range(1, n_epochs + 1):
            train_loss = self.do_train()

            if epoch == 1 or epoch % self.wandb.config.validation_intervals == 0:
                validation_loss = self.do_validation()
                elapsed_time = datetime.now() - training_start_time
                epoch_per_second = 1000 * epoch / elapsed_time.microseconds
```

RNN with PyTorch

◇ Regression Trainer - train_loop

```
class CustomRegressionTrainer:
    ...
    def train_loop(self):
        ...
        for epoch in range(1, n_epochs + 1):
            ...
            if epoch == 1 or epoch % self.wandb.config.validation_intervals == 0:
                ...
                message, early_stop = early_stopping.check_and_save(validation_loss, self.model)

            print(
                f"[Epoch {epoch:>3}] "
                f"T_loss: {train_loss:6.4f}, "
                f"V_loss: {validation_loss:6.4f}, "
                f"{message} | "
                f"T_time: {strfdelta(elapsed_time, '%H:%M:%S')}, "
                f"T_speed: {epoch_per_second:4.3f}"
            )
```

RNN with PyTorch

◆ Regression Trainer - train_loop

```
class CustomRegressionTrainer:
    ...
    def train_loop(self):
        ...
        for epoch in range(1, n_epochs + 1):
            ...
            if epoch == 1 or epoch % self.wandb.config.validation_intervals == 0:
                ...
                self.wandb.log({
                    "Epoch": epoch, "Training loss": train_loss, "Validation loss": validation_loss,
                    "Training speed (epochs/sec.)": epoch_per_second,
                })

            if early_stop:
                break

elapsed_time = datetime.now() - training_start_time
print(f"Final training time: {strfdelta(elapsed_time, '%H:%M:%S')}")
```


RNN with PyTorch

◆ RNN with Bikes Dataset

```
import torch
from torch import nn, optim
from torch.utils.data import DataLoader, random_split
from datetime import datetime
import os
import wandb
from pathlib import Path

BASE_PATH = str(Path(__file__).resolve().parent.parent.parent) # BASE_PATH: /Users/yhhan/git/link_dl
import sys
sys.path.append(BASE_PATH)

CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")
if not os.path.isdir(CHECKPOINT_FILE_PATH):
    os.makedirs(os.path.join(CURRENT_FILE_PATH, "checkpoints"))

from _01_code._03_real_world_data_to_tensors.o_hourly_bikes_sharing_dataset_data_loader import
get_hourly_bikes_data, HourlyBikesDataset
from _01_code._10_rnn.g_rnn_trainer import RegressionTrainer
from _01_code._10_rnn.f_arg_parser import get_parser
```

RNN with PyTorch

◆ RNN with Bikes Dataset

```
def get_train_bikes_data():
    X_train, X_validation, X_test, y_train, y_validation, y_test = get_hourly_bikes_data(
        sequence_size=24, validation_size=96, test_size=24, y_normalizer=100
    )

    train_hourly_bikes_dataset = HourlyBikesDataset(X=X_train, y=y_train)
    validation_hourly_bikes_dataset = HourlyBikesDataset(X=X_validation, y=y_validation)

    train_data_loader = DataLoader(
        dataset=train_hourly_bikes_dataset, batch_size=wandb.config.batch_size, shuffle=True
    )

    validation_data_loader = DataLoader(
        dataset=validation_hourly_bikes_dataset, batch_size=wandb.config.batch_size, shuffle=True
    )

    return train_data_loader, validation_data_loader
```

RNN with PyTorch

◆ RNN with Bikes Dataset

```
def get_model():  
    class MyModel(nn.Module):  
        def __init__(self, n_input, n_output):  
            super().__init__()  
  
            self.rnn = nn.RNN(  
                input_size=n_input, hidden_size=128, num_layers=2, batch_first=True  
            )  
            self.fcn = nn.Linear(in_features=128, out_features=n_output)  
  
        def forward(self, x):  
            x, hidden = self.rnn(x)  
            x = x[:, -1, :] # x.shape: [32, 128]  
            x = self.fcn(x)  
            return x  
  
    my_model = MyModel(n_input=18, n_output=1)  
    return my_model
```

RNN with PyTorch

◆ RNN with Bikes Dataset

```
def main(args):
    run_time_str = datetime.now().astimezone().strftime('%Y-%m-%d_%H-%M-%S')
    config = {
        'epochs': args.epochs, 'batch_size': args.batch_size,
        'validation_intervals': args.validation_intervals,
        'learning_rate': args.learning_rate,
        'early_stop_patience': args.early_stop_patience
        'early_stop_delta': args.early_stop_delta,
    }
    project_name = "rnn_bikes"
    wandb.init(
        mode="online" if args.wandb else "disabled",
        project=project_name,    notes="bikes experiment with rnn",
        tags=["rnn", "bikes"],  name=run_time_str,
        config=config
    )
    print(args)
    print(wandb.config)
```

RNN with PyTorch

◆ RNN with Bikes Dataset

```
def main(args):  
    ...  
    train_data_loader, validation_data_loader = get_train_bikes_data()  
    device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")  
    print(f"Training on device {device}.")  
    model = get_model()  
    model.to(device)  
    wandb.watch(model)  
  
    optimizer = optim.Adam(model.parameters(), lr=wandb.config.learning_rate)  
  
    classification_trainer = CustomRegressionTrainer(  
        project_name, model, optimizer, train_data_loader, validation_data_loader, None,  
        run_time_str, wandb, device, CHECKPOINT_FILE_PATH  
    )  
  
    classification_trainer.train_loop()  
    wandb.finish()
```

RNN with PyTorch

◆ RNN with Bikes Dataset

```
if __name__ == "__main__":  
    parser = get_parser()  
    args = parser.parse_args()  
    main(args)  
# python _01_code/_10_rnn/h_bikes_train_rnn.py -v 100
```

RNN Test with PyTorch

◆ RNN Test with Test Bikes Dataset

```
def test_main(test_model):  
    _, _, X_test, _, _, y_test = get_hourly_bikes_data(  
        sequence_size=24, validation_size=96, test_size=24, y_normalizer=100  
    )  
  
    test_hourly_bikes_dataset = HourlyBikesDataset(X=X_test, y=y_test)  
  
    test_data_loader = DataLoader(  
        dataset=test_hourly_bikes_dataset, batch_size=len(test_hourly_bikes_dataset)  
    )  
  
    test_model.eval()  
  
    y_normalizer = 100  
  
    print("[TEST DATA]")
```

RNN Test with PyTorch

◆ RNN Test with Test Bikes Dataset

```
def test_main(test_model):  
    ...  
    with torch.no_grad():  
        for test_batch in test_data_loader:  
            input_test, target_test = test_batch  
  
            output_test = test_model(input_test)  
  
            for idx, (output, target) in enumerate(zip(output_test, target_test)):  
                output = round(output.item() * y_normalizer)  
                target = target.item() * y_normalizer  
  
                print("{0:2}: {1:6,.2f} <--> {2:6,.2f} (Loss: {3:>13,.2f})".format(  
                    idx, output, target, abs(output - target)  
                ))
```


RNN Test with PyTorch

◆ RNN Test with Test Bikes Dataset

```
def predict_all(test_model):
    y_normalizer = 100
    X_train, X_validation, X_test, y_train, y_validation, y_test = get_hourly_bikes_data(
        sequence_size=24, validation_size=96, test_size=24, y_normalizer=100]
    )
    train_hourly_bikes_dataset = HourlyBikesDataset(X=X_train, y=y_train)
    validation_hourly_bikes_dataset = HourlyBikesDataset(X=X_validation, y=y_validation)
    test_hourly_bikes_dataset = HourlyBikesDataset(X=X_test, y=y_test)

    dataset_list = [
        train_hourly_bikes_dataset, validation_hourly_bikes_dataset, test_hourly_bikes_dataset
    ]
    dataset_labels = [
        "train", "validation", "test"
    ]

    num = 0
    fig, axs = plt.subplots(3, 1, figsize=(6, 9))
```

RNN Test with PyTorch

◆ RNN Test with Test Bikes Dataset

```
def predict_all(test_model):
    for i in range(3):
        X = []
        TARGET_Y = []
        PREDICTION_Y = []
        for data in dataset_list[i]:
            input, target = data
            prediction = test_model(input.unsqueeze(0)).squeeze(-1).squeeze(-1)
            X.append(num)
            TARGET_Y.append(target.item() * y_normalizer)
            PREDICTION_Y.append(prediction.item() * y_normalizer)
            num += 1
        axs[i].plot(X, TARGET_Y, label='target')
        axs[i].plot(X, PREDICTION_Y, label='prediction')
        axs[i].set_title(dataset_labels[i])
        axs[i].legend()
    plt.tight_layout()
    plt.show()
```