

# Transfer Learning

December 2023

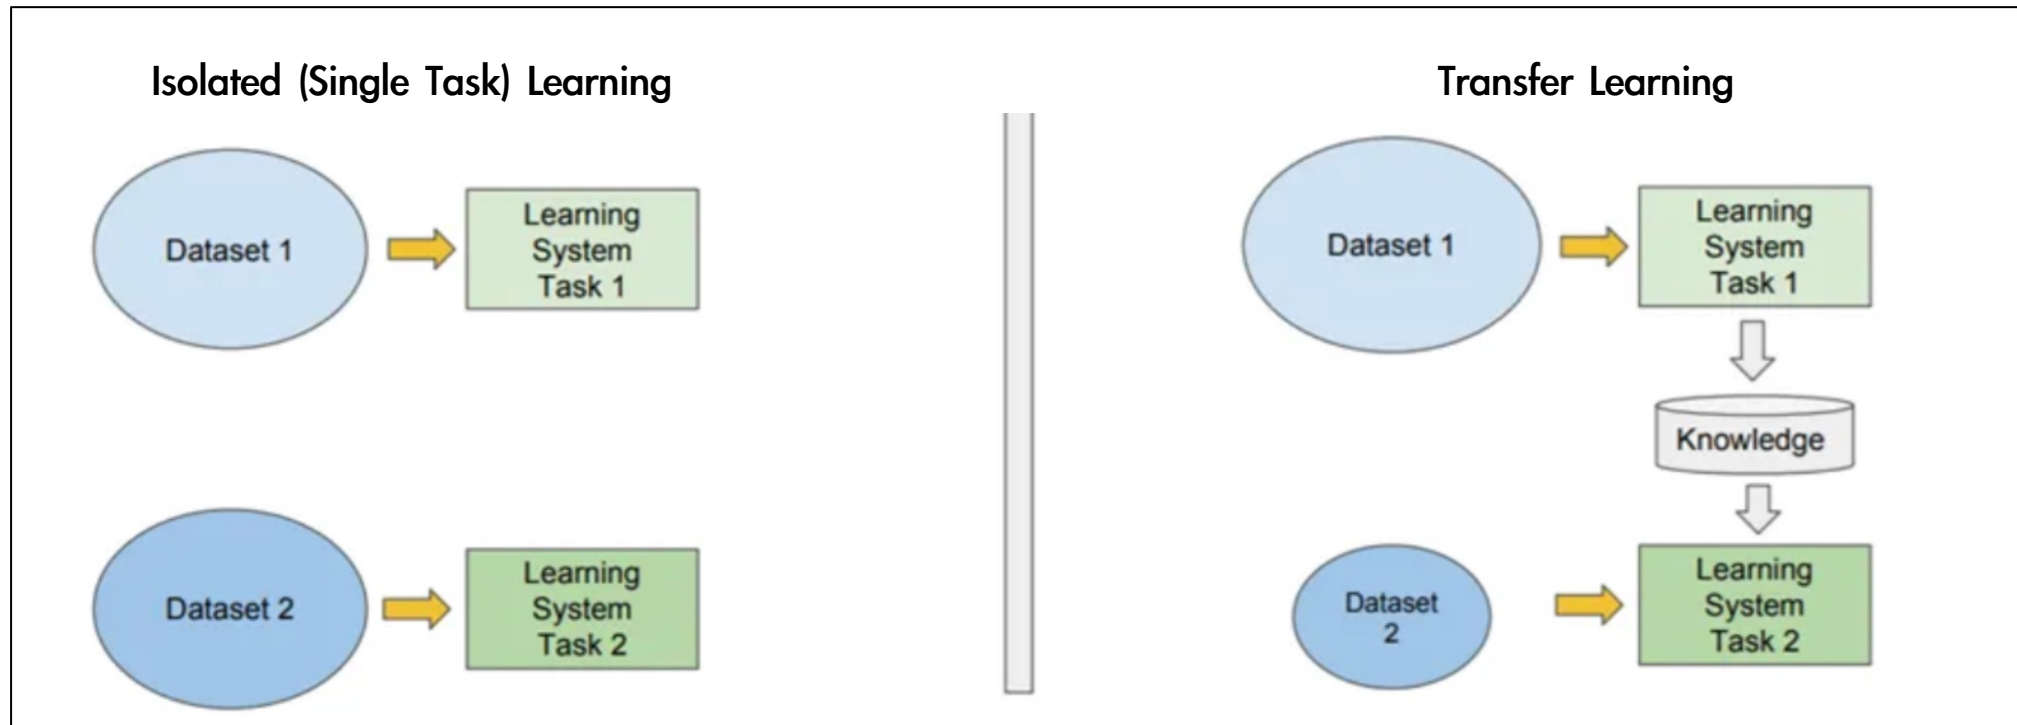
<http://link.koreatech.ac.kr>

# Transfer Learning

# Transfer Learning

## ◆ What is Transfer Learning?

- Transferring the knowledge of one model to perform a new task
- Reuse of a pre-trained model on a new task
- a.k.a "Domain Adaptation"

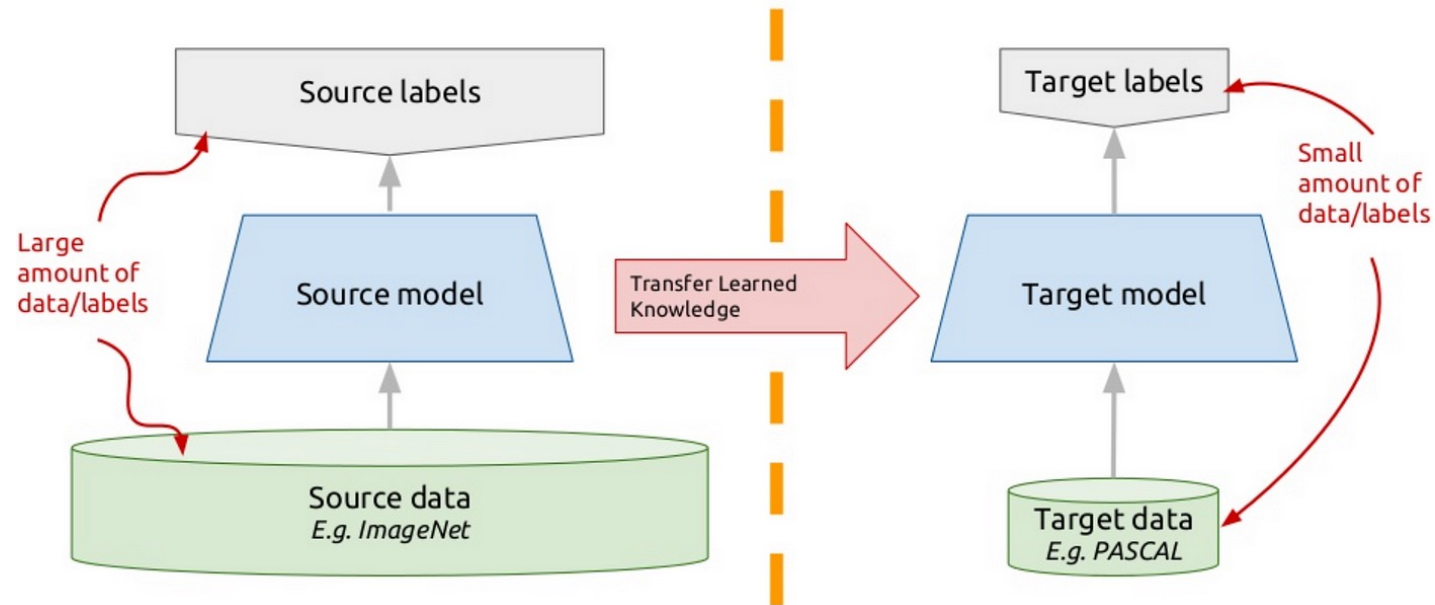


# Transfer Learning

## ◆ Motivation of Transfer Learning

- Lots of data, time, resources needed to train and tune a neural network from scratch
  - An ImageNet deep neural network can take weeks to train and fine-tune from scratch
  - Unless you have 256 GPUs, it is not possible to achieve in 1 hour
- Cheaper & faster way of adapting a neural network by exploiting their generalization properties


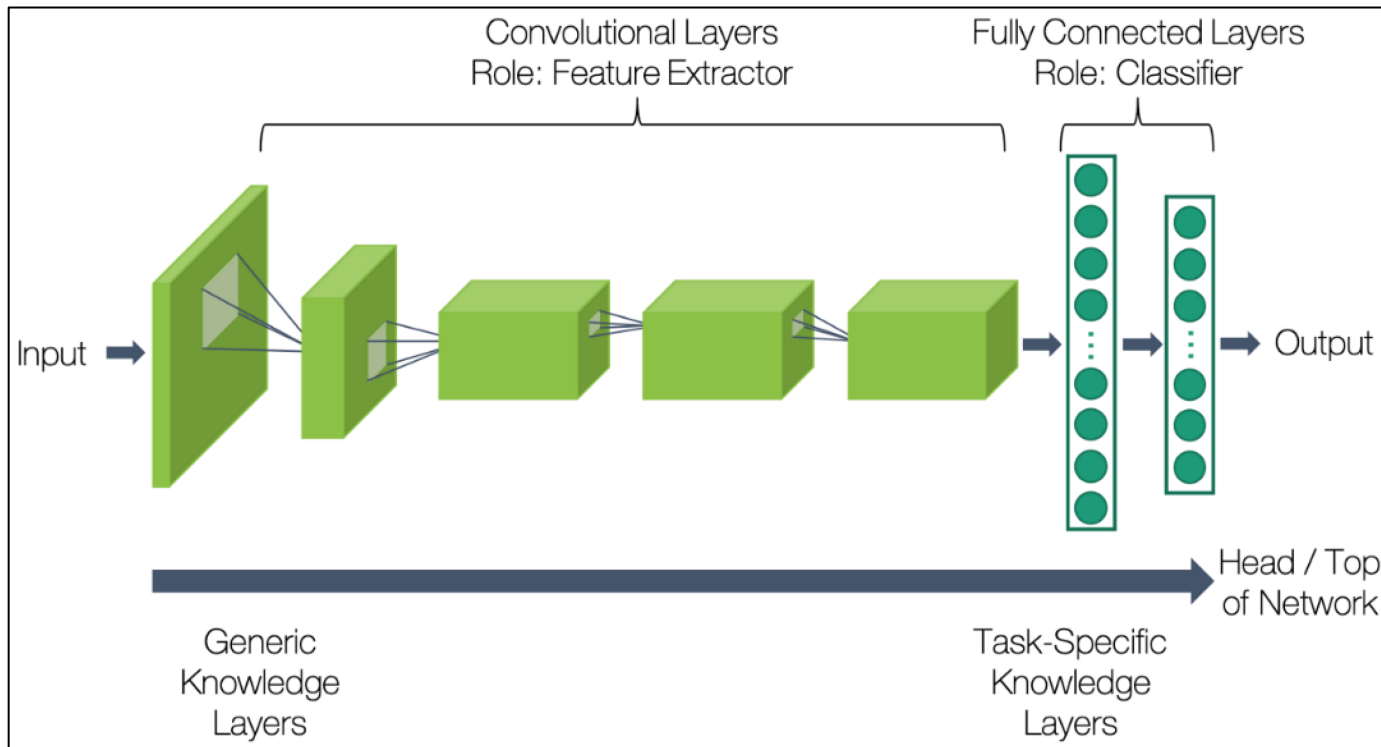
### Transfer learning: idea



# Transfer Learning

## ◆ Neural Network Layers: General to Specific

- Bottom/first/earlier layers: general learners
  - Low-level notions of edges, visual shapes
- Top/last/later layers: specific learners
  - High-level features such as eyes, feathers

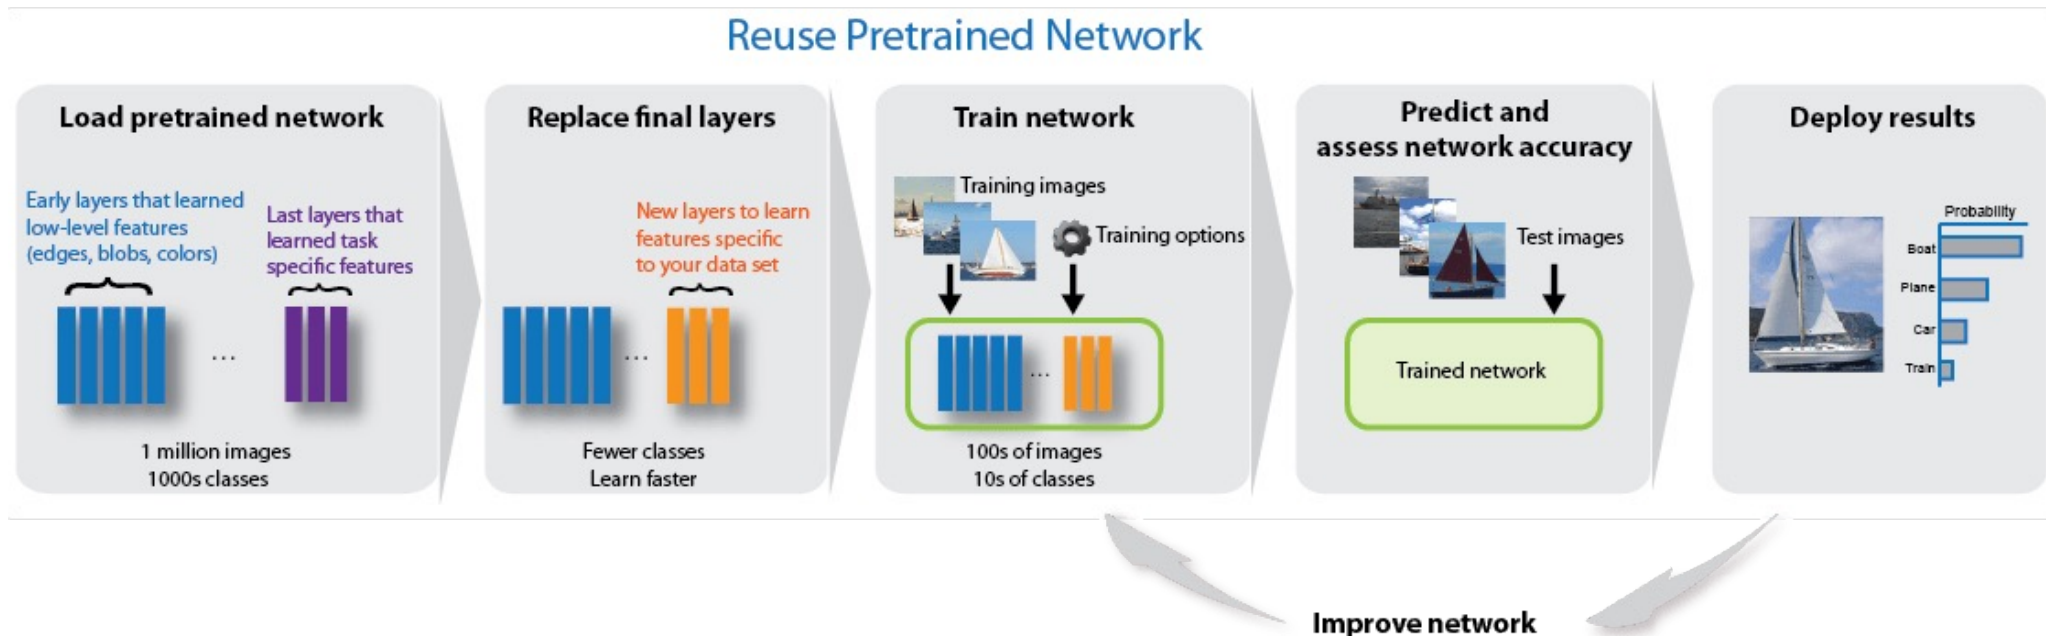


**Example:**  
**VGG 16 Filters**

# Transfer Learning

## ◆ Transfer Learning Process (1/2)

- 1. Start with a pre-trained network
- 2. Partition network into:
  - **Featurizers**: identify which layers to keep (feature extraction layers)
  - **Classifiers**: identify which layers to replace
    - In many cases, we replace the old layers with new ones (with different numbers of outputs) on top of featurizers





# Transfer Learning

## ◆ Transfer Learning Process (2/2)

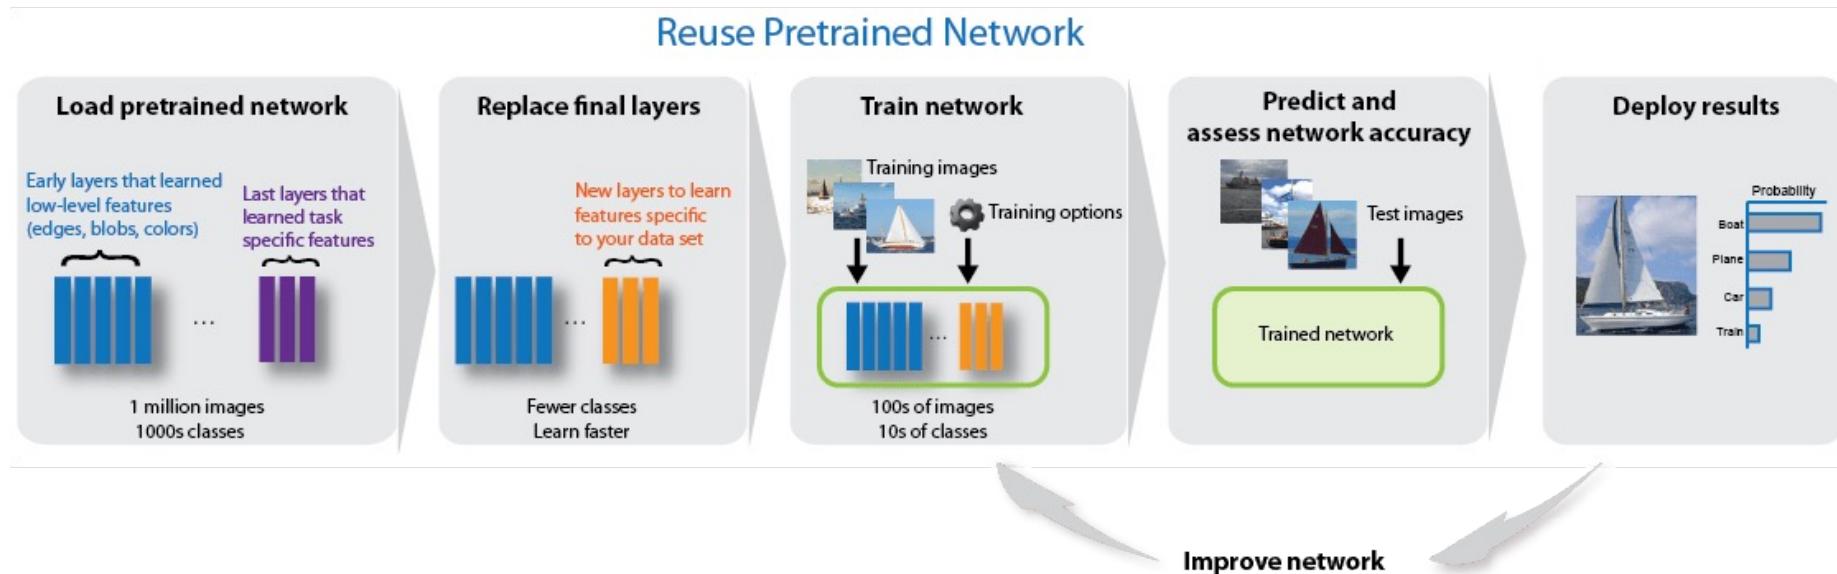
### – 3. Two major transfer learning methods

1) Frozen the **featurizers** and Train the **classifiers** (including new layers) with new data

- It is common to pretrain a ConvNet on a very large dataset (e.g. ImageNet, which contains 1.2 million images with 1000 categories), and then use the ConvNet as a **featurizers** for the task of interest

2) **Fine-tune** the whole model with new data

- Unfreeze weights and fine-tune the whole network with smaller learning rate
- It is not common method because it is relatively rare to have a dataset of sufficient size



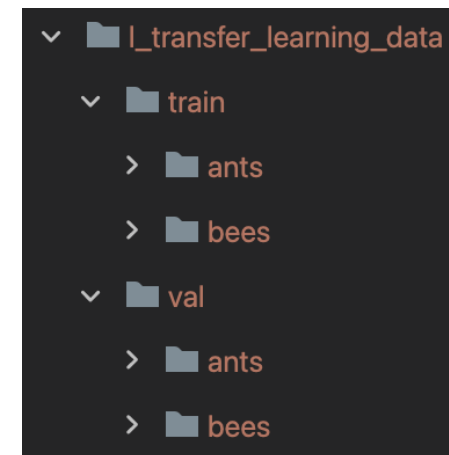
# **Transfer Learning with PyTorch**



# Data Preparation

## ◆ Data Preparation

- We will use torchvision and torch.utils.data packages for loading the data
- The problem we're going to solve today is to train a model to classify **ants** and **bees**
  - There are about 120 training images each for ants and bees
  - There are 75 validation images each for ants and bees
- Usually, this is a very small dataset to generalize upon, if trained from scratch.
- Since we are using transfer learning, we should be able to generalize reasonably well
- Download the data from here and extract it to the predefined folder
  - [https://download.pytorch.org/tutorial/hymenoptera\\_data.zip](https://download.pytorch.org/tutorial/hymenoptera_data.zip)
  - Make "\_00\_data/l\_transfer\_learning\_data" and store the unzipped files into the new folder



# Transfer Learning with PyTorch

## ◆ Preparation

```
import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
import numpy as np
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
import time, os, sys
from pathlib import Path

BASE_PATH = str(Path(__file__).resolve().parent.parent.parent) # BASE_PATH: /Users/yhhan/git/link_dl
sys.path.append(BASE_PATH)

CURRENT_FILE_PATH = os.path.dirname(os.path.abspath(__file__))
CHECKPOINT_FILE_PATH = os.path.join(CURRENT_FILE_PATH, "checkpoints")

if not os.path.isdir(CHECKPOINT_FILE_PATH):
    os.makedirs(os.path.join(CURRENT_FILE_PATH, "checkpoints"))
```

# Transfer Learning with PyTorch

## ◆ Data Transforms

```
data_transforms = {  
    'train': transforms.Compose([  
        transforms.RandomResizedCrop(224),  
        transforms.RandomHorizontalFlip(),  
        transforms.ToTensor(),  
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])  
    ]),  
  
    'val': transforms.Compose([  
        transforms.Resize(256),  
        transforms.CenterCrop(224),  
        transforms.ToTensor(),  
        transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])  
    ])  
}
```

# Transfer Learning with PyTorch

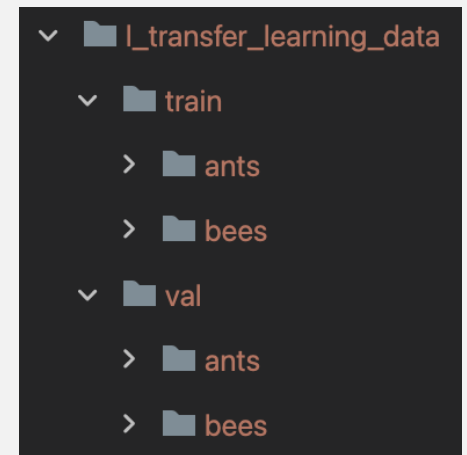
## ◆ Get New Data

```
def get_new_data():
    new_data_path = os.path.join(BASE_PATH, "_00_data", "l_transfer_learning_data")
    image_datasets = {
        x: datasets.ImageFolder(os.path.join(new_data_path, x), data_transforms[x])
        for x in ['train', 'val']
    }

    dataset_sizes = {
        x: len(image_datasets[x]) for x in ['train', 'val']
    }

    dataloaders = {
        x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4, shuffle=True, num_workers=0)
        for x in ['train', 'val']
    }

    class_names = image_datasets['train'].classes
    return image_datasets, dataset_sizes, dataloaders, class_names
```



# Transfer Learning with PyTorch

## ◆ Train

```
def train_model(
    dataloaders, dataset_sizes, model, loss_fn, optimizer, scheduler, num_epochs=25, device=torch.device("cpu")
):
    since = time.time()

    best_model_params_path = os.path.join(CHECKPOINT_FILE_PATH, 'best_model_params.pt')

    best_acc = 0.0

    for epoch in range(num_epochs):
        print(f'Epoch {epoch}/{num_epochs - 1}')
        print('-' * 10)

        # Each epoch has a training and validation phase
        for phase in ['train', 'val']:
            if phase == 'train':
                model.train()  # Set model to training mode
            else:
                model.eval()  # Set model to evaluate mode
```

# Transfer Learning with PyTorch

## ◆ Train

```
def train_model(
    dataloaders, dataset_sizes, model, loss_fn, optimizer, scheduler, num_epochs=25, device=torch.device("cpu")
):
    ...
    for epoch in range(num_epochs):
        ...
        for phase in ['train', 'val']:
            ...
            running_loss = 0.0
            running_corrects = 0

            # Iterate over data
            for inputs, labels in dataloaders[phase]:
                inputs = inputs.to(device)
                labels = labels.to(device)

            # zero the parameter gradients
            optimizer.zero_grad()
```



# Transfer Learning with PyTorch

## ◆ Train

```
def train_model(
    dataloaders, dataset_sizes, model, loss_fn, optimizer, scheduler, num_epochs=25, device=torch.device("cpu")
):
    ...
    for epoch in range(num_epochs):
        ...
        for phase in ['train', 'val']:
            ...
            for inputs, labels in dataloaders[phase]:
                ...
                # track history if only in train
                with torch.set_grad_enabled(phase == 'train'):
                    outputs = model(inputs)
                    _, preds = torch.max(outputs, 1)
                    loss = criterion(outputs, labels)
                # backward + optimize only if in training phase
                if phase == 'train':
                    loss.backward()
                    optimizer.step()
```

# Transfer Learning with PyTorch

## ◆ Train

```
def train_model(
    dataloaders, dataset_sizes, model, loss_fn, optimizer, scheduler, num_epochs=25, device=torch.device("cpu")
):
    ...
    for epoch in range(num_epochs):
        ...
        for phase in ['train', 'val']:
            ...
            for inputs, labels in dataloaders[phase]:
                ...
                # statistics
                running_loss += loss.item() * inputs.size(0)
                running_corrects += torch.sum(preds == labels.data)

            if phase == 'train':
                scheduler.step() # Scheduling the learning rate

        epoch_loss = running_loss / dataset_sizes[phase]
        epoch_acc = running_corrects.double() / dataset_sizes[phase]
```

# Transfer Learning with PyTorch

## ◆ Train

```
def train_model(
    dataloaders, dataset_sizes, model, loss_fn, optimizer, scheduler, num_epochs=25, device=torch.device("cpu")
):
    ...
    for epoch in range(num_epochs):
        ...
        for phase in ['train', 'val']:
            ...
            print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')
            # deep copy the model
            if phase == 'val' and epoch_acc > best_acc:
                best_acc = epoch_acc
                torch.save(model.state_dict(), best_model_params_path)

        print()
    time_elapsed = time.time() - since
    print(f'Training complete in {time_elapsed // 60:.0f}m {time_elapsed % 60:.0f}s')
    print(f'Best val Acc: {best_acc:4f}\n')
```

# Transfer Learning with PyTorch

## ◆ Train

```
def train_model(
    dataloaders, dataset_sizes, model, loss_fn, optimizer, scheduler, num_epochs=25, device=torch.device("cpu")
):
    ...

    # load best model weights
    model.load_state_dict(torch.load(best_model_params_path))

    return model
```

# Transfer Learning with PyTorch

## ◆ Visualization

```
def imshow(input, label=None):  
    """Display image for Tensor."""  
    input = input.numpy().transpose((1, 2, 0))  
  
    mean = np.array([0.485, 0.456, 0.406])  
    std = np.array([0.229, 0.224, 0.225])  
  
    input = std * input + mean  
    input = np.clip(input, 0, 1)  
  
    plt.imshow(input)  
  
    if label is not None:  
        plt.title(label)  
  
    plt.show()
```

# Transfer Learning with PyTorch

## ◆ Visualization

```
def visualize_model(dataloaders, class_names, model, num_images=6):
    model.eval()

    images_so_far = 0

    for i, (inputs, labels) in enumerate(dataloaders['val']):
        outputs = model(inputs)
        _, preds = torch.max(outputs, dim=1)

        for j in range(inputs.size()[0]):
            imshow(
                inputs.data[j], label="Prediction: {0} - Label: {1}".format(
                    class_names[preds[j].item()], class_names[labels[j].item()]
                )
            )
        images_so_far += 1
    if images_so_far == num_images:
        return
```



# Transfer Learning with PyTorch

## ◆ Get Model

```
def get_model(method, device=torch.device("cpu")):  
    model_ft = models.resnet18(weights='IMAGENET1K_V1')
```

```
    if method == "frozen_and_train_new_classifier":  
        for param in model_ft.parameters():  
            param.requires_grad = False
```

We need to freeze all the network except the final layer.  
We need to set `requires_grad = False` to freeze the parameters so that the gradients are not computed in `backward()`

```
    print(model_ft)  
    print("#" * 100)
```

```
    # Here the size of each output sample is set to 2  
    num_ftrs = model_ft.fc.in_features  
    model_ft.fc = nn.Linear(in_features=num_ftrs, out_features=2)
```

```
    print(model_ft)  
    print("#" * 100)
```

```
    model_ft = model_ft.to(device)  
    return model_ft
```

# Transfer Learning with PyTorch

## ◆ Main

```
def main(method):
    device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

    image_datasets, dataset_sizes, dataloaders, class_names = get_new_data()
    model_ft = get_model(method, device)
    loss_fn = nn.CrossEntropyLoss()

    # Observe that all parameters are being optimized
    optimizer_ft = optim.Adam(model_ft.parameters(), lr=0.001)

    # Decay LR by a factor of 0.1 every 7 epochs
    exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)

    model_ft = train_model(
        dataloaders, dataset_sizes, model_ft, loss_fn, optimizer_ft, exp_lr_scheduler,
        num_epochs=25, device=device
    )
    visualize_model(dataloaders, class_names, model_ft)
```

# Transfer Learning with PyTorch

## ◆ Main

```
if __name__ == "__main__":
    method_idx = 0
    methods = [
        "frozen_and_train_new_classifier",
        "fine_tune_the_whole_model"
    ]
    main(methods[method_idx])
```

...

...

Epoch 24/24

-----

train Loss: 0.4411 Acc: 0.7951

val Loss: 0.2509 Acc: 0.9346

Training complete in 2m 28s

Best val Acc: 0.941176

```
if __name__ == "__main__":
    method_idx = 1
    methods = [
        "frozen_and_train_new_classifier",
        "fine_tune_the_whole_model"
    ]
    main(methods[method_idx])
```

...

...

Epoch 24/24

-----

train Loss: 0.6055 Acc: 0.6721

val Loss: 0.5691 Acc: 0.6732

Training complete in 5m 26s

Best val Acc: 0.699346